

Uncovering Library Features from API Usage on Stack Overflow

The 29th IEEE International Conference on Software Analysis, Evolution and Reengineering

Camilo Velázquez-Rodríguez, Eleni Constantinou and Coen De Roover

March 15th, 2022

Introduction

JSoup library

```
String text = EntityUtils.toString(getResponse.getEntity());
Document doc = Jsoup.parse(text);
Element form = doc.select("form").first();

for (Element input : form.select("input")) {
    String value = input.attr("value");
}
```

Introduction

JSoup library

```
String text = EntityUtils.toString(getResponse.getEntity());  
Document doc = Jsoup.parse(text);  
Element form = doc.select("form").first();  
  
for (Element input : form.select("input")) {  
    String value = input.attr("value");  
}
```

Aspect Oriented

Actor Frameworks

Application Metrics

Build Tools

Bytecode Libraries

Command Line Parsers

Cache Implementations

Cloud Computing

Code Analyzers

Collections

Configuration Libraries

Core Utilities

Date and Time Utilities

Dependency Injection

Embedded SQL Databases

HTML Parsers

HTTP Clients

I/O Utilities

...

Introduction

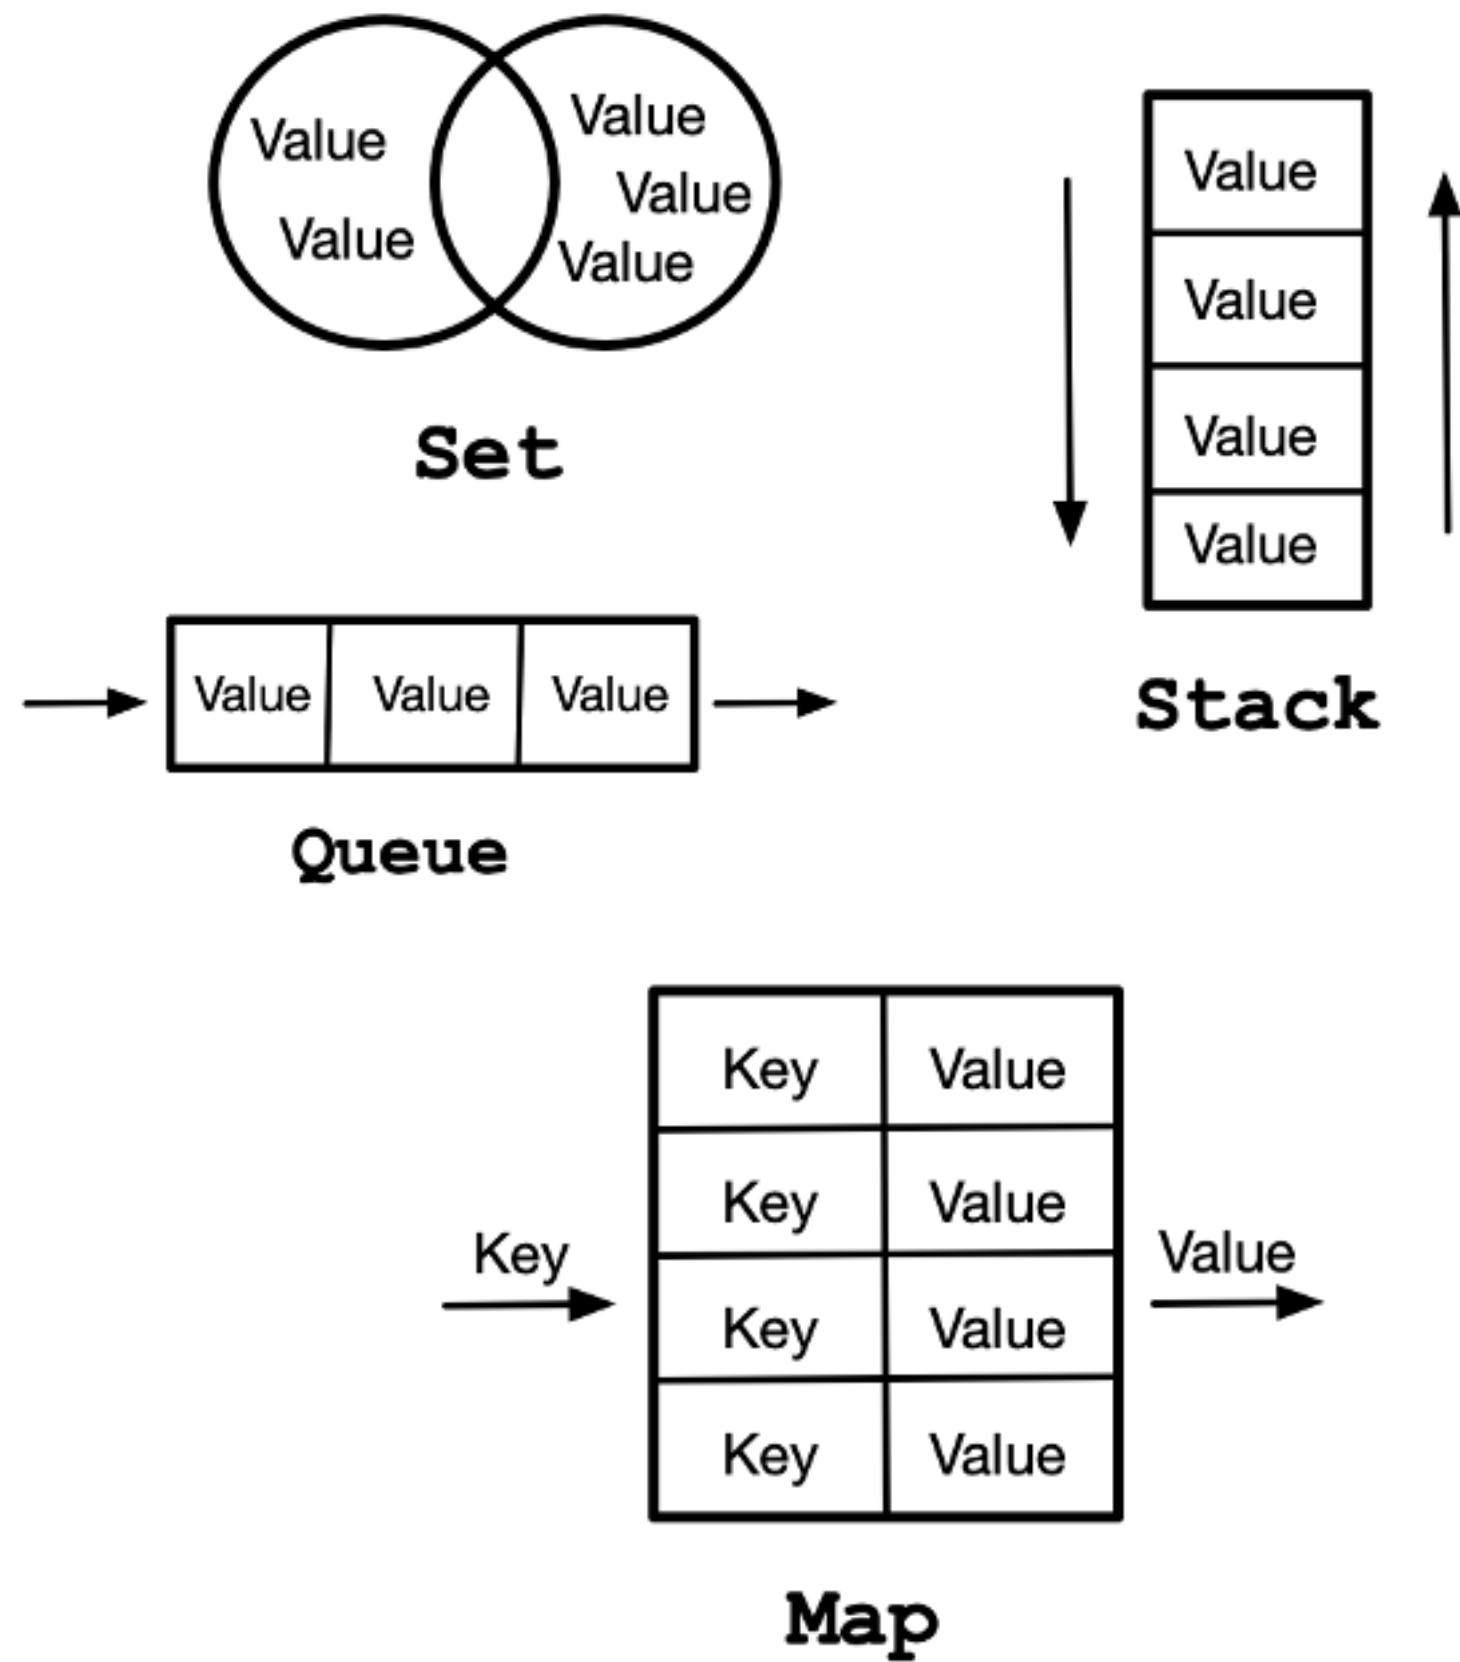
Collections (72)

<div>1. Apache Commons Collections</div> <div>commons-collections » commons-collections</div> <div>Types that extend and augment the Java Collections Framework.</div> <div>Last Release on Nov 12, 2015</div>	<div>5,836 usages</div> <div>Apache</div>
<div>2. Apache Commons Collections</div> <div>org.apache.commons » commons-collections4</div> <div>The Apache Commons Collections package contains types that extend and augment the Java Collections Framework.</div> <div>Last Release on Jul 9, 2019</div>	<div>3,304 usages</div> <div>Apache</div>
<div>3. Fastutil</div> <div>it.unimi.dsi » fastutil</div> <div>fastutil extends the Java Collections Framework by providing type-specific maps, sets, lists, and queues with a small memory footprint and fast access and insertion; it provides also big (64-bit) arrays, sets and lists, sorting algorithms, fast, practical I/O classes for binary and text files, and facilities for memory mapping large files. Note that if you have both this jar and fastutil-core.jar in your dependencies, fastutil-core.jar should be excluded.</div> <div>Last Release on Feb 10, 2022</div>	<div>627 usages</div> <div>Apache</div>
<div>4. GNU Trove</div> <div>net.sf.trove4j » trove4j</div> <div>The Trove library provides high speed regular and primitive collections for Java.</div> <div>Last Release on Jun 4, 2012</div>	<div>542 usages</div> <div>LGPL</div>
<div>5. Google Collections Library</div> <div>com.google.collections » google-collections</div> <div>Google Collections Library is a suite of new collections and collection-related goodness for Java 5.0</div> <div>Last Release on Dec 30, 2009</div>	<div>454 usages</div> <div>Apache</div>

Introduction

Collections (72)

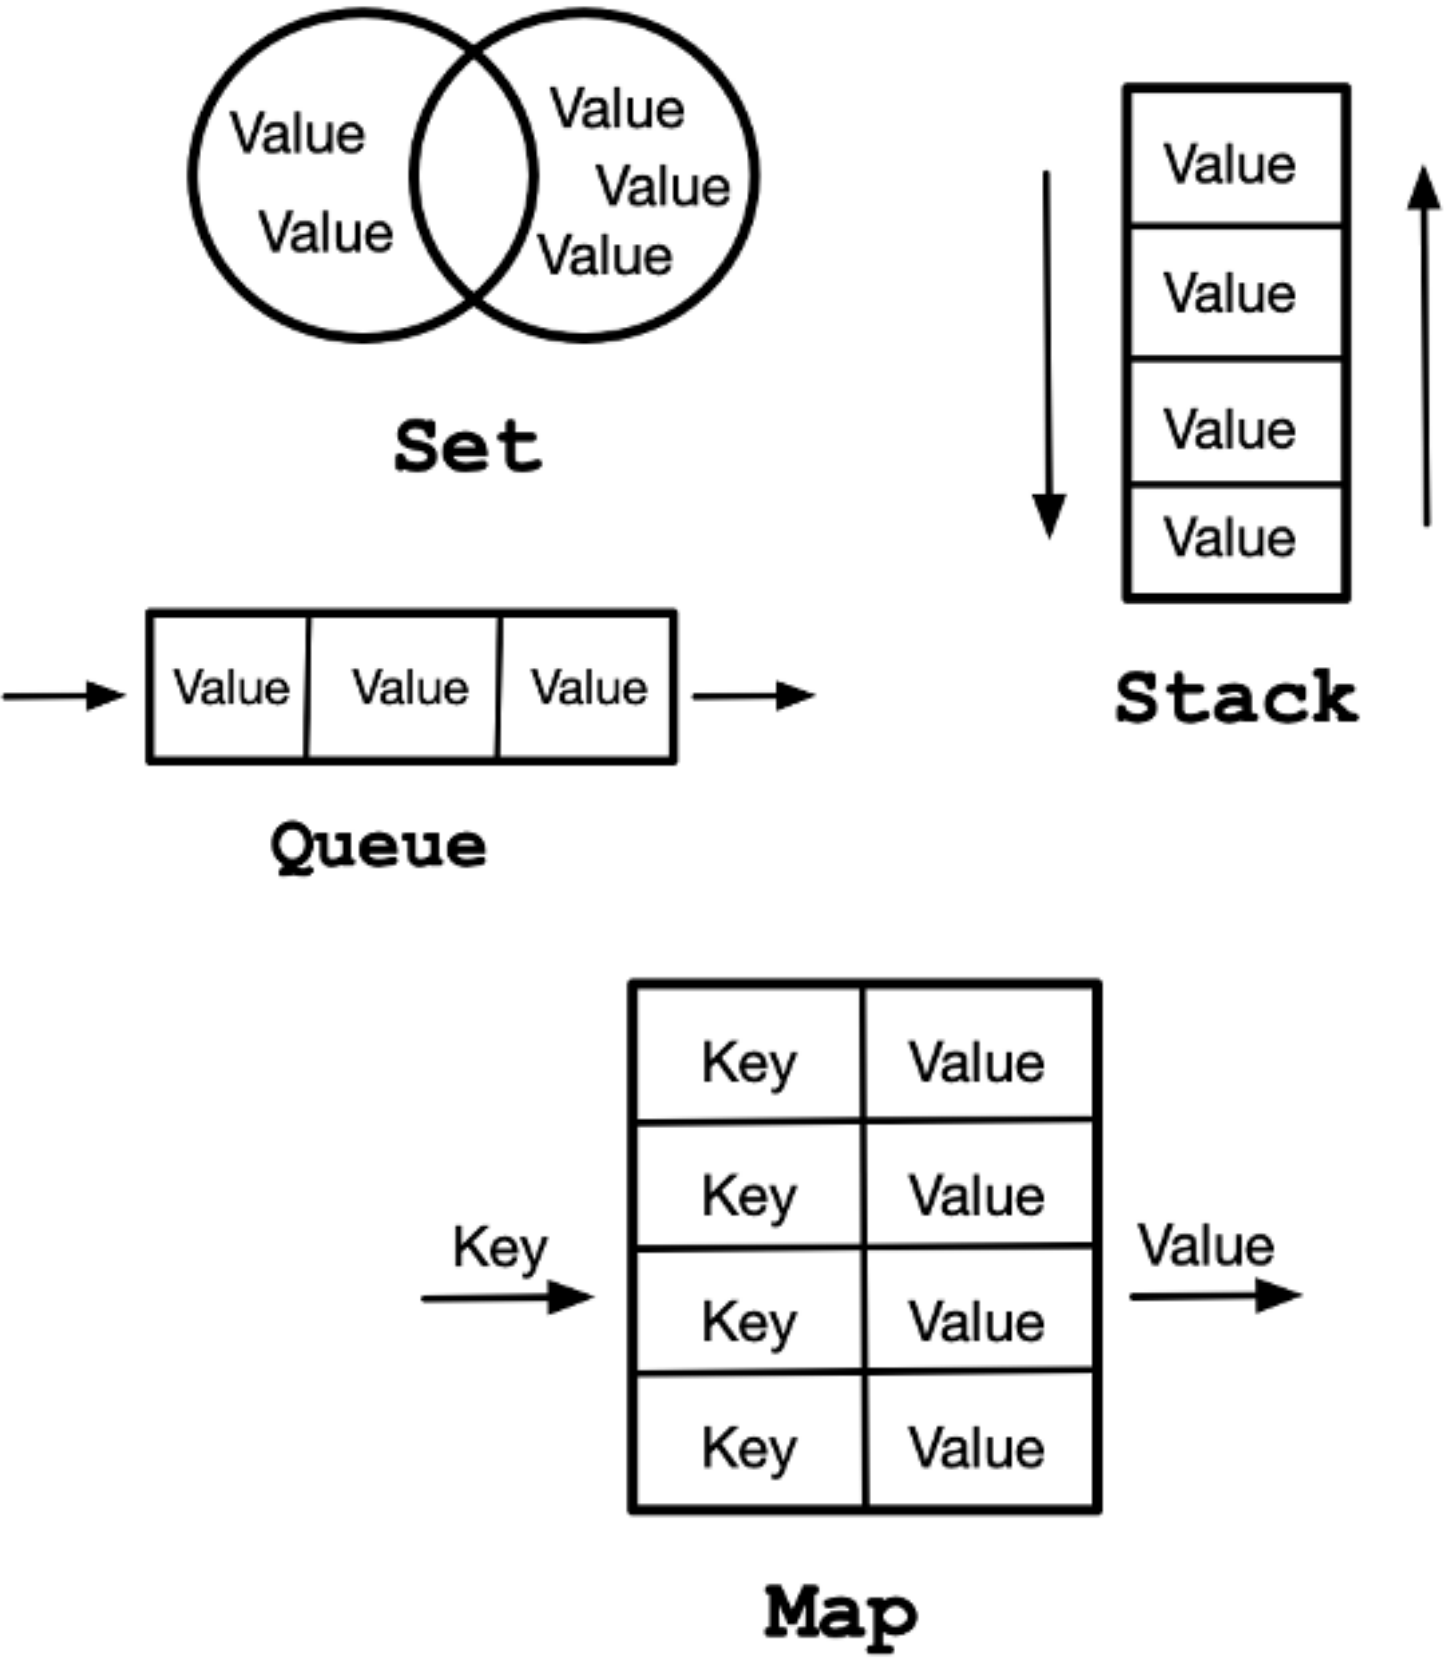
API Features



Introduction

Collections (72)

API Features



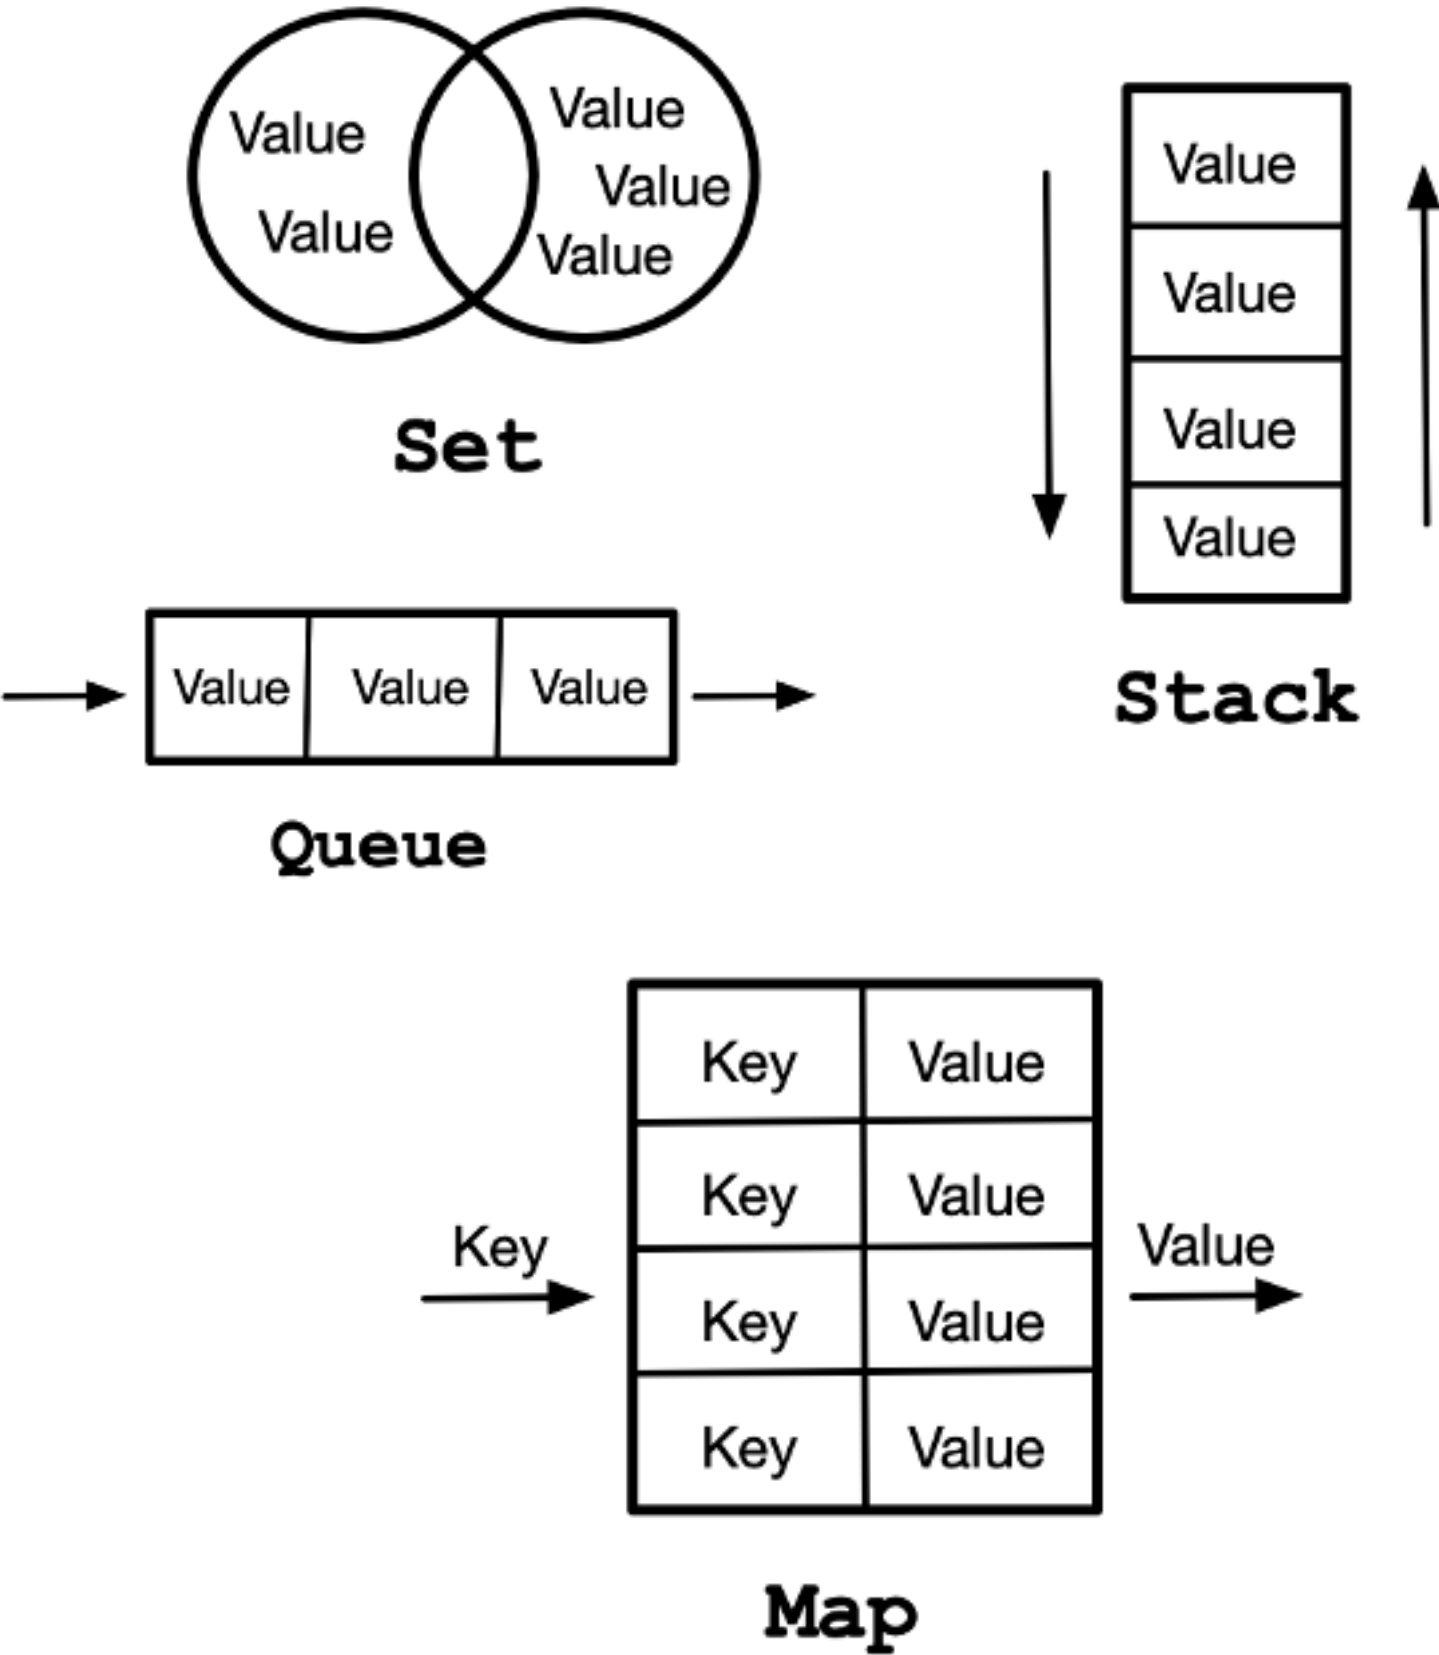
Features

- Persistent or thread-safe?
- Methods for sorting?
- Methods for reversing?
- Filter a collection?
- Transform collections?

Introduction

Collections (72)

API Features



Features

- Persistent or thread-safe?
- Methods for sorting?
- Methods for reversing?
- Filter a collection?
- Transform collections?

Use of Features

- Single method call?
- Group of method calls?
- Static or instance method calls?
- How do libraries compare feature-wise?

Introduction

Feature Identification: A Novel Approach and a Case Study

Giuliano Antoniol^{1,3}

Yann-Gaël Guéhéneuc³

¹RCOST, University of Sannio, Italy

²GEODES, DIRO, University of Montreal, Canada

³Département de Génie Informatique, École Polytechnique de Montréal, Canada

antoniol@ieee.org, guehene@iro.umontreal.ca

Abstract

Feature identification is a well-known technique to identify subsets of a program source code activated when exercising a functionality. Several approaches have been proposed to identify features. We present an approach to feature identification and comparison for large object-oriented multi-threaded programs using both static and dynamic data. We use processor emulation, knowledge filtering, and probabilistic ranking to overcome the difficulties of collecting dynamic data, i.e., imprecision and noise. We use model transformations to compare and to visualise identified features. We compare our approach with a naive approach and a concept analysis-based approach using a case study on a real-life large object-oriented multi-threaded program, Mozilla, to show the advantages of our approach. We also use the case study to compare processor emulation with sta-

higher-level abstractions beyond those obtained by examining the system itself [4]. We propose an approach to support the recovery of higher-level abstractions through program feature identification and comparison. We define a feature of a program as a set of data structures (i.e., fields and classes) and operations (i.e., functions and methods) participating in the realisation of a user-observable functionality in a given scenario. The scenario details the particular conditions of realisation of the functionality: For example in a web browser, accessing a page from the bookmarks corresponds to one feature, adding a Uniform Resource Locator (URL) to the bookmarks is another feature.

In this paper, we describe our approach of building micro-architectures, subsets of a program architecture [10], from feature identification using static and dynamic data. We assume that the program source code is available and

“... a set of data structures (i.e., fields and classes) and operations (i.e., functions and methods) participating in the realisation of a (...) functionality.”

Semi-Automatically Extracting Features from Source Code of Android Applications

Tetsuya KANDA^{†a)}, Yuki MANABE^{†b)}, Nonmembers, Takashi ISHIO^{†c)}, Member, Makoto MATSUSHITA^{†d)}, Nonmember, and Katsuro INOUE^{†e)}, Fellow

SUMMARY It is not always easy for an Android user to choose the most suitable application for a particular task from the great number of applications available. In this paper, we propose a semi-automatic approach to extract feature names from Android applications. The case study verifies that we can associate common sequences of Android API calls with feature names.

key words: Android, feature extraction, software categorization, API

1. Introduction

Android is one of the most popular platforms for mobile phones and tablets. A user can search and choose from more than 600,000 Android applications in Google Play [1]. Because there are so many choices, however, selecting an appropriate application is not a trivial task. For example, in November 2012, at least 1,000 applications could be found when searching with the keyword “calculator” on Google Play.

A simple but important criterion for selection of an application is the set of features it provides. Investigating the features by trying each application, however, is time consuming. Although documentation is an important source of information, many applications are less than adequate in this area.

MUDABlue [2] and LACT [3] are the solutions that en-

though software developers can use arbitrary sequences of API calls, we hypothesize that a popular feature of an application is likely to be implemented by the same sequence of API calls, since similar applications use the same set of APIs [4]. In our proposed solution, therefore, we automatically extract common sequences of API calls in two or more applications, and manually associate each of these with a feature name. We use the associations as a knowledge-base. We then automatically extract API calls from other target applications and, using our knowledge-base, output feature names associated with the API calls.

It should be noted that this paper is a revised version of our technical report [5]. The technical report describes the detailed implementation of our method. We compare extracted features from applications in this paper while we compared sequence of API calls of applications directly in the technical report, yet it does not affect the result of the case study.

The next section describes the technical details of our approach. Section 3 shows the results of a case study. In Sect. 4, we conclude this article with a discussion of future work.

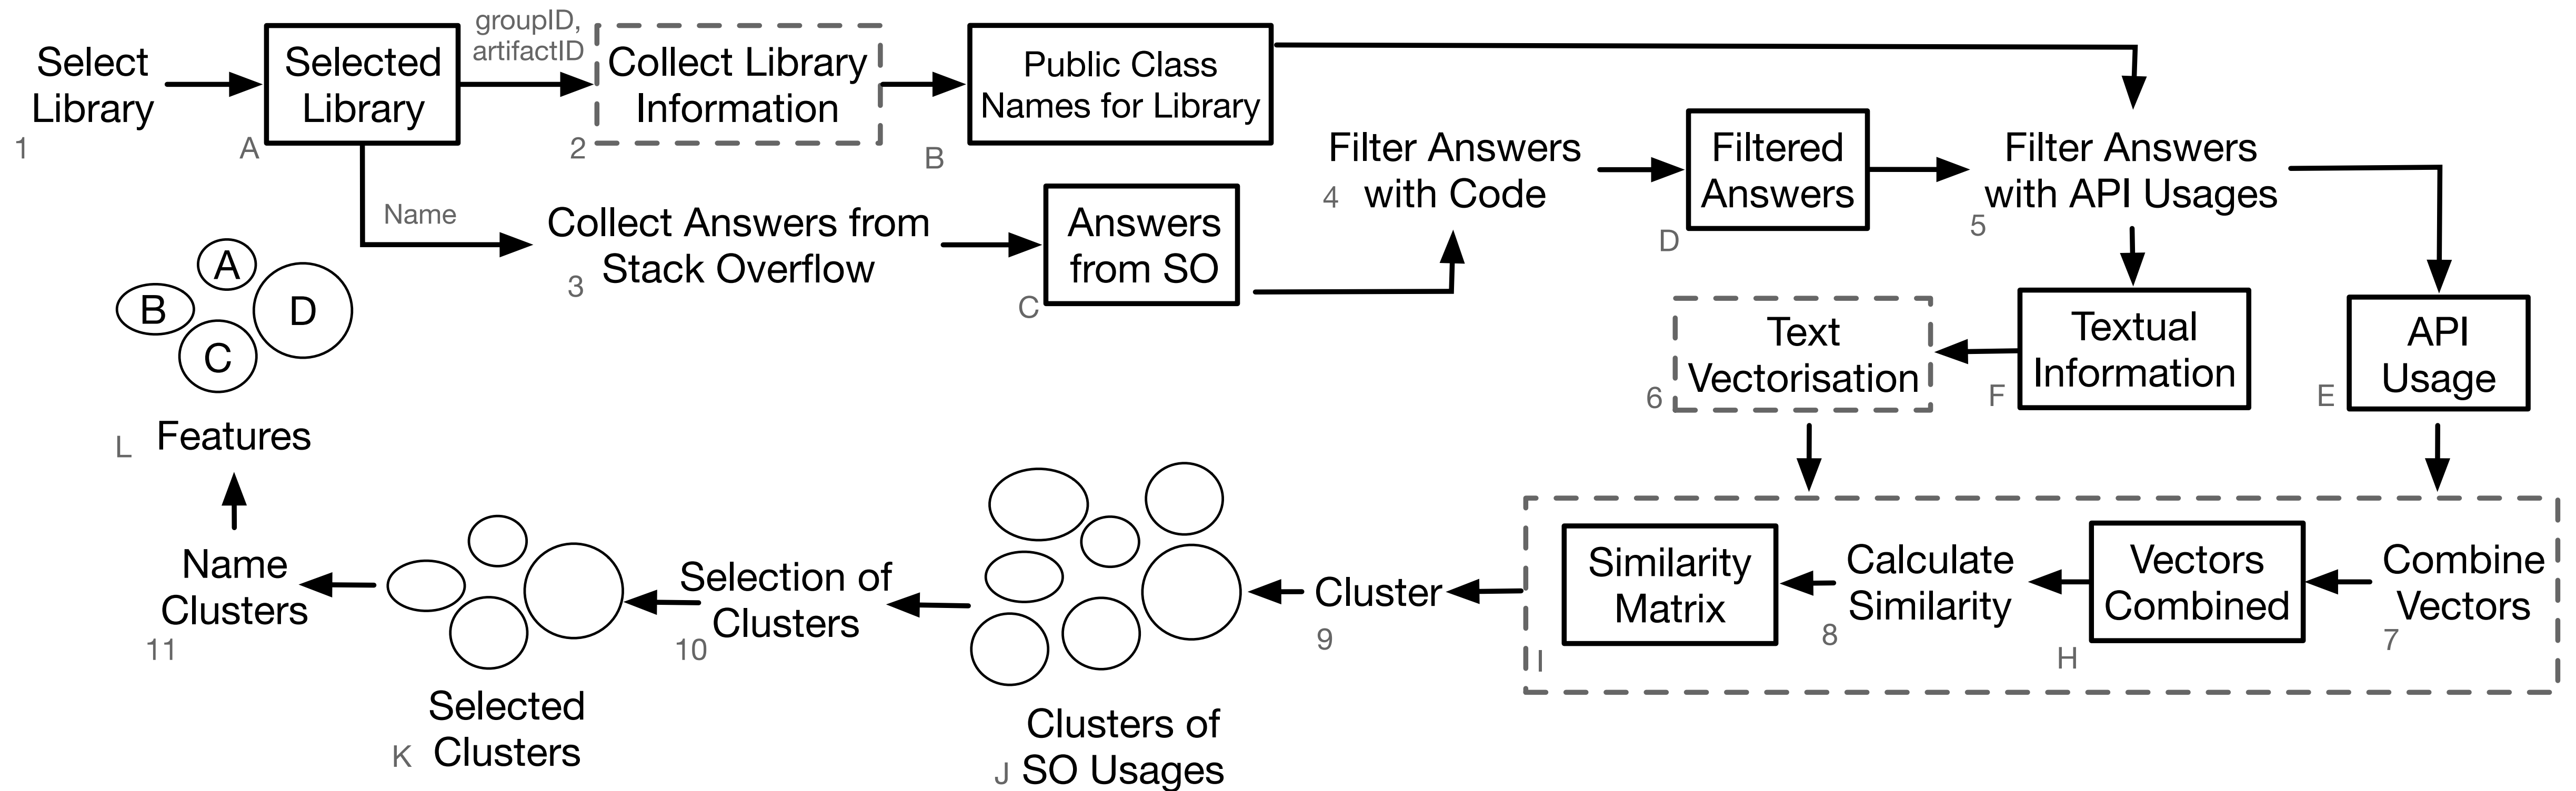
2. Associating API Calls with Feature Names

“... a set of API calls with a corresponding name.”

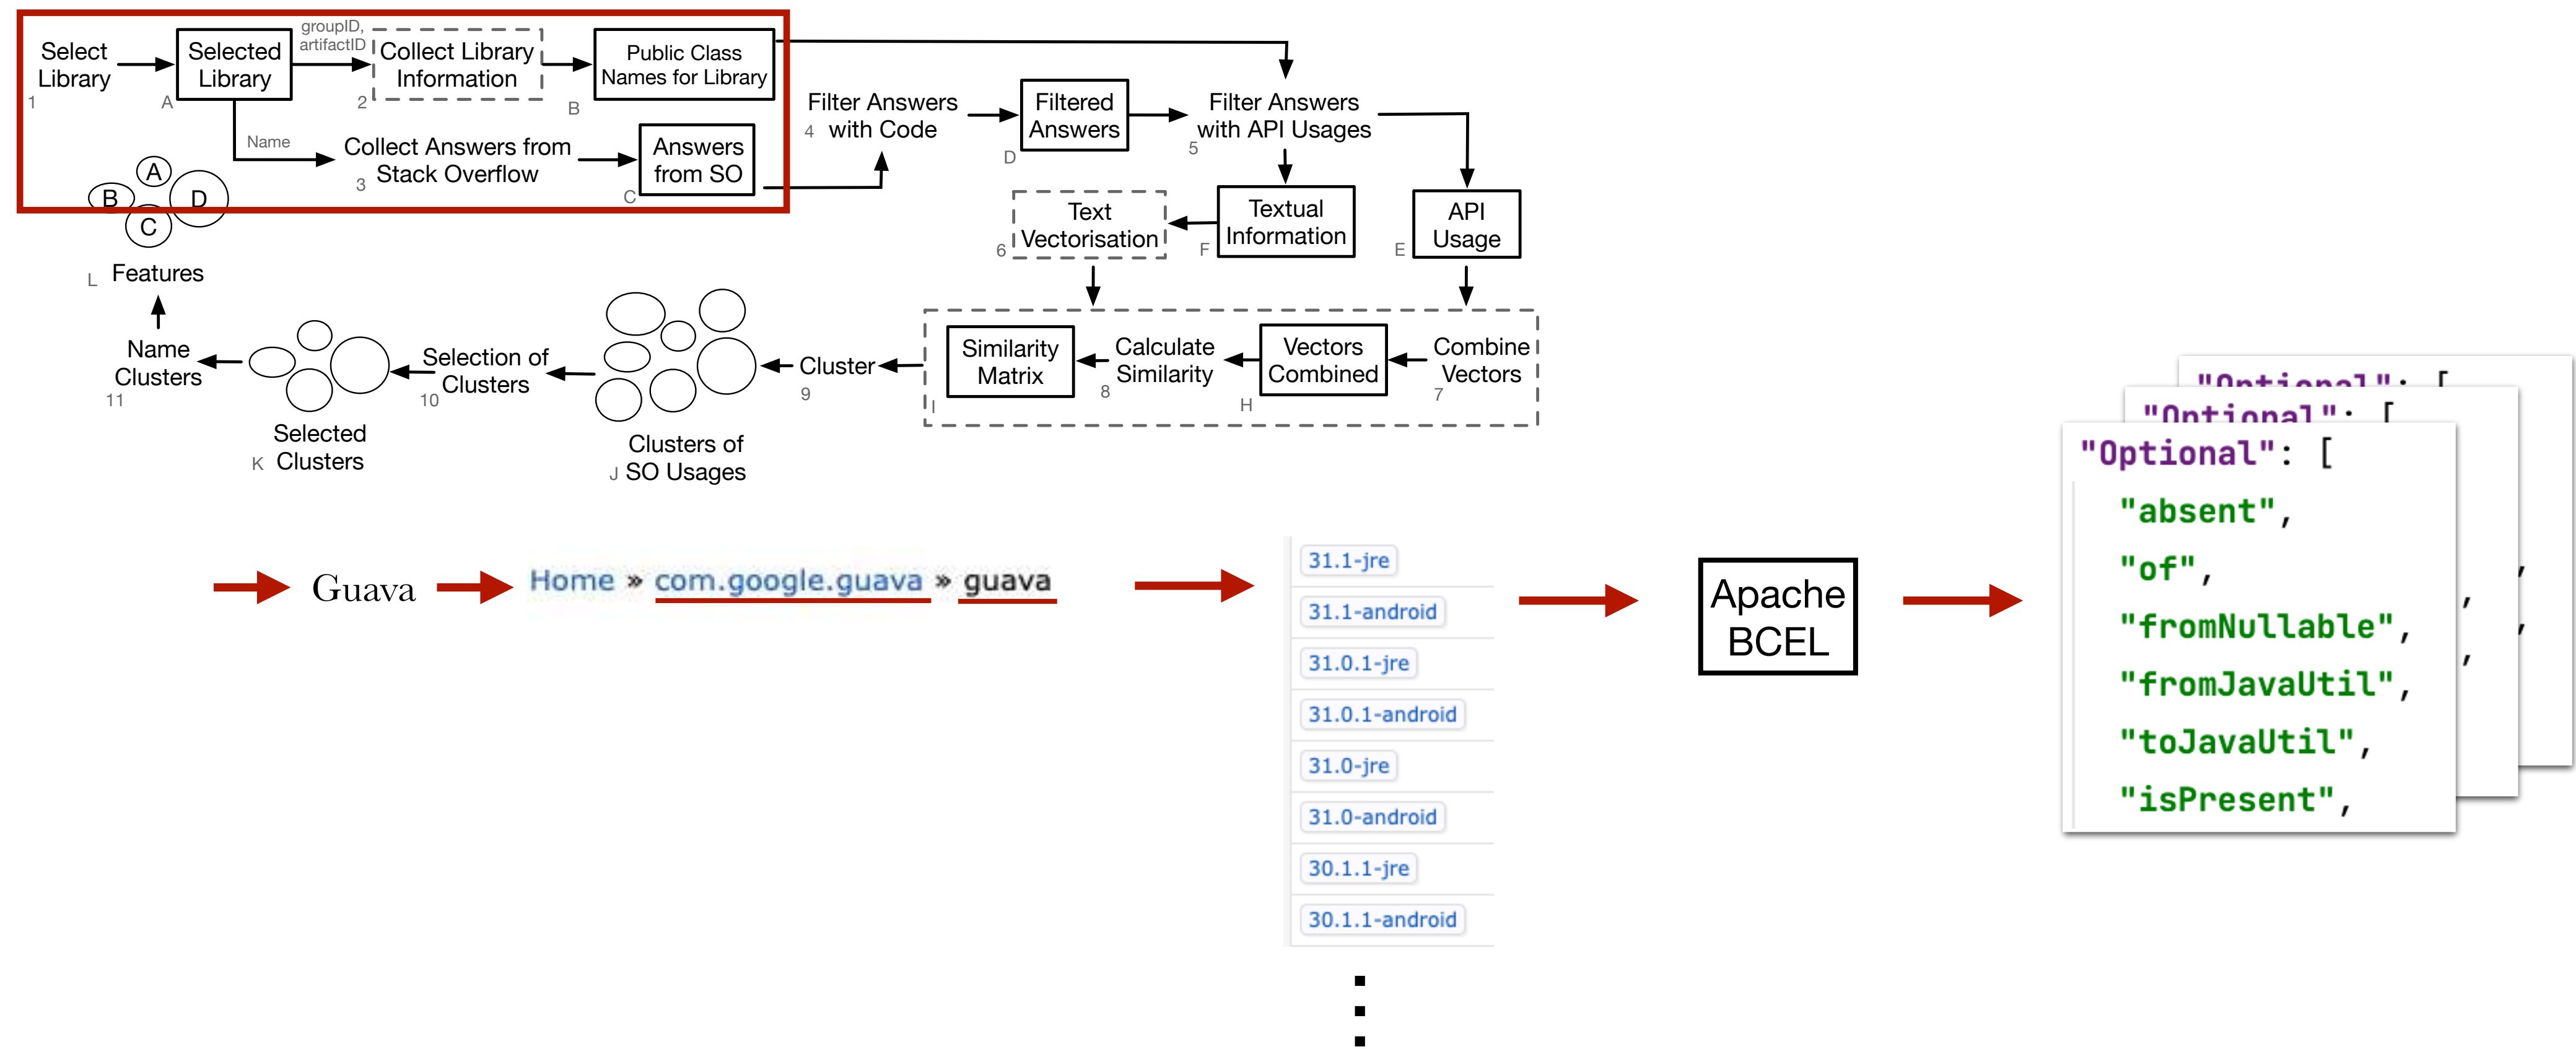
[1] - Antoniol, G. & Guéhéneuc, Y.-G. Feature Identification: A Novel Approach and a Case Study. *21st Ieee Int Conf Softw Maintenance Icsm'05* 1–10 (2005).

[2] - Kanda, T., Manabe, Y., Ishio, T., Matsushita, M. & Inoue, K. Semi-Automatically Extracting Features from Source Code of Android Applications. *Ieice T Inf Syst*, 2857–2859 (2013).

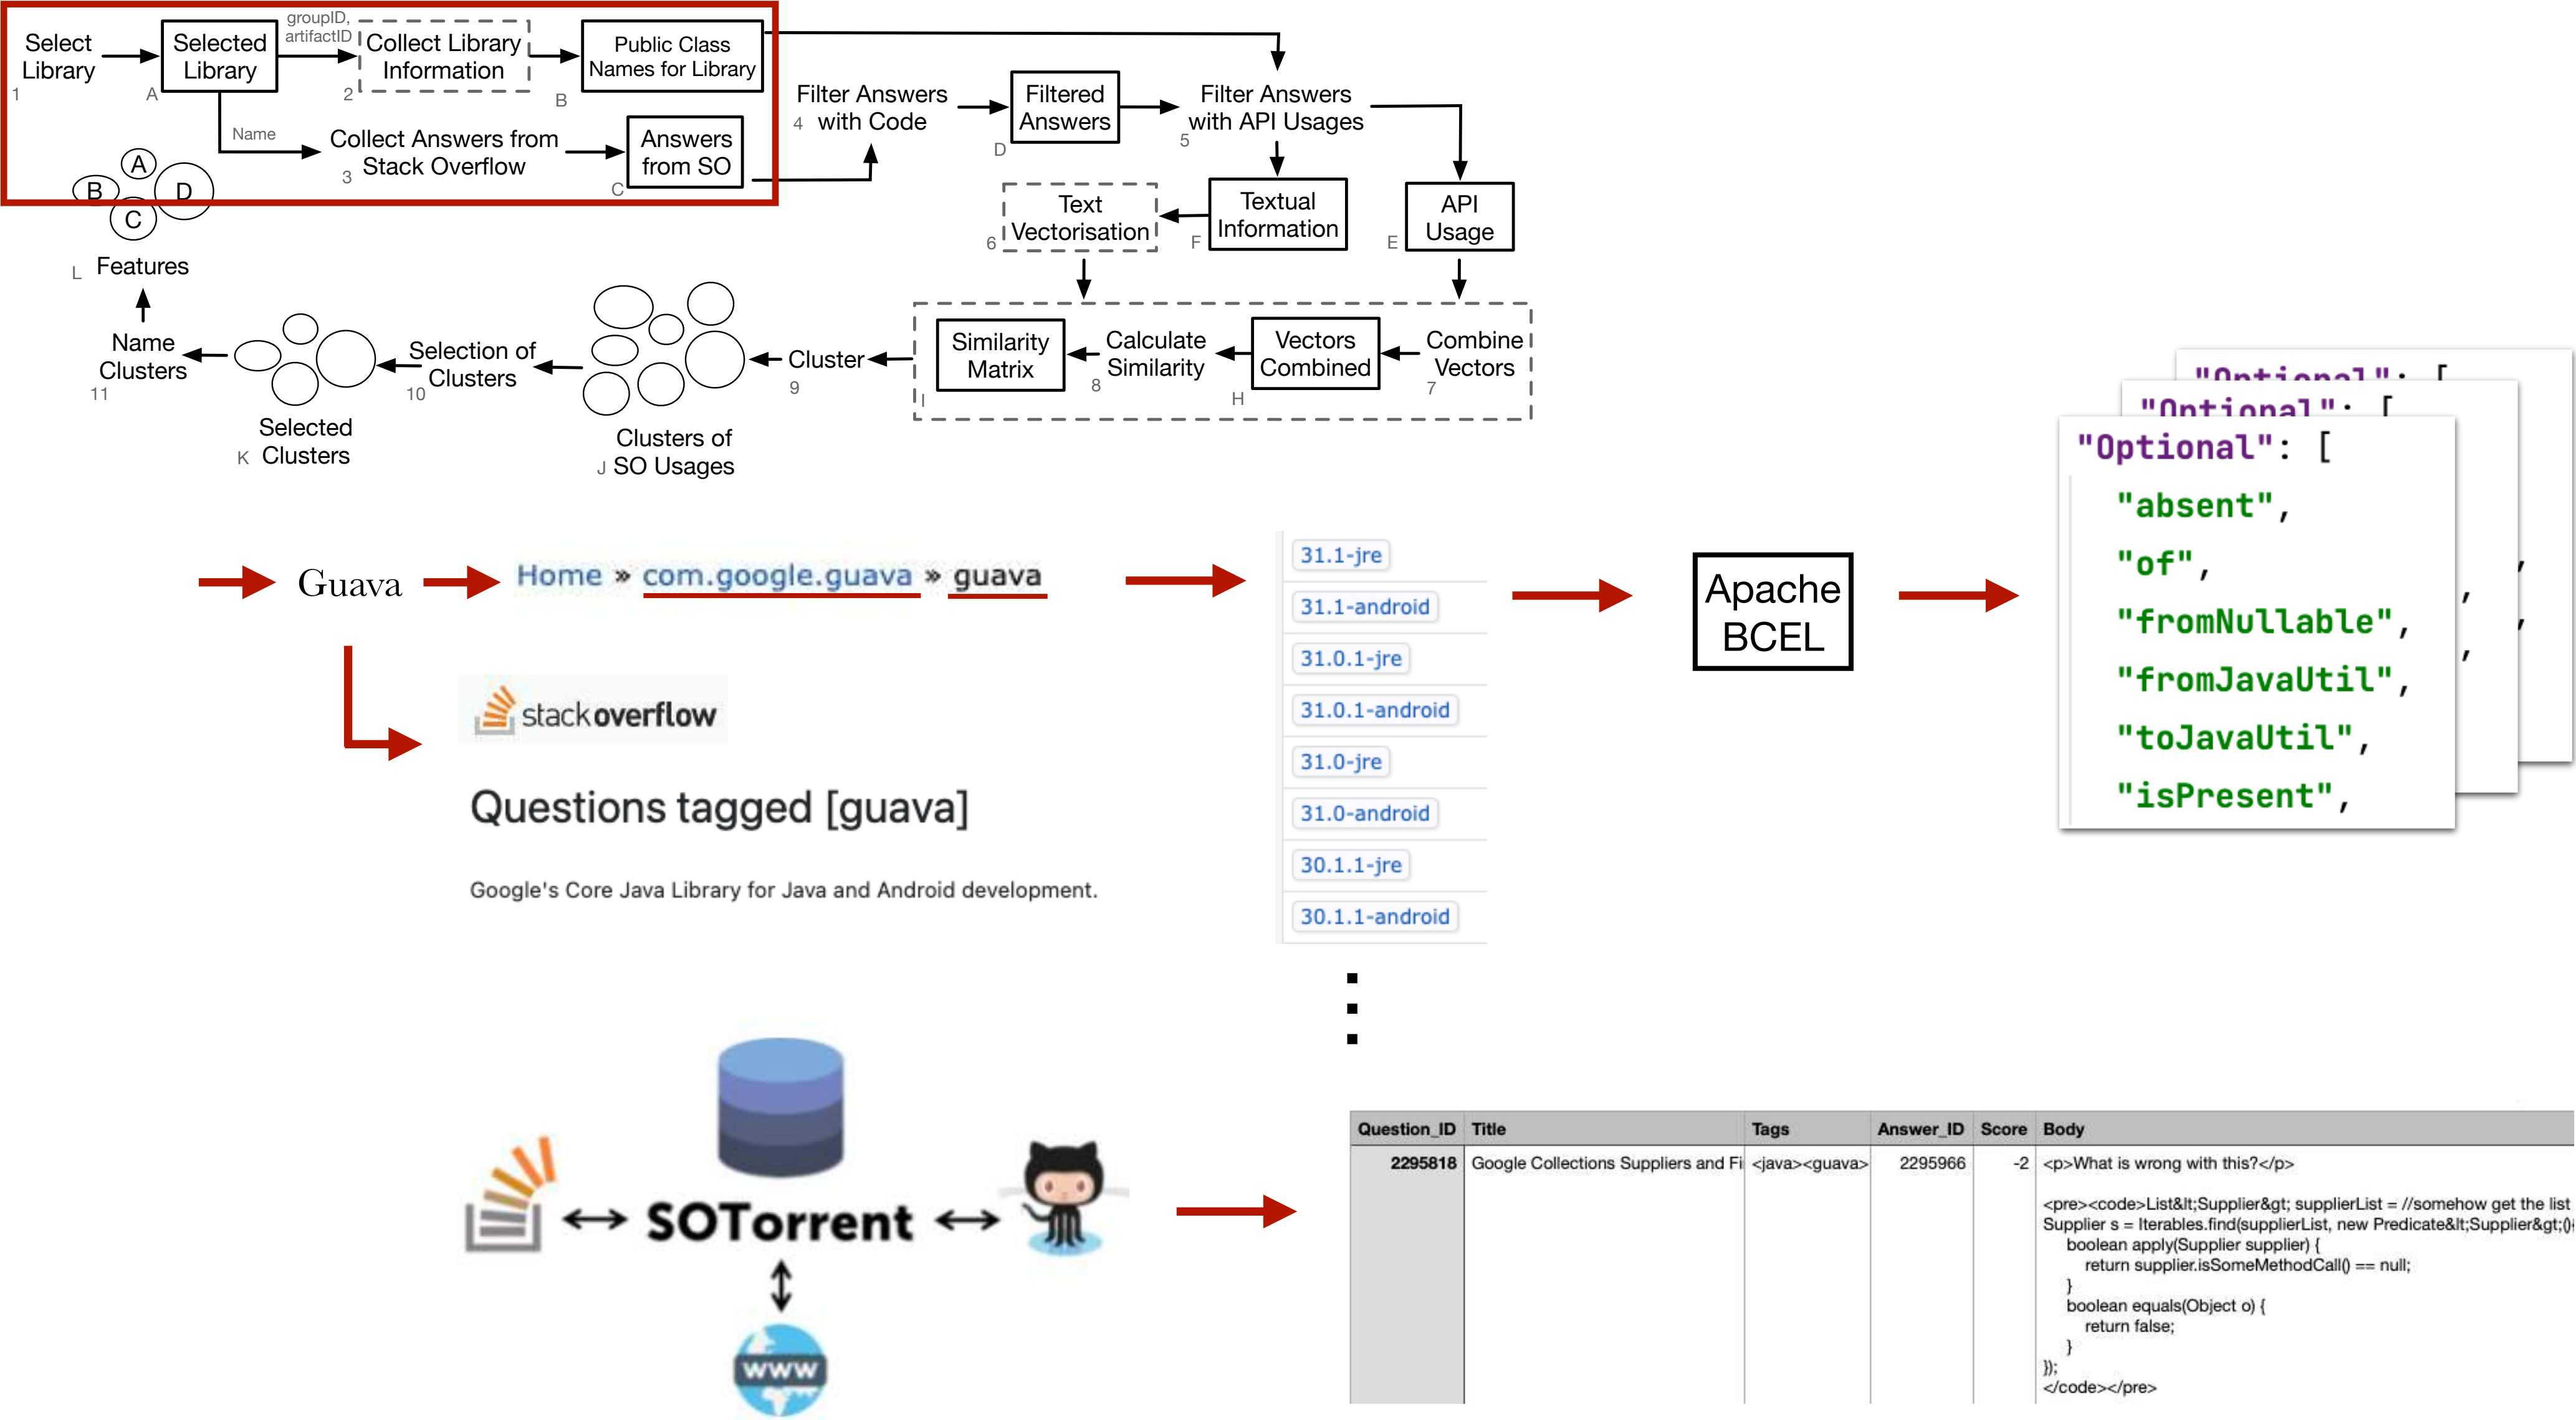
Approach



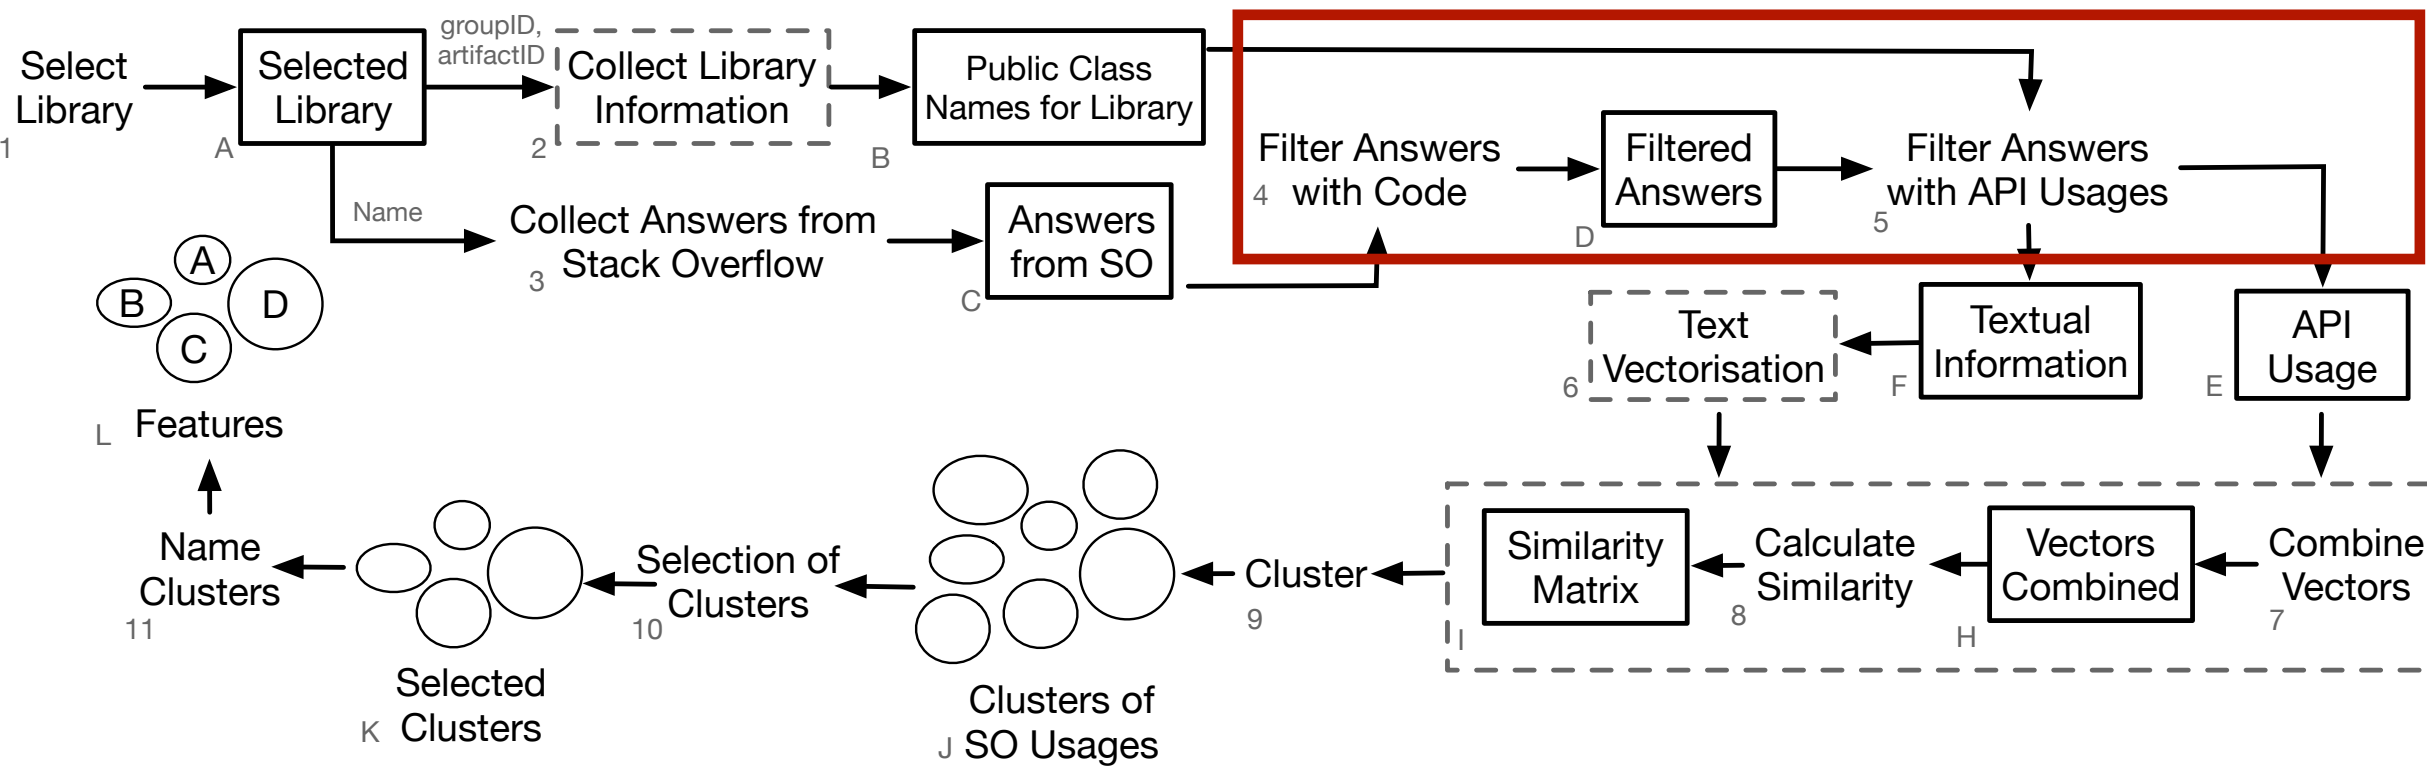
Approach



Approach



Approach

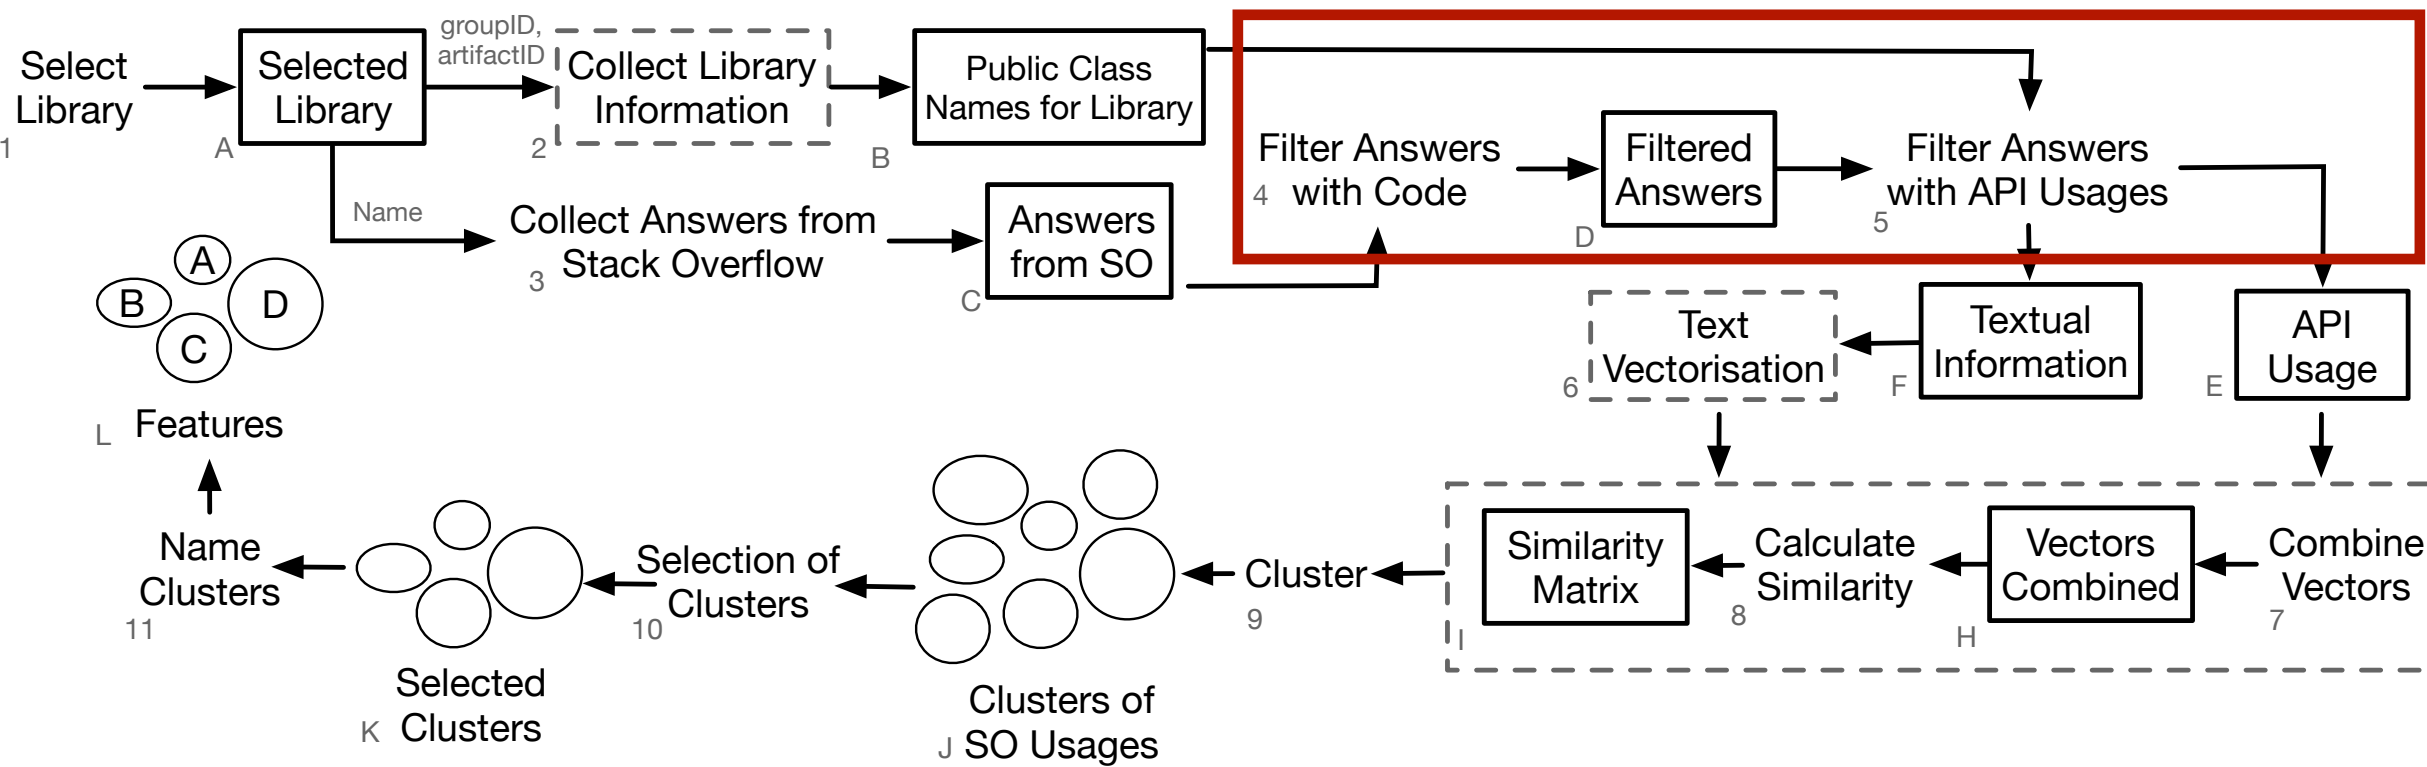


I would have answered in a comment but my reputation still doesn't allow me to:

11 In addition to Rushdi's accepted answer, I want to emphasize that the models which are created for the cross-validation fold sets are all discarded after the performance measurements have been carried out and averaged.

The resulting model is *a/ways* based on the *entire* training set, regardless of your test options. Since M-T-A was asking for an update to the quoted line, here it is: <https://web.archive.org/web/20170519061006/http://list.waikato.ac.nz/pipermail/wekalist/2009-December/046633.html>. It's an answer from one of the WEKA maintainers, pointing out just what I wrote.

Approach



I would have answered in a comment but my reputation still doesn't allow me to:

11 In addition to Rushdi's accepted answer, I want to emphasize that the models which are created for the cross-validation fold sets are all discarded after the performance measurements have been carried out and averaged.

The resulting model is *always* based on the *entire* training set, regardless of your test options. Since M-T-A was asking for an update to the quoted line here it is: <https://web.archive.org/web/20170519061006/http://list.waikato.ac.nz/pipermail/wekalist/2009-December/046633.html>. It's an answer from one of the WEKA maintainers, pointing out just what I wrote.

Change Comparator in MinMaxPriorityQueue guava

The documentation of Guava's MinMaxPriorityQueue state:

... the queue's specified comparator. If no comparator is given at construction time, the natural order of elements is used.

It should be

```
Comparator<Long> comp = Ordering.natural().reverse();
MinMaxPriorityQueue<Long> pq = MinMaxPriorityQueue.orderedBy(comp)
    .maximumSize(11).create();
```

As you can see [in the docs](#) `MinMaxPriorityQueue#orderedBy` is a static method and returns a `Builder`. You should only call that method via the class and it should tell you that the method *doesn't manipulate an instance*.

Guava works with the Builder pattern: Rather than offering methods to change the behaviour of an instance, it offers the possibility to set that behaviour before instantiating it (through a Builder). You set your attributes on the Builder and call `Builder#create` to get the instance.

For more information see [this question](#) and [this one](#) or read on [wikipedia](#).

```
Comparator<Long> comp = Ordering.natural().reverse();
MinMaxPriorityQueue<Long> pq = MinMaxPriorityQueue.orderedBy(comp)
    .maximumSize(11).create();
```

`MinMaxPriorityQueue#orderedBy`

`Builder#create`

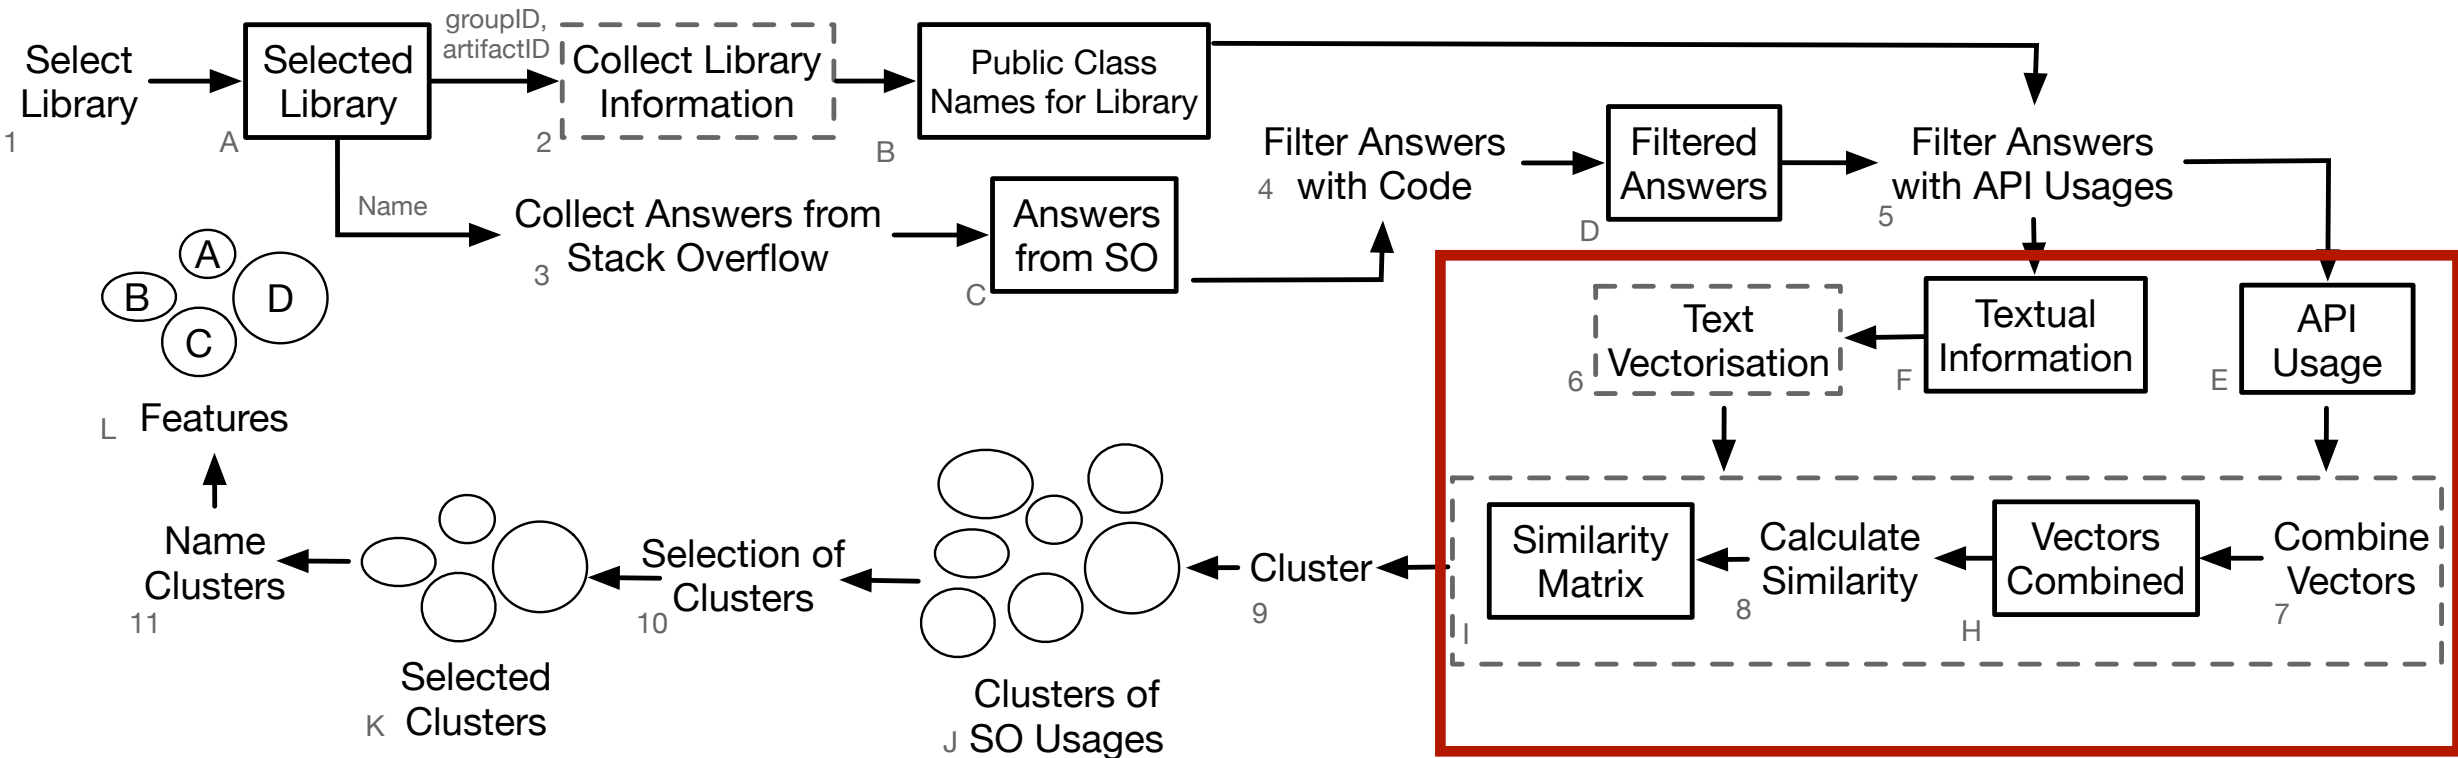
`Ordering.natural().reverse`

`MinMaxPriorityQueue.orderedBy.maximumSize.create`



```
"Optional": [
  "absent",
  "of",
  "fromNullable",
  "fromJavaUtil",
  "toJavaUtil",
  "isPresent",
```


Approach



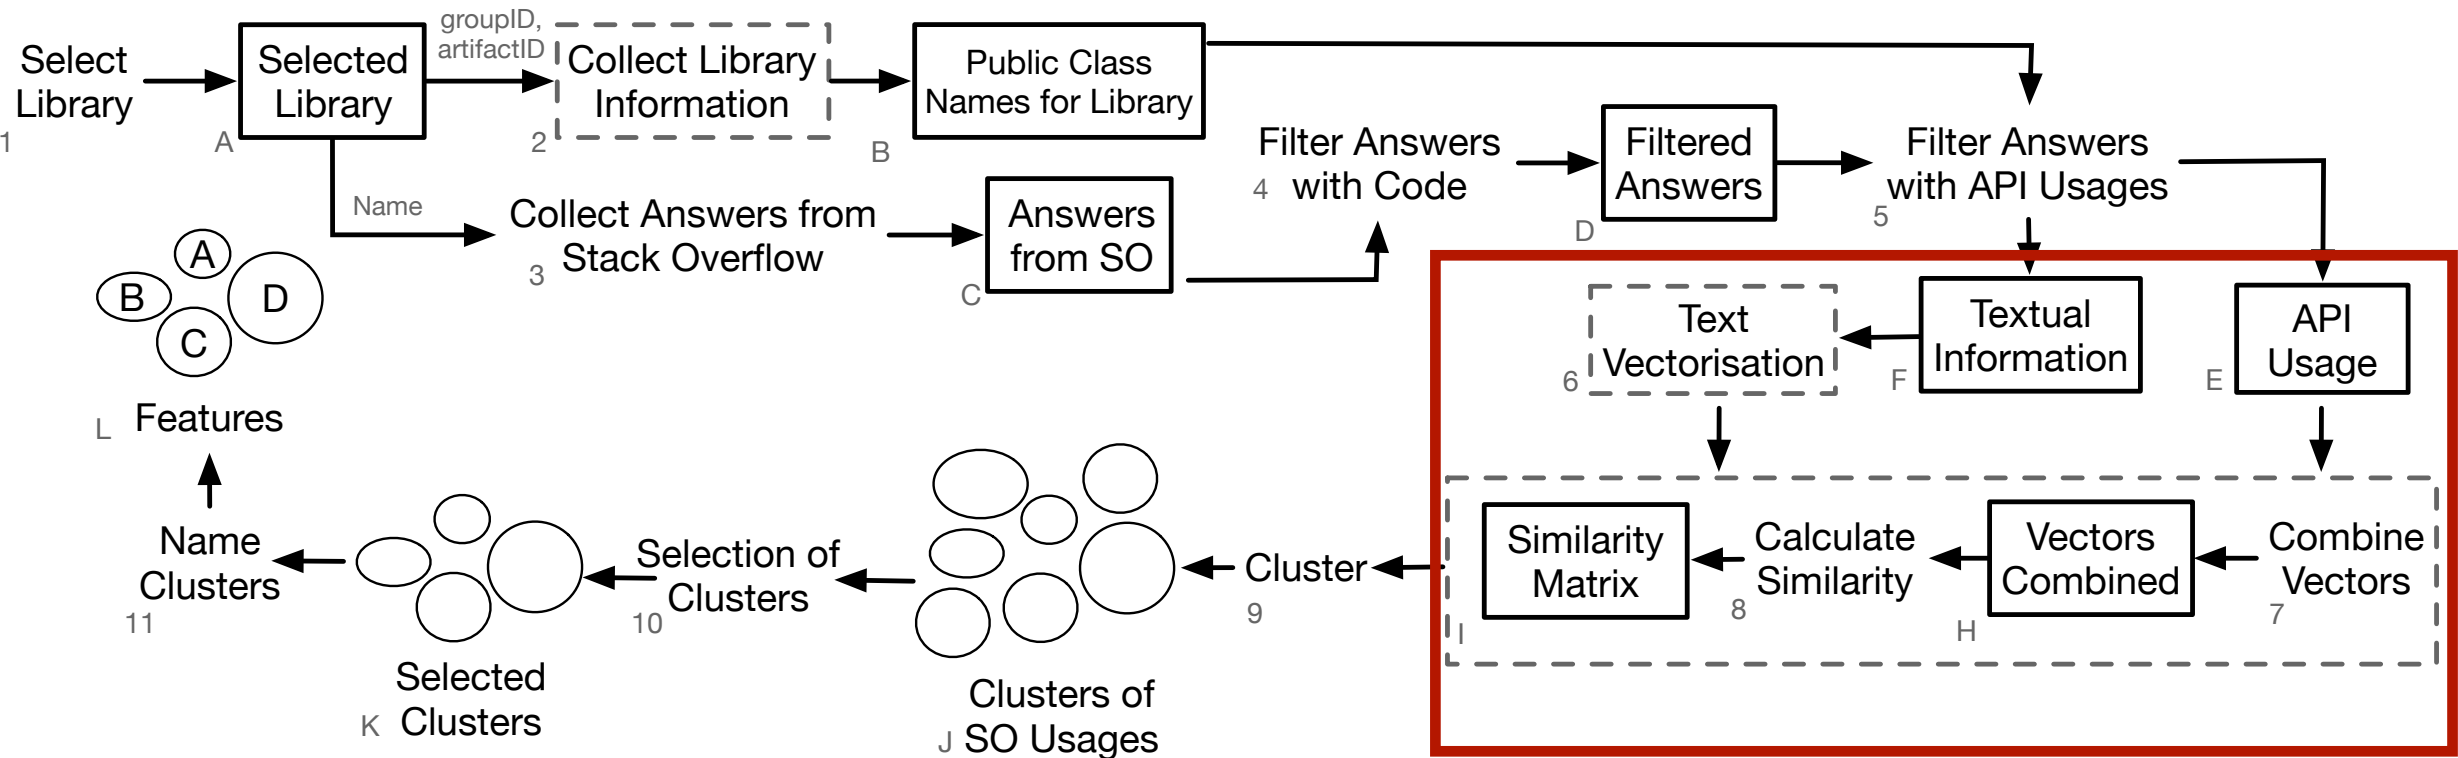
Textual Information

- Post Titles
- Body of the Question
- Body of the Answer

API Usage Information

- Method Names
- Method Names split by Camel Case
- API Calls

Approach

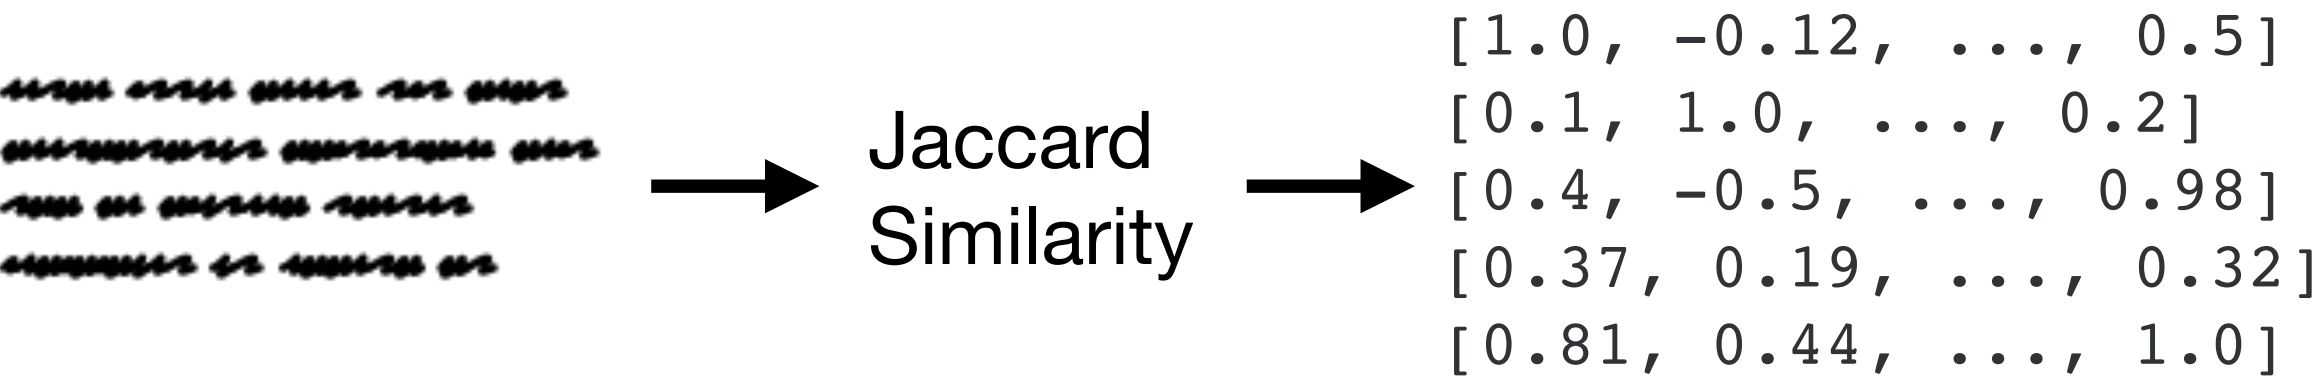
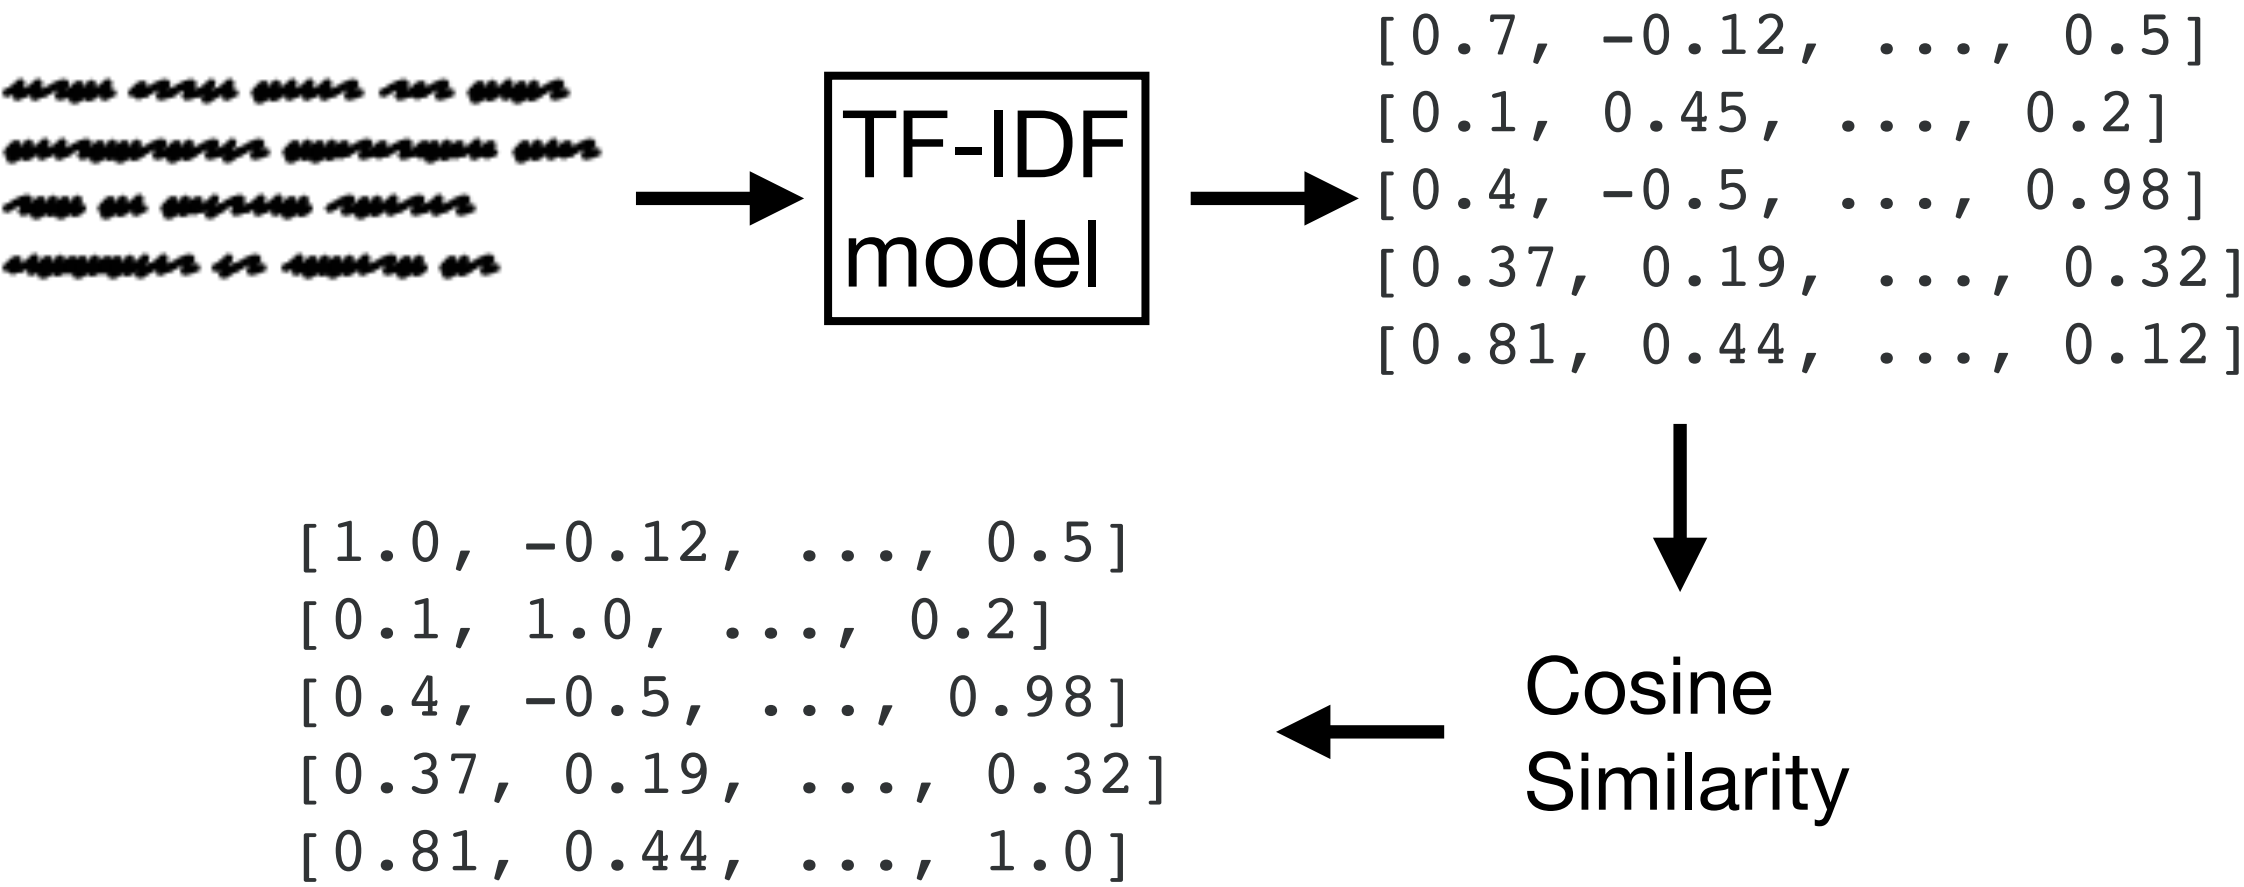


Textual Information

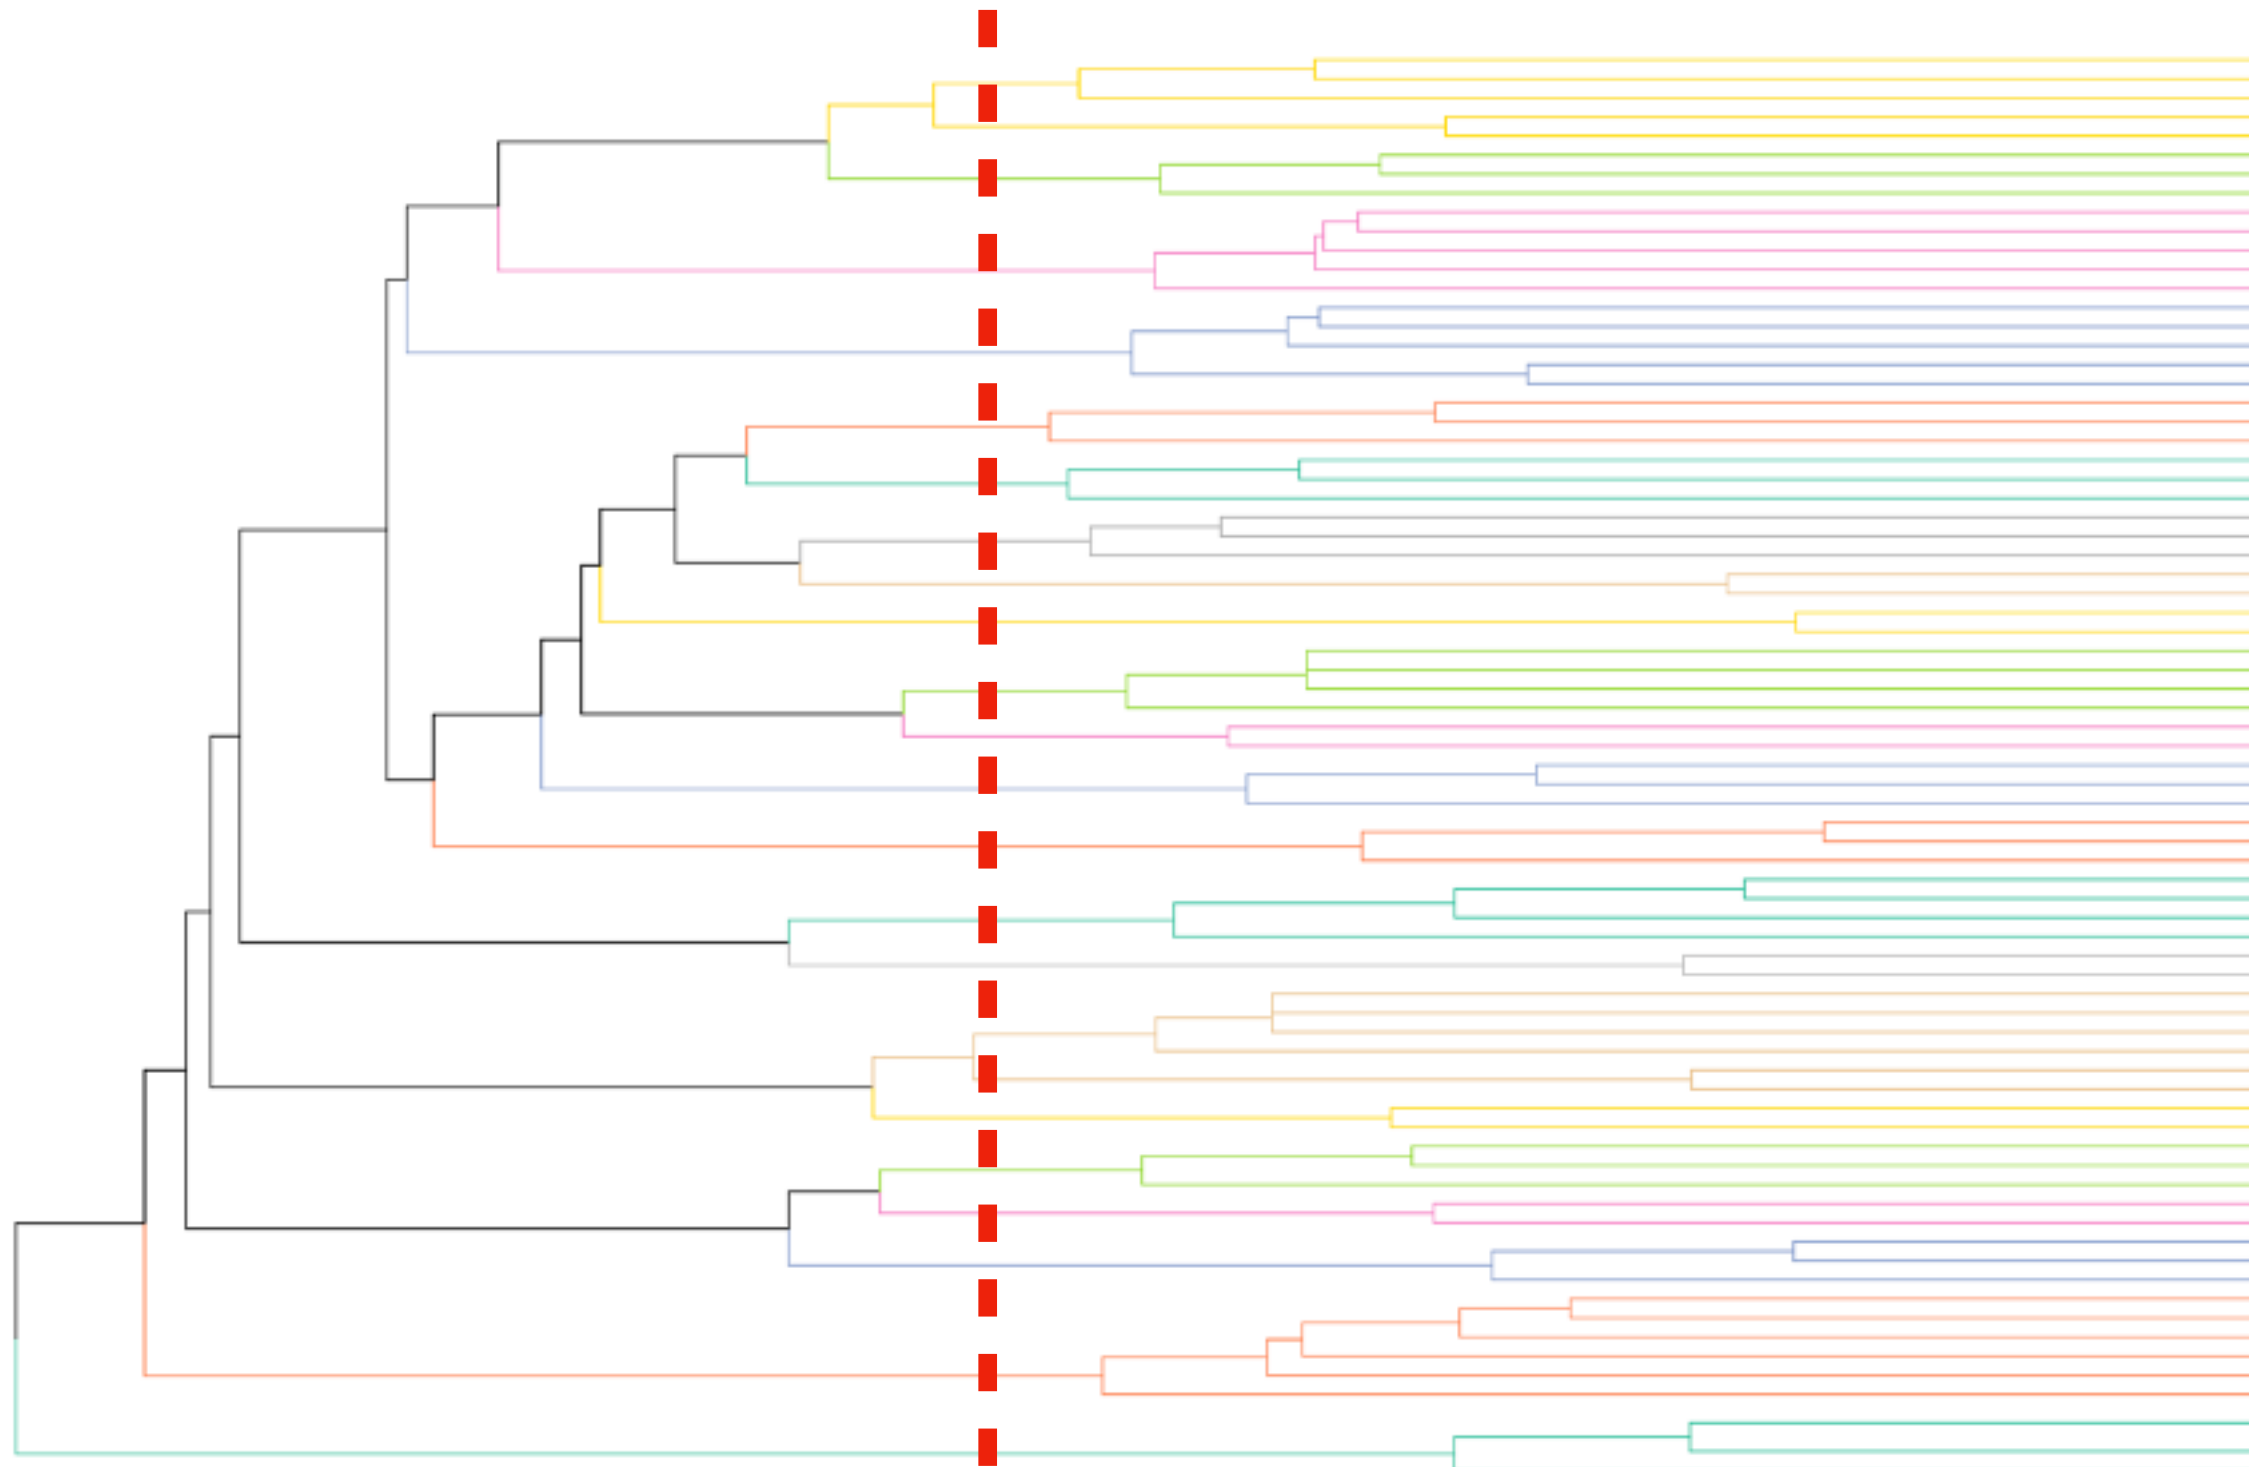
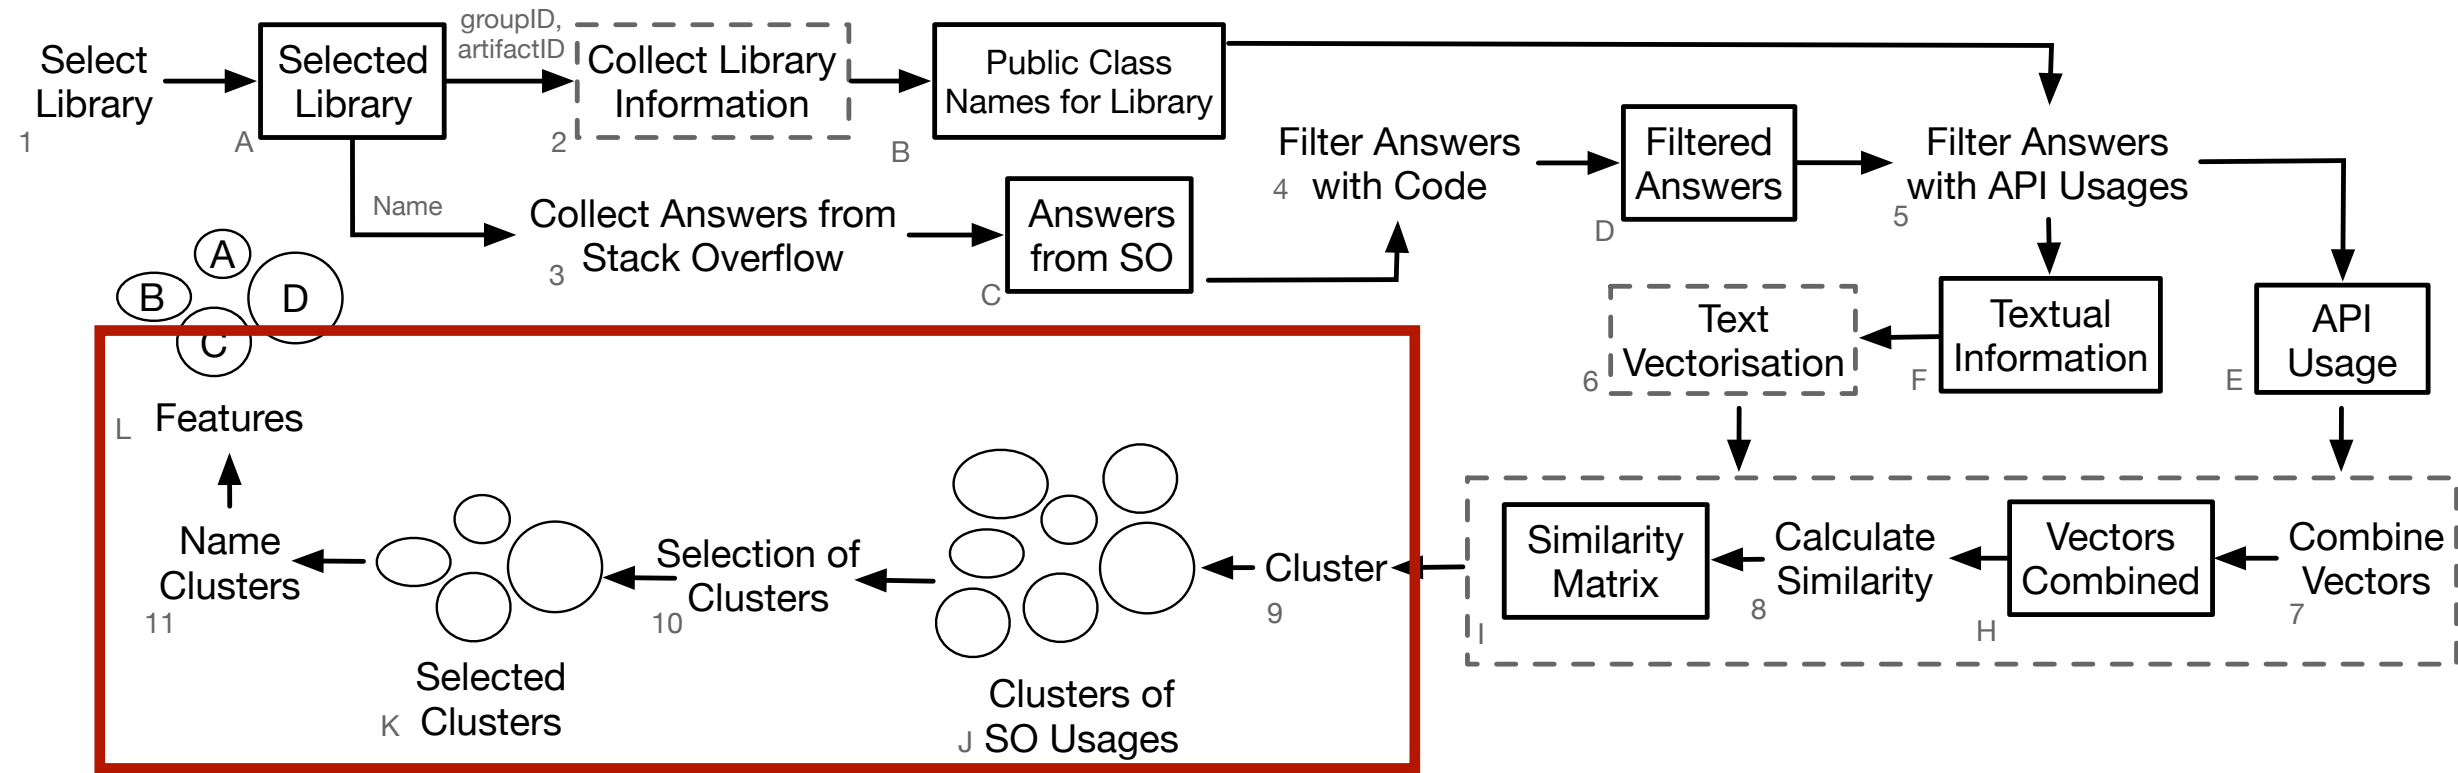
- Post Titles
- Body of the Question
- Body of the Answer

API Usage Information

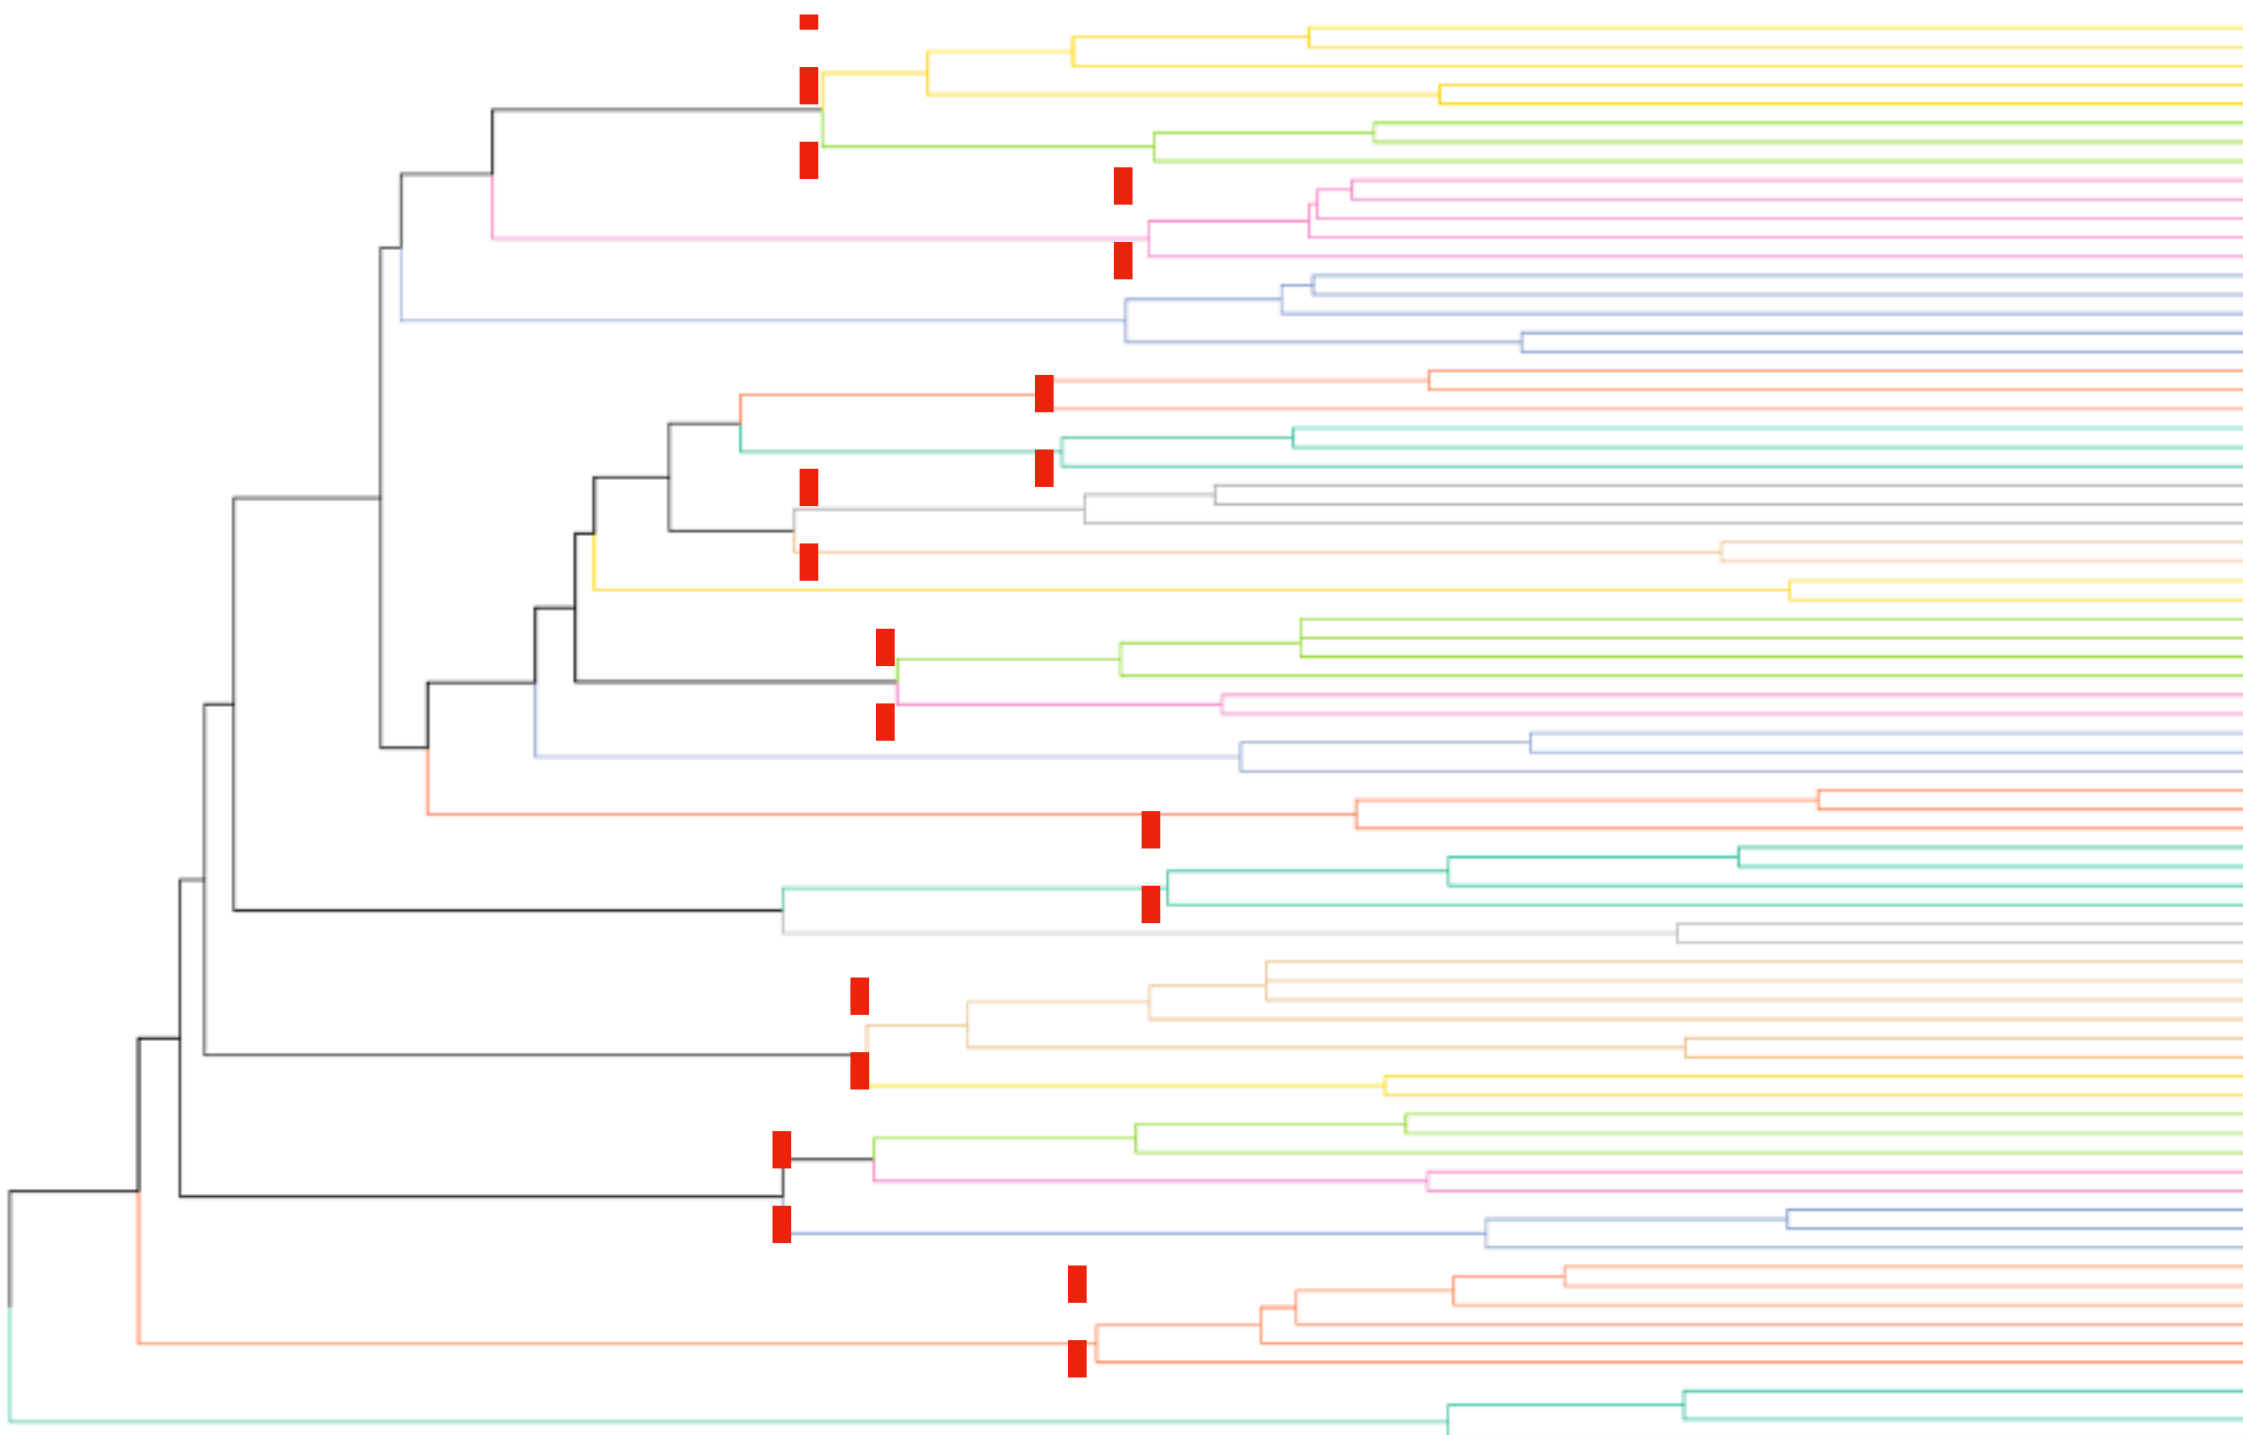
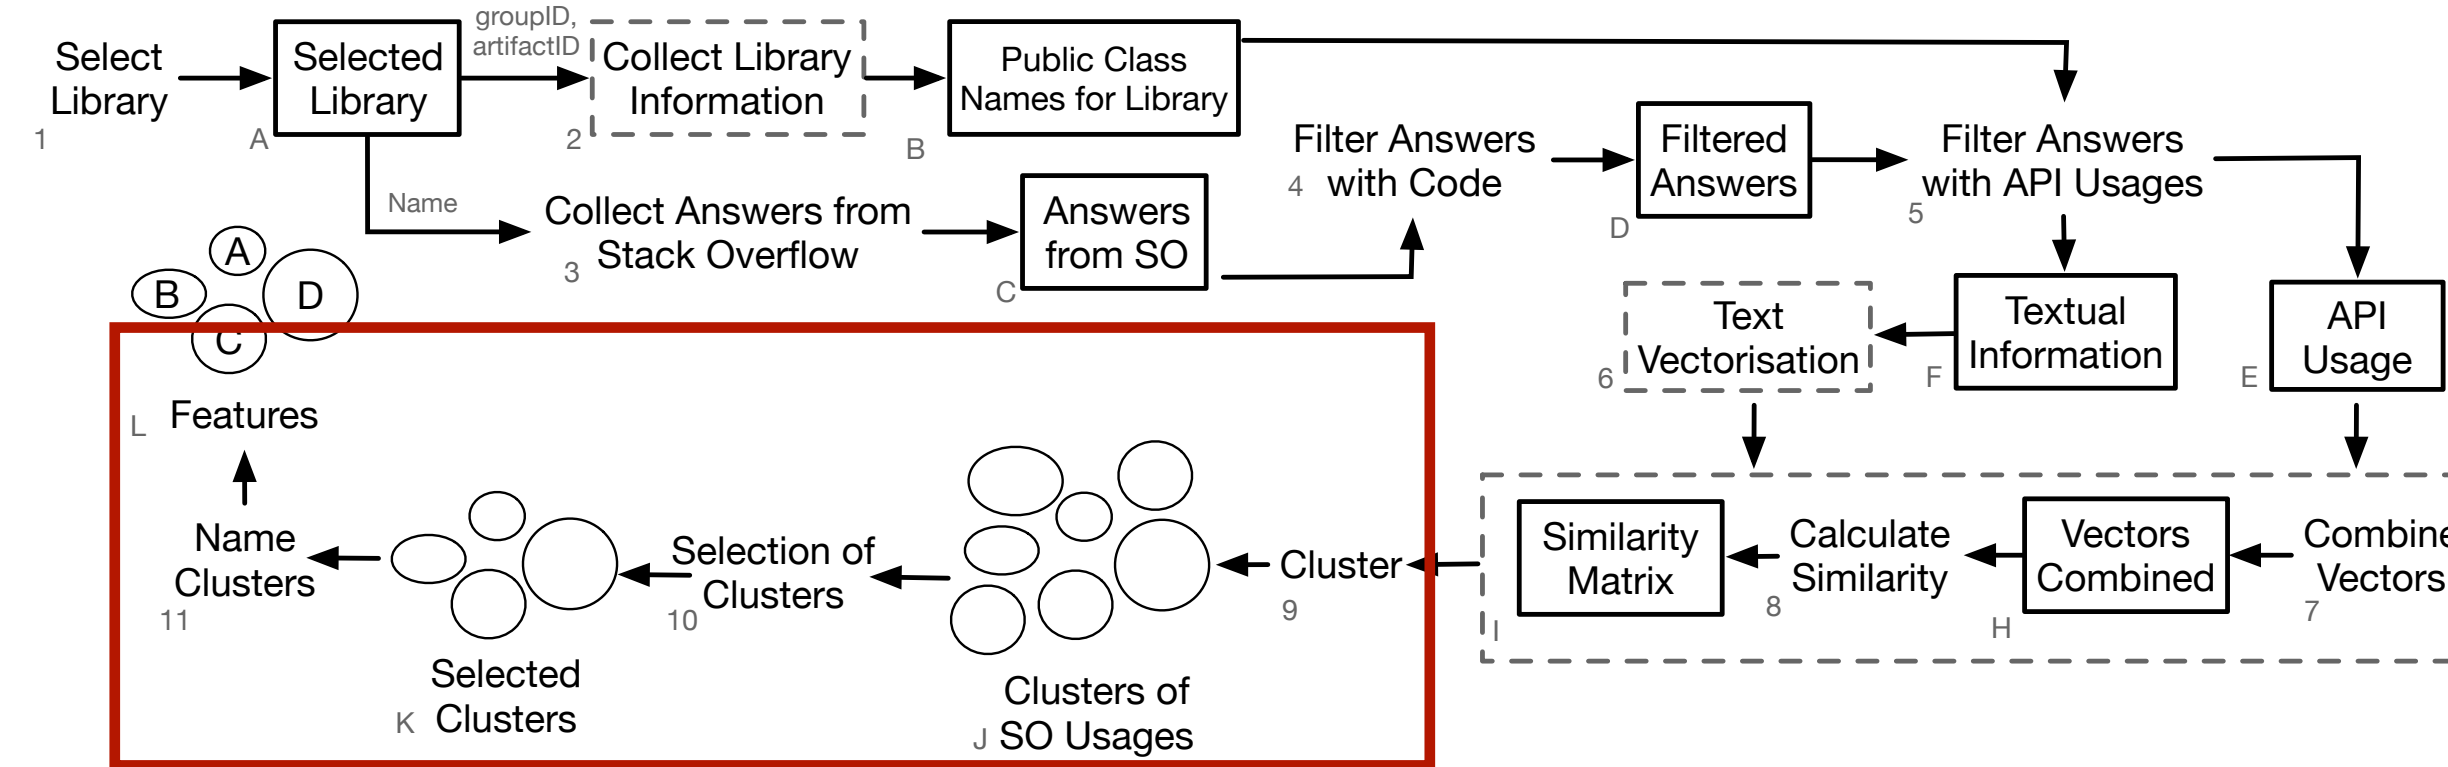
- Method Names
- Method Names split by Camel Case
- API Calls



Approach



Approach



Gene expression

Defining clusters from a hierarchical cluster tree: the Dynamic Tree Cut package for R

Peter Langfelder^{1,†}, Bin Zhang^{2,†} and Steve Horvath^{1,*}

¹Department of Human Genetics, University of California at Los Angeles, CA 90095-7088 and

²Rosetta Inpharmatics-Merck Research Laboratories, Seattle, WA, USA

Received and revised on September 12, 2007; accepted on November 6, 2007

Advance Access publication November 16, 2007

Associate Editor: Trey Ideker

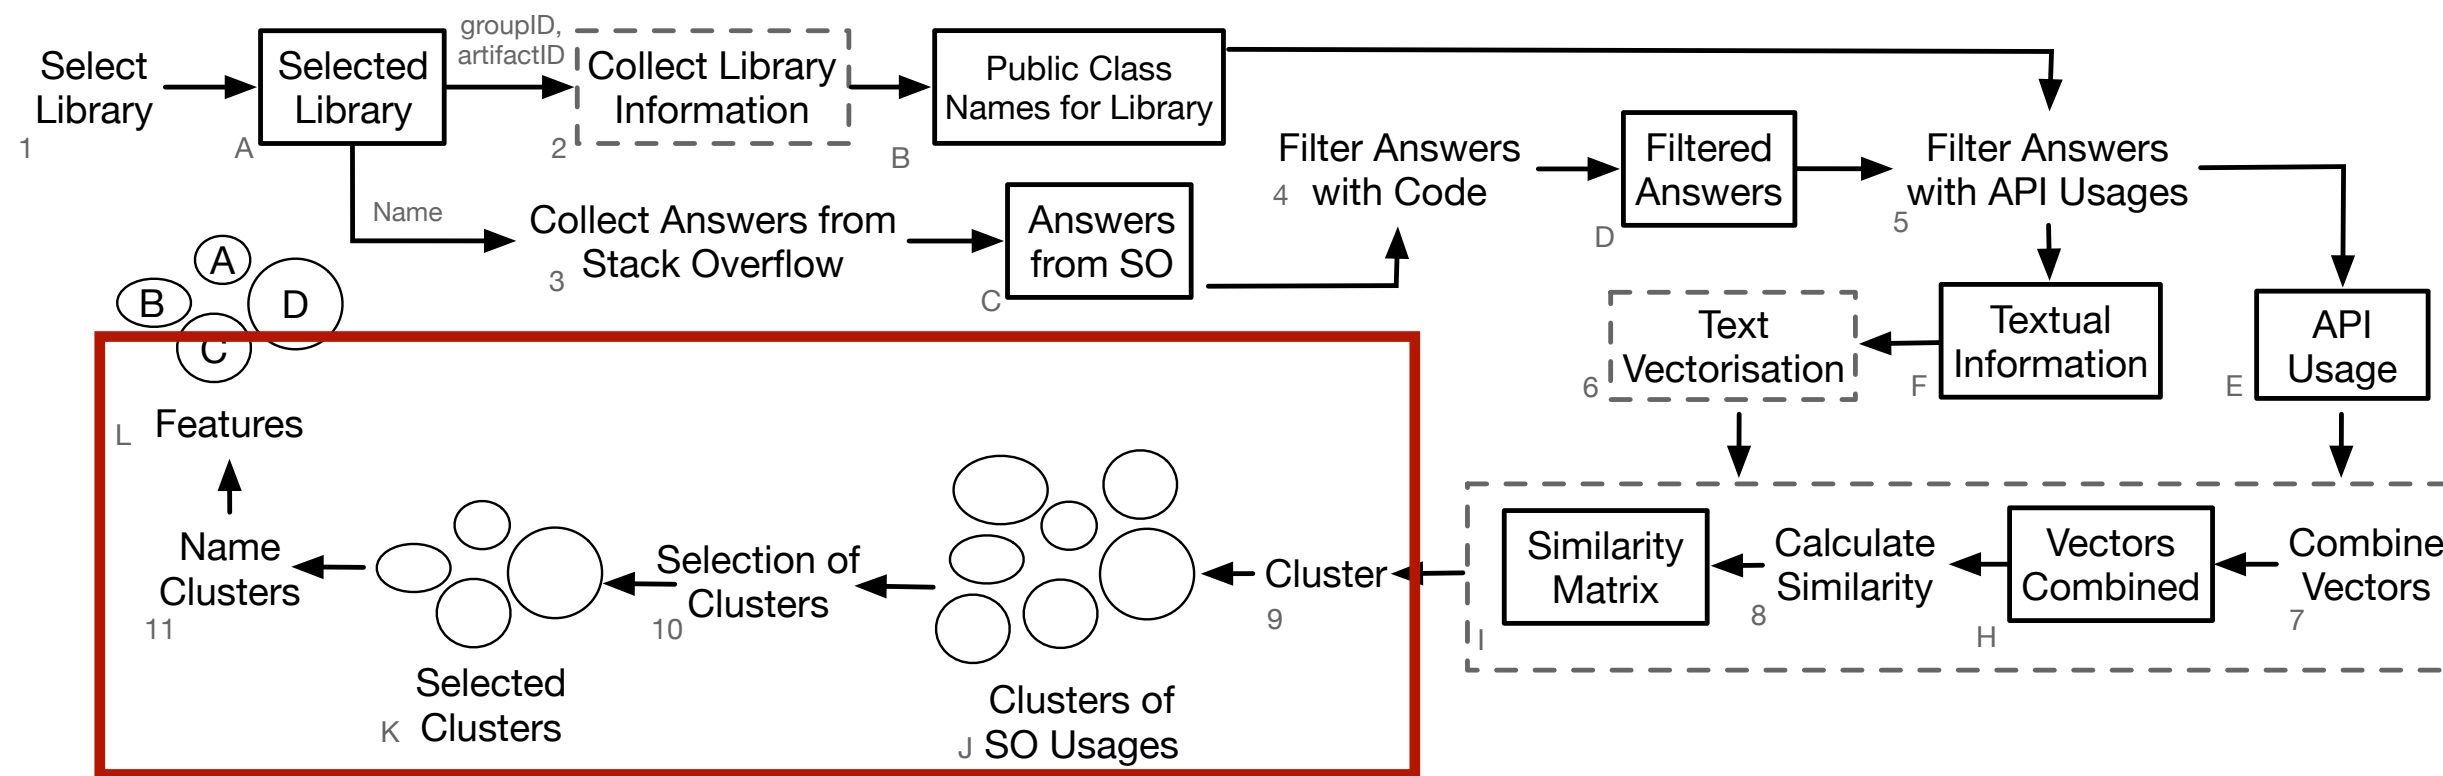
ABSTRACT

Summary: Hierarchical clustering is a widely used method for detecting clusters in genomic data. Clusters are defined by cutting branches off the dendrogram. A common but inflexible method uses a constant height cutoff value; this method exhibits suboptimal performance on complicated dendrograms. We present the Dynamic Tree Cut R package that implements novel dynamic branch cutting methods for detecting clusters in a dendrogram depending on their shape. Compared to the constant height cutoff method, our techniques offer the following advantages: (1) they are capable of identifying nested clusters; (2) they are flexible—cluster shape parameters can be tuned to suit the application at hand; (3) they are suitable for automation; and (4) they can optionally combine the advantages of hierarchical clustering and partitioning around medoids, giving better detection of outliers. We illustrate the use of these methods by applying them to protein-protein interaction network data and to a simulated gene expression data set.

cluster joining heights often poses a challenge to cluster definition. While distinct clusters may be recognizable by visual inspection, computational cluster definition by a static cut does not always identify clusters correctly. To address this challenge, we have developed a novel tree cut method based on analyzing the shape of the branches of a dendrogram. As a motivating example, consider Figure 1A that shows a dendrogram for cluster detection in a protein-protein interaction network in *Drosophila*. The Dynamic Tree Cut method succeeds at identifying branches that could not have been identified using the static cut method. The found clusters are highly significantly enriched with known gene ontologies (Dong and Horvath, 2007) which provides indirect evidence that the resulting clusters are biologically meaningful.

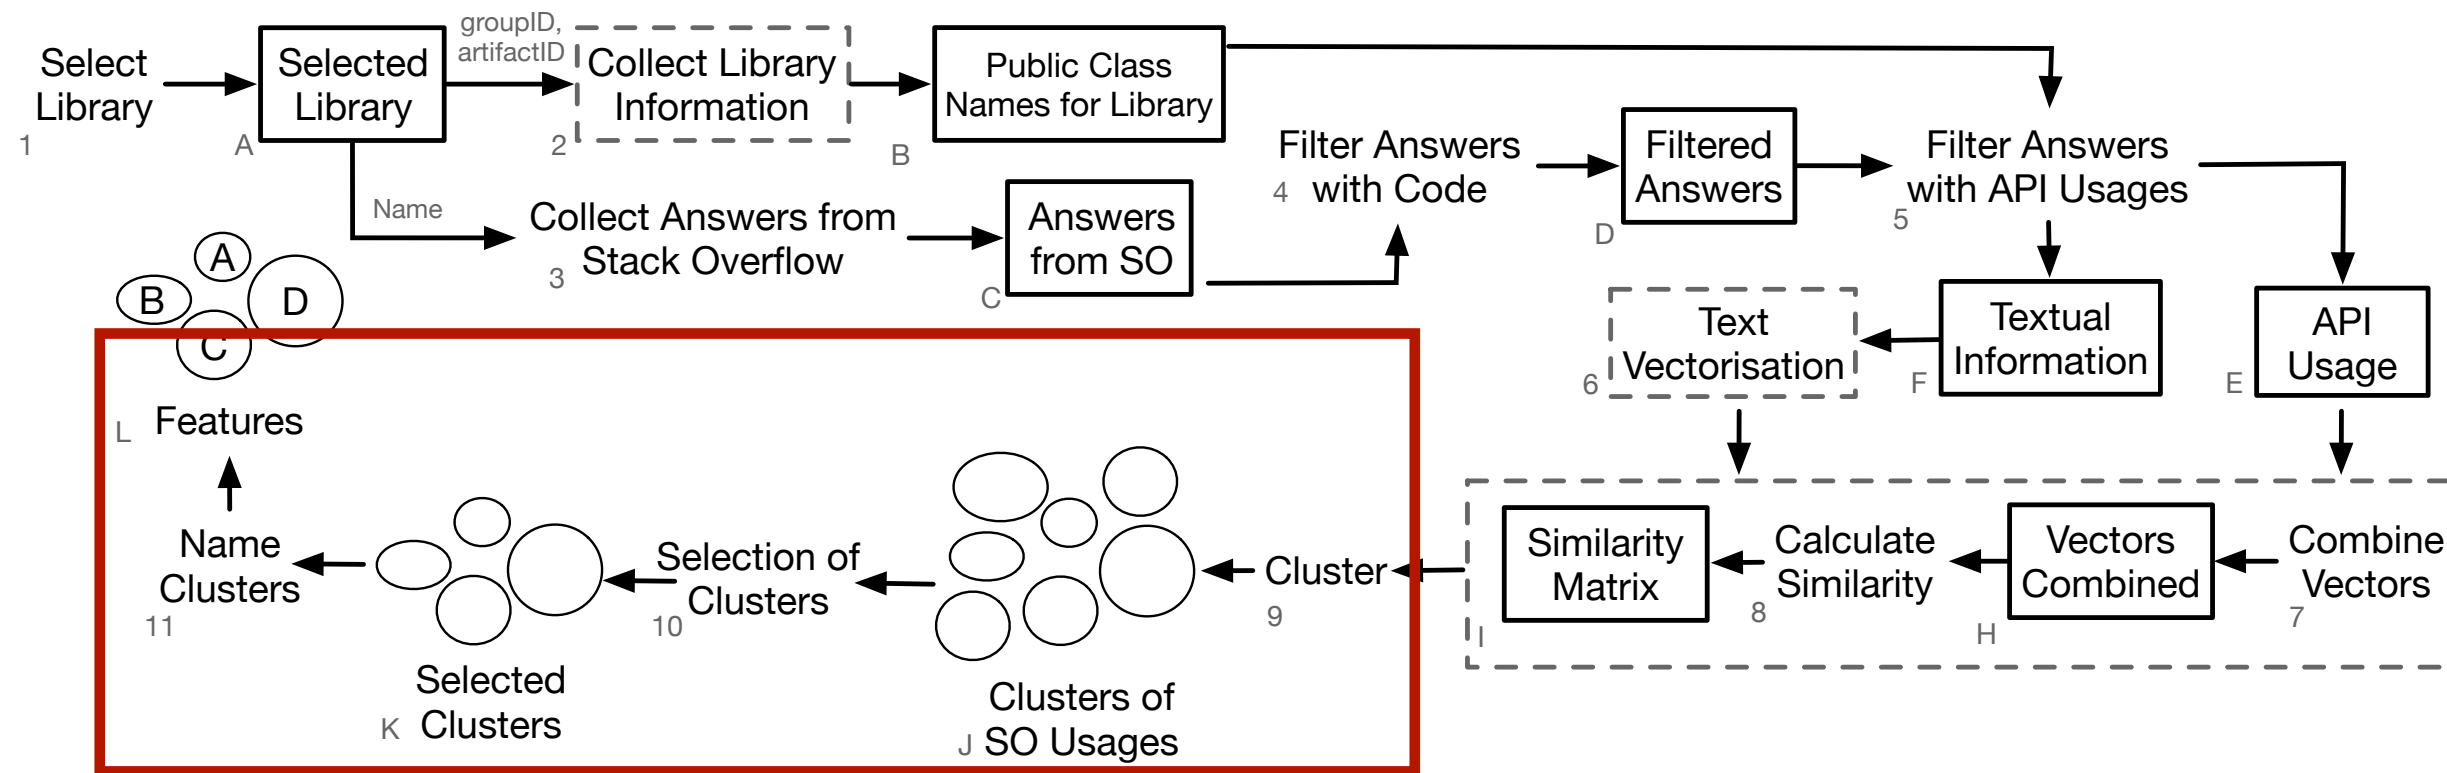
2 ALGORITHM AND IMPLEMENTATION

Approach



Formed clusters would ideally have frequent API calls, characterising the code-part of the cluster.

Approach



Formed clusters would ideally have frequent API calls, characterising the code-part of the cluster.

LOF: Identifying Density-Based Local Outliers

Markus M. Breunig[†], Hans-Peter Kriegel[†], Raymond T. Ng[‡], Jörg Sander[†]

† Institute for Computer Science
University of Munich
Oettingenstr. 67, D-80538 Munich, Germany

{ breunig | kriegel | sander }
@dbs.informatik.uni-muenchen.de

Department of Computer Science
University of British Columbia
Vancouver, BC V6T 1Z4 Canada

rng@cs.ubc.ca

ABSTRACT

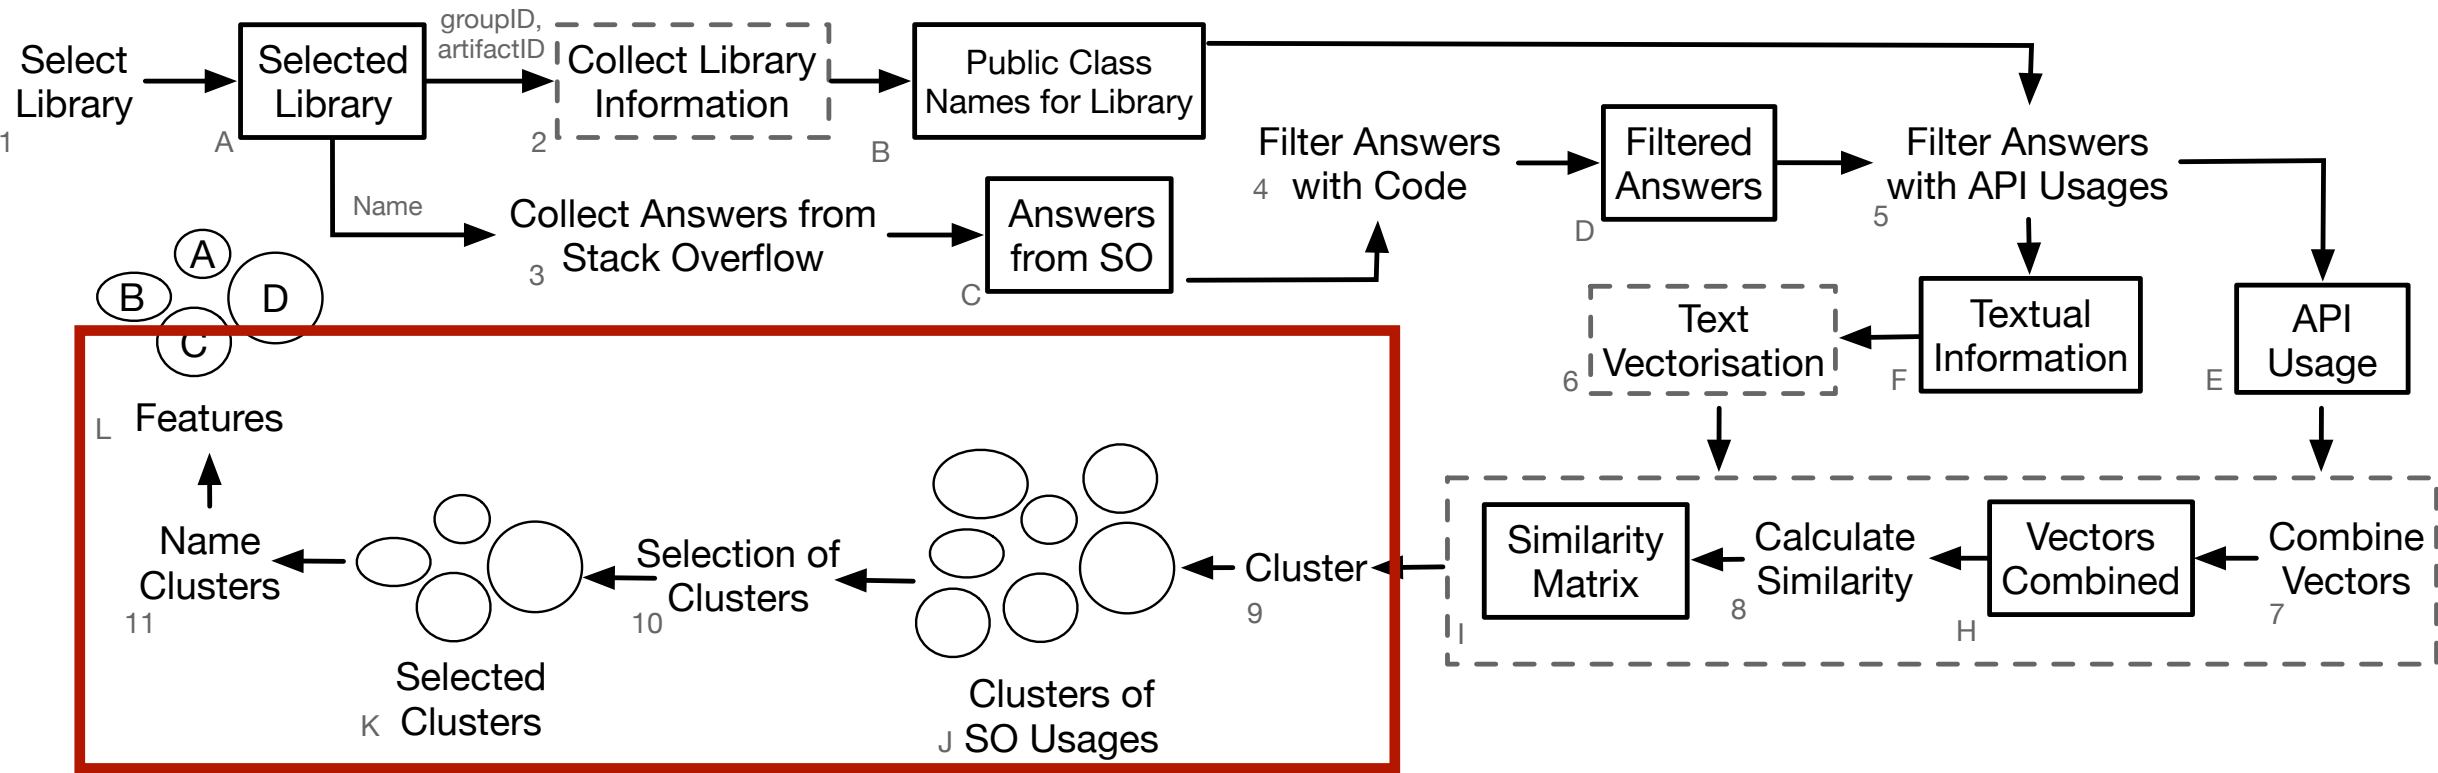
For many KDD applications, such as detecting criminal activities in E-commerce, finding the rare instances or the outliers, can be more interesting than finding the common patterns. Existing work in outlier detection regards being an outlier as a binary property. In this paper, we contend that for many scenarios, it is more meaningful to assign to each object a *degree* of being an outlier. This degree is called the *local outlier factor* (LOF) of an object. It is *local* in that the degree depends on how isolated the object is with respect to the surrounding neighborhood. We give a detailed formal analysis showing that LOF enjoys many desirable properties. Using real-world datasets, we demonstrate that LOF can be used to find outliers which appear to be meaningful, but can otherwise not be identified with existing approaches. Finally, a careful performance evaluation of our algorithm confirms we show that our approach of finding local outliers can be practical.

mon cases. Finding such exceptions and outliers, however, has not yet received as much attention in the KDD community as some other topics have, e.g. association rules.

Recently, a few studies have been conducted on outlier detection for large datasets (e.g. [18], [1], [13], [14]). While a more detailed discussion on these studies will be given in section 2, it suffices to point out here that most of these studies consider being an outlier as a binary property. That is, either an object in the dataset is an outlier or not. For many applications, the situation is more complex. And it becomes more meaningful to assign to each object a *degree* of being an outlier.

Also related to outlier detection is an extensive body of work on clustering algorithms. From the viewpoint of a clustering algorithm, outliers are objects not located in clusters of a dataset, usually called noise. The set of noise produced by a clustering algorithm, however, is highly dependent on the particular algorithm and on its

Approach



Formed clusters would ideally have frequent API calls, characterising the code-part of the cluster.

LOF: Identifying Density-Based Local Outliers

Markus M. Breunig[†], Hans-Peter Kriegel[†], Raymond T. Ng[‡], Jörg Sander[†]

[†] Institute for Computer Science
University of Munich
Oettingenstr. 67, D-80538 Munich, Germany
{ breunig | kriegel | sander }
@dbs.informatik.uni-muenchen.de

Department of Computer Science
University of British Columbia
Vancouver, BC V6T 1Z4 Canada
rng@cs.ubc.ca

ABSTRACT

For many KDD applications, such as detecting criminal activities in E-commerce, finding the rare instances or the outliers, can be more interesting than finding the common patterns. Existing work in outlier detection regards being an outlier as a binary property. In this paper, we contend that for many scenarios, it is more meaningful to assign to each object a *degree* of being an outlier. This degree is called the *local outlier factor* (LOF) of an object. It is *local* in that the degree depends on how isolated the object is with respect to the surrounding neighborhood. We give a detailed formal analysis showing that LOF enjoys many desirable properties. Using real-world datasets, we demonstrate that LOF can be used to find outliers which appear to be meaningful, but can otherwise not be identified with existing approaches. Finally, a careful performance evaluation of our algorithm confirms we show that our approach of finding local outliers can be practical.

mon cases. Finding such exceptions and outliers, however, has not yet received as much attention in the KDD community as some other topics have, e.g. association rules.

Recently, a few studies have been conducted on outlier detection for large datasets (e.g. [18], [1], [13], [14]). While a more detailed discussion on these studies will be given in section 2, it suffices to point out here that most of these studies consider being an outlier as a binary property. That is, either an object in the dataset is an outlier or not. For many applications, the situation is more complex. And it becomes more meaningful to assign to each object a *degree* of being an outlier.

Also related to outlier detection is an extensive body of work on clustering algorithms. From the viewpoint of a clustering algorithm, outliers are objects not located in clusters of a dataset, usually called noise. The set of noise produced by a clustering algorithm, however, is highly dependent on the particular algorithm and on its

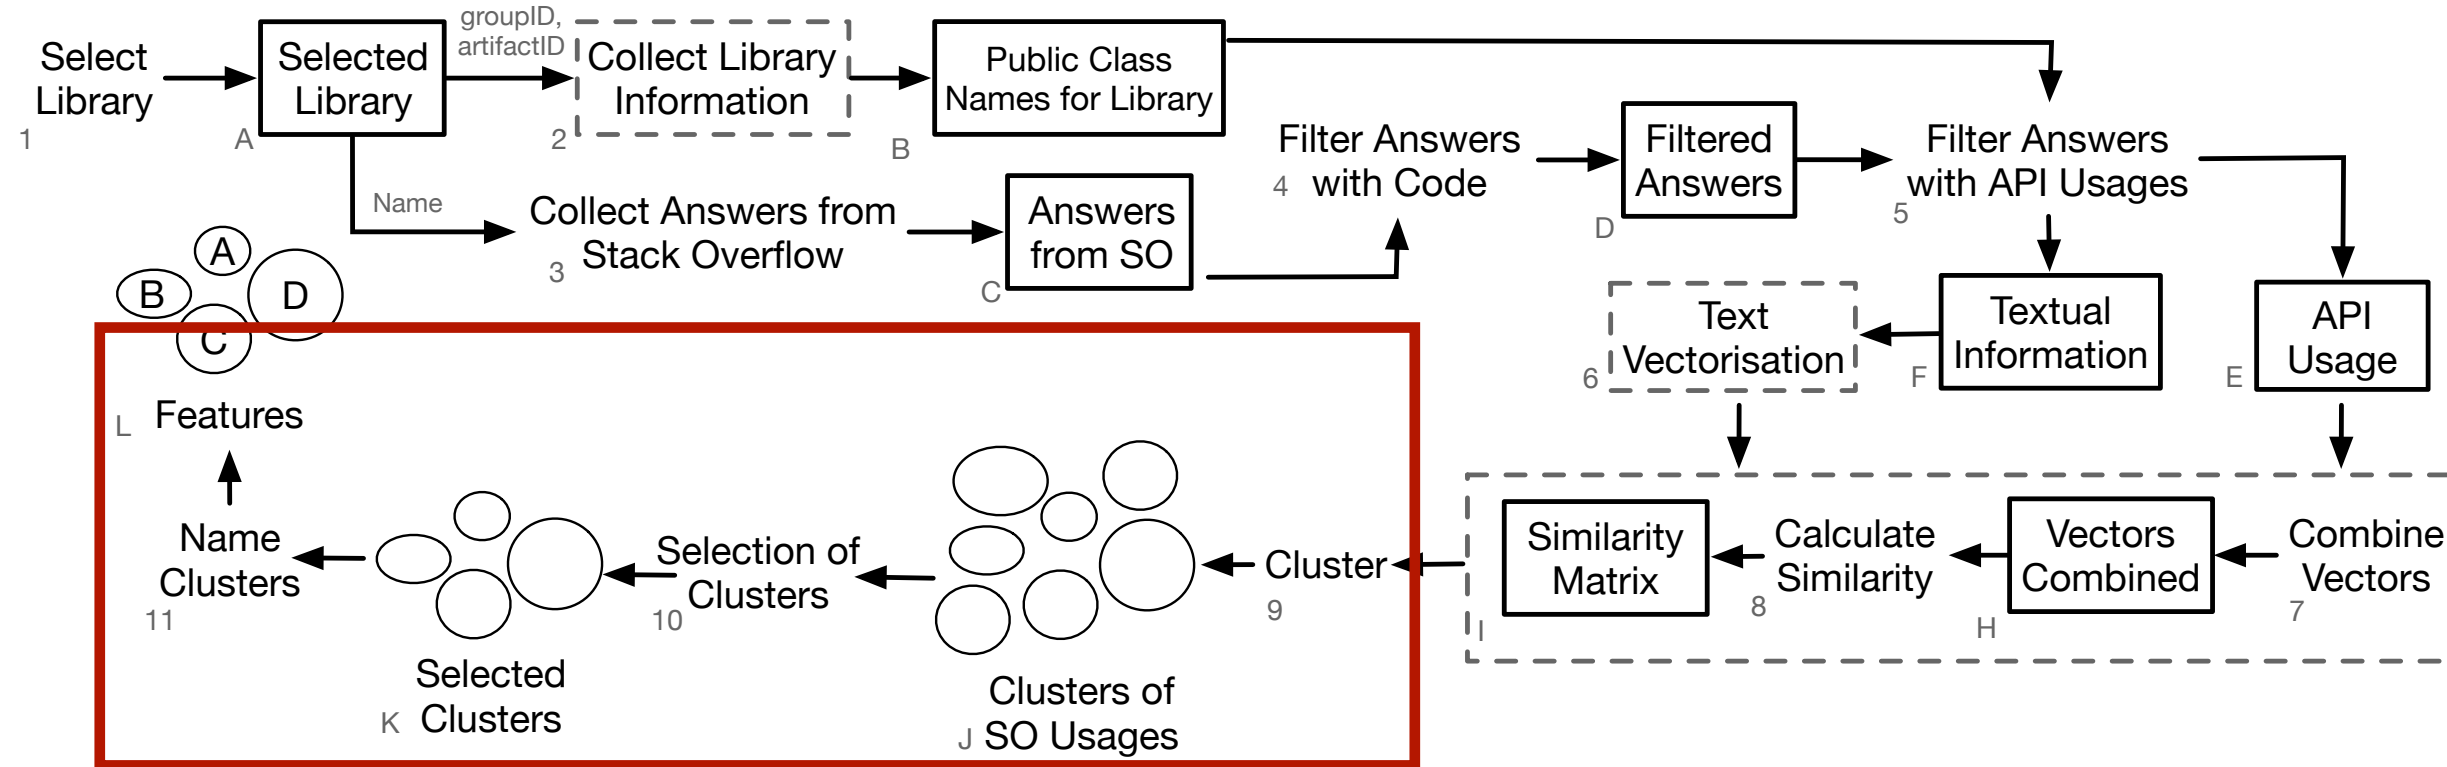
```
API calls: 5
ImmutableMultiset.of
ImmutableMultiset.entrySet
ImmutableMultiset.entrySet.asList.get.getCount
ImmutableMultiset.entrySet.iterator.next.getCount
Multisets.copyHighestCountFirst.copyHighestCountFirst
```

unique.methods	freq.methods	percent.methods
entrySet	6	100.00
Entry	4	100.00
getCount	4	100.00
iterator	3	100.00
copyHighestCountFirst	3	100.00
getElement	2	67.00

[4] - Breunig, Markus M., Hans-Peter Kriegel, Raymond T. Ng, and Jörg Sander. "LOF: identifying density-based local outliers." In *Proceedings of the 2000 ACM SIGMOD international conference on Management of data*, pp. 93-104. 2000.

[5] - C. D. Manning, M. Surdeanu, J. Bauer, J. R. Finkel, S. Bethard, and D. McClosky, "The Stanford CoreNLP natural language processing toolkit," in *The 52nd Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pp. 55-60. 2014.

Approach



Formed clusters would ideally have frequent API calls, characterising the code-part of the cluster.

LOF: Identifying Density-Based Local Outliers

Markus M. Breunig[†], Hans-Peter Kriegel[†], Raymond T. Ng[‡], Jörg Sander[†]

† Institute for Computer Science
University of Munich
Oettingenstr. 67, D-80538 Munich, Germany

{ breunig | kriegel | sander }
@dbs.informatik.uni-muenchen.de

Department of Computer Science
University of British Columbia
Vancouver, BC V6T 1Z4 Canada

rng@cs.ubc.ca

ABSTRACT

For many KDD applications, such as detecting criminal activities in E-commerce, finding the rare instances or the outliers, can be more interesting than finding the common patterns. Existing work in outlier detection regards being an outlier as a binary property. In this paper, we contend that for many scenarios, it is more meaningful to assign to each object a *degree* of being an outlier. This degree is called the *local outlier factor* (LOF) of an object. It is *local* in that the degree depends on how isolated the object is with respect to the surrounding neighborhood. We give a detailed formal analysis showing that LOF enjoys many desirable properties. Using real-world datasets, we demonstrate that LOF can be used to find outliers which appear to be meaningful, but can otherwise not be identified with existing approaches. Finally, a careful performance evaluation of our algorithm confirms we show that our approach of finding local outliers can be practical.

mon cases. Finding such exceptions and outliers, however, has not yet received as much attention in the KDD community as some other topics have, e.g. association rules.

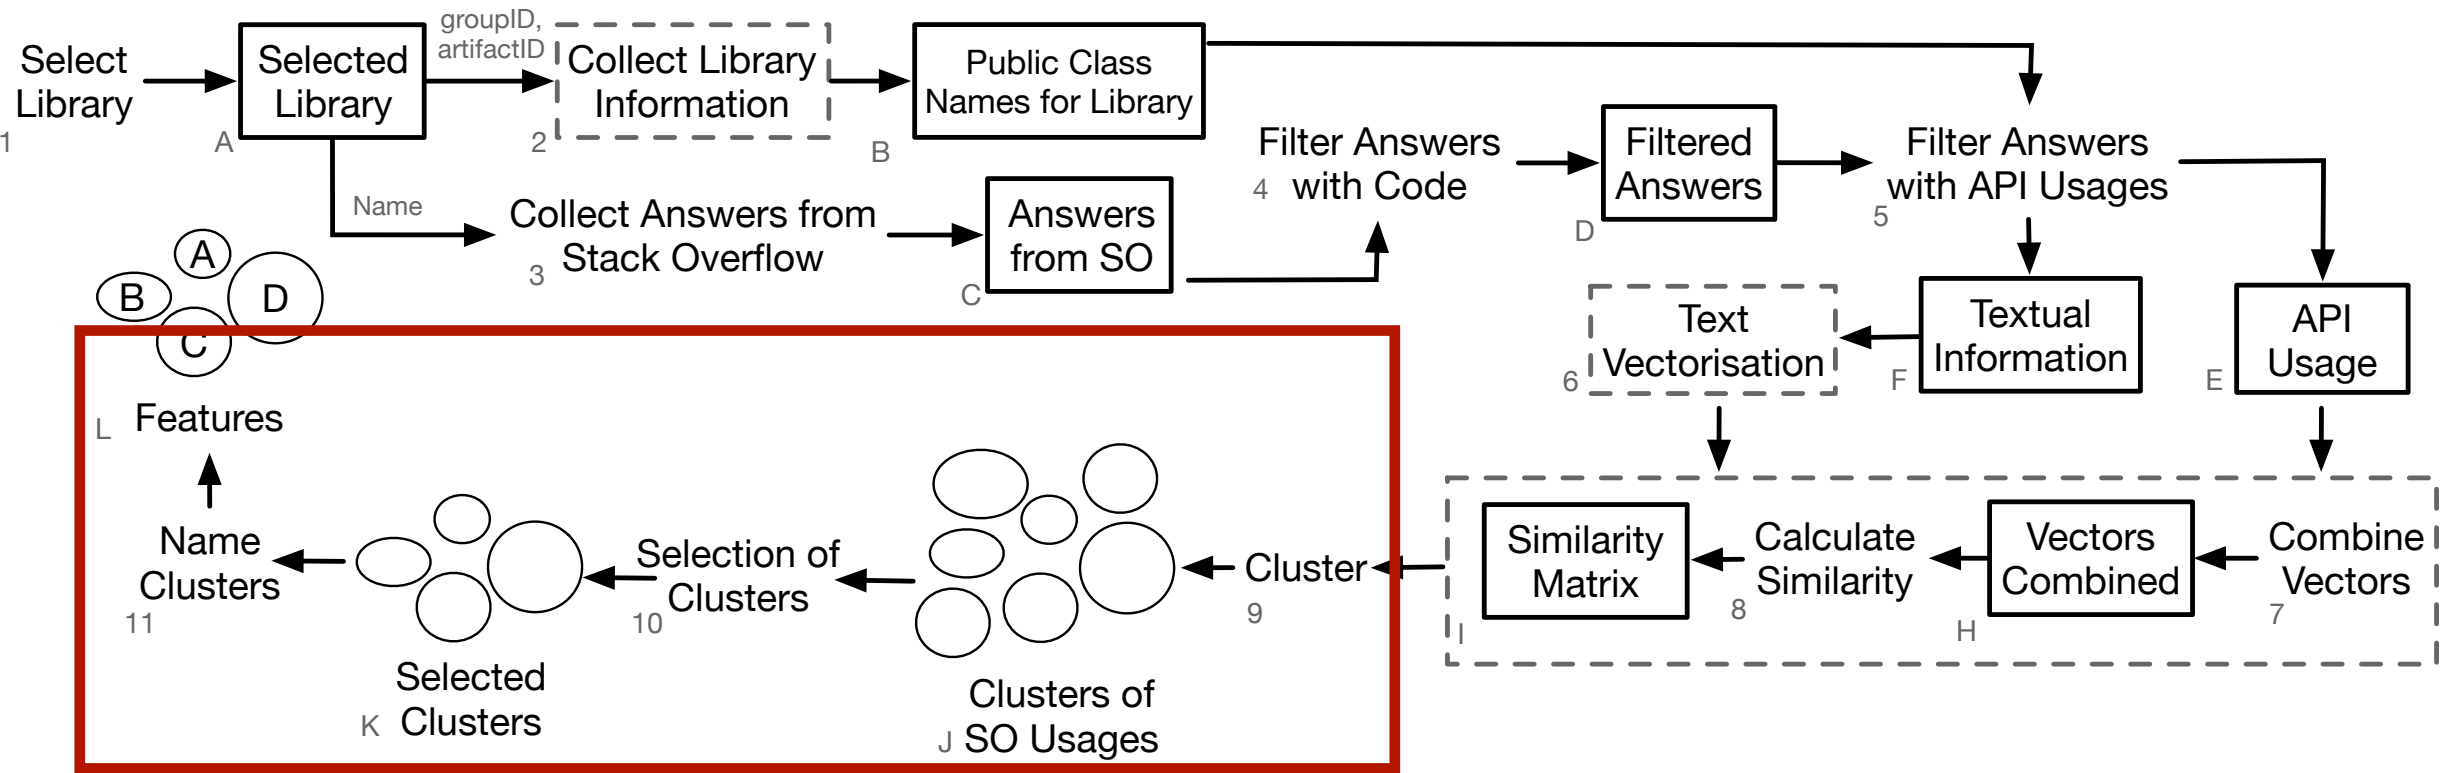
Recently, a few studies have been conducted on outlier detection for large datasets (e.g. [18], [1], [13], [14]). While a more detailed discussion on these studies will be given in section 2, it suffices to point out here that most of these studies consider being an outlier as a binary property. That is, either an object in the dataset is an outlier or not. For many applications, the situation is more complex. And it becomes more meaningful to assign to each object a *degree* of being an outlier.

Also related to outlier detection is an extensive body of work on clustering algorithms. From the viewpoint of a clustering algorithm, outliers are objects not located in clusters of a dataset, usually called noise. The set of noise produced by a clustering algorithm, however, is highly dependent on the particular algorithm and on its

```
API calls: 5
ImmutableMultiset.of
ImmutableMultiset.entrySet
ImmutableMultiset.entrySet.asList.get.getCount
ImmutableMultiset.entrySet.iterator.next.getCount
Multisets.copyHighestCountFirst.copyHighestCountFirst
```

unique.methods	req.methods	percent.methods
entrySet	6	100.00
Entry	4	100.00
getCount	4	100.00
iterator	3	100.00
copyHighestCountFirst	3	100.00
getElement	2	67.00

Approach



Formed clusters would ideally have frequent API calls, characterising the code-part of the cluster.

LOF: Identifying Density-Based Local Outliers

Markus M. Breunig[†], Hans-Peter Kriegel[†], Raymond T. Ng[‡], Jörg Sander[†]

[†] Institute for Computer Science
University of Munich
Oettingenstr. 67, D-80538 Munich, Germany

{ breunig | kriegel | sander }
@dbs.informatik.uni-muenchen.de

Department of Computer Science
University of British Columbia
Vancouver, BC V6T 1Z4 Canada

rng@cs.ubc.ca

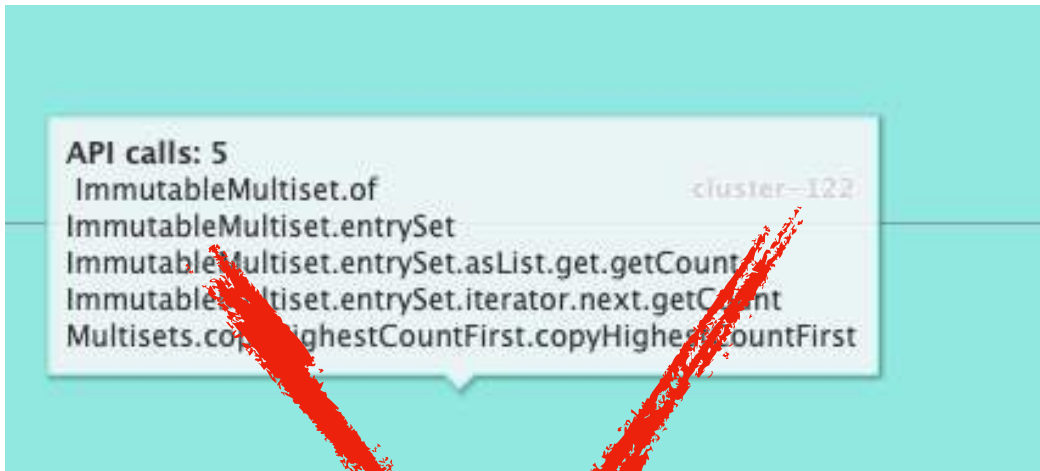
ABSTRACT

For many KDD applications, such as detecting criminal activities in E-commerce, finding the rare instances or the outliers, can be more interesting than finding the common patterns. Existing work in outlier detection regards being an outlier as a binary property. In this paper, we contend that for many scenarios, it is more meaningful to assign to each object a *degree* of being an outlier. This degree is called the *local outlier factor* (LOF) of an object. It is *local* in that the degree depends on how isolated the object is with respect to the surrounding neighborhood. We give a detailed formal analysis showing that LOF enjoys many desirable properties. Using real-world datasets, we demonstrate that LOF can be used to find outliers which appear to be meaningful, but can otherwise not be identified with existing approaches. Finally, a careful performance evaluation of our algorithm confirms we show that our approach of finding local outliers can be practical.

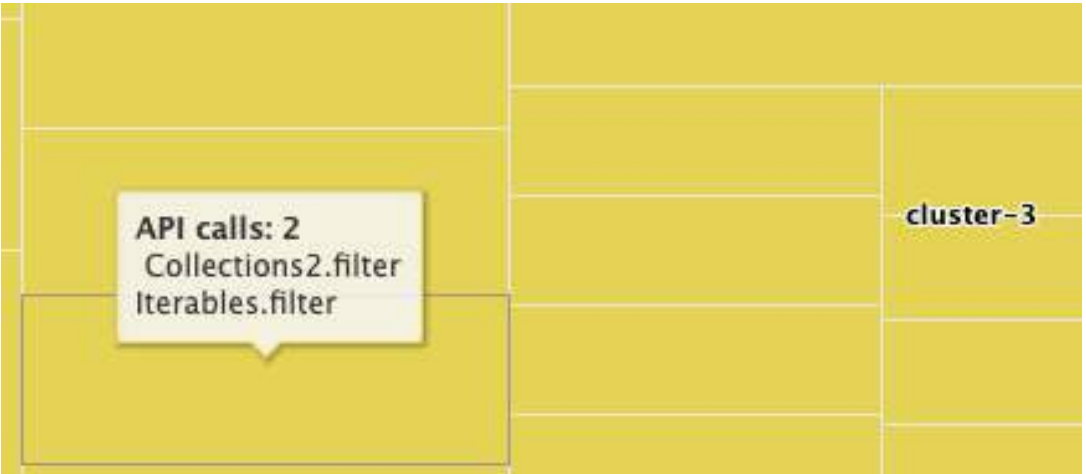
mon cases. Finding such exceptions and outliers, however, has not yet received as much attention in the KDD community as some other topics have, e.g. association rules.

Recently, a few studies have been conducted on outlier detection for large datasets (e.g. [18], [1], [13], [14]). While a more detailed discussion on these studies will be given in section 2, it suffices to point out here that most of these studies consider being an outlier as a binary property. That is, either an object in the dataset is an outlier or not. For many applications, the situation is more complex. And it becomes more meaningful to assign to each object a *degree* of being an outlier.

Also related to outlier detection is an extensive body of work on clustering algorithms. From the viewpoint of a clustering algorithm, outliers are objects not located in clusters of a dataset, usually called noise. The set of noise produced by a clustering algorithm, however, is highly dependent on the particular algorithm and on its



unique.methods	freq.methods	percent.methods
entrySet	6	100.00
Entry	4	100.00
getCount	4	100.00
iterator	3	100.00
copyHighestCountFirst	3	100.00
getElement	2	67.00

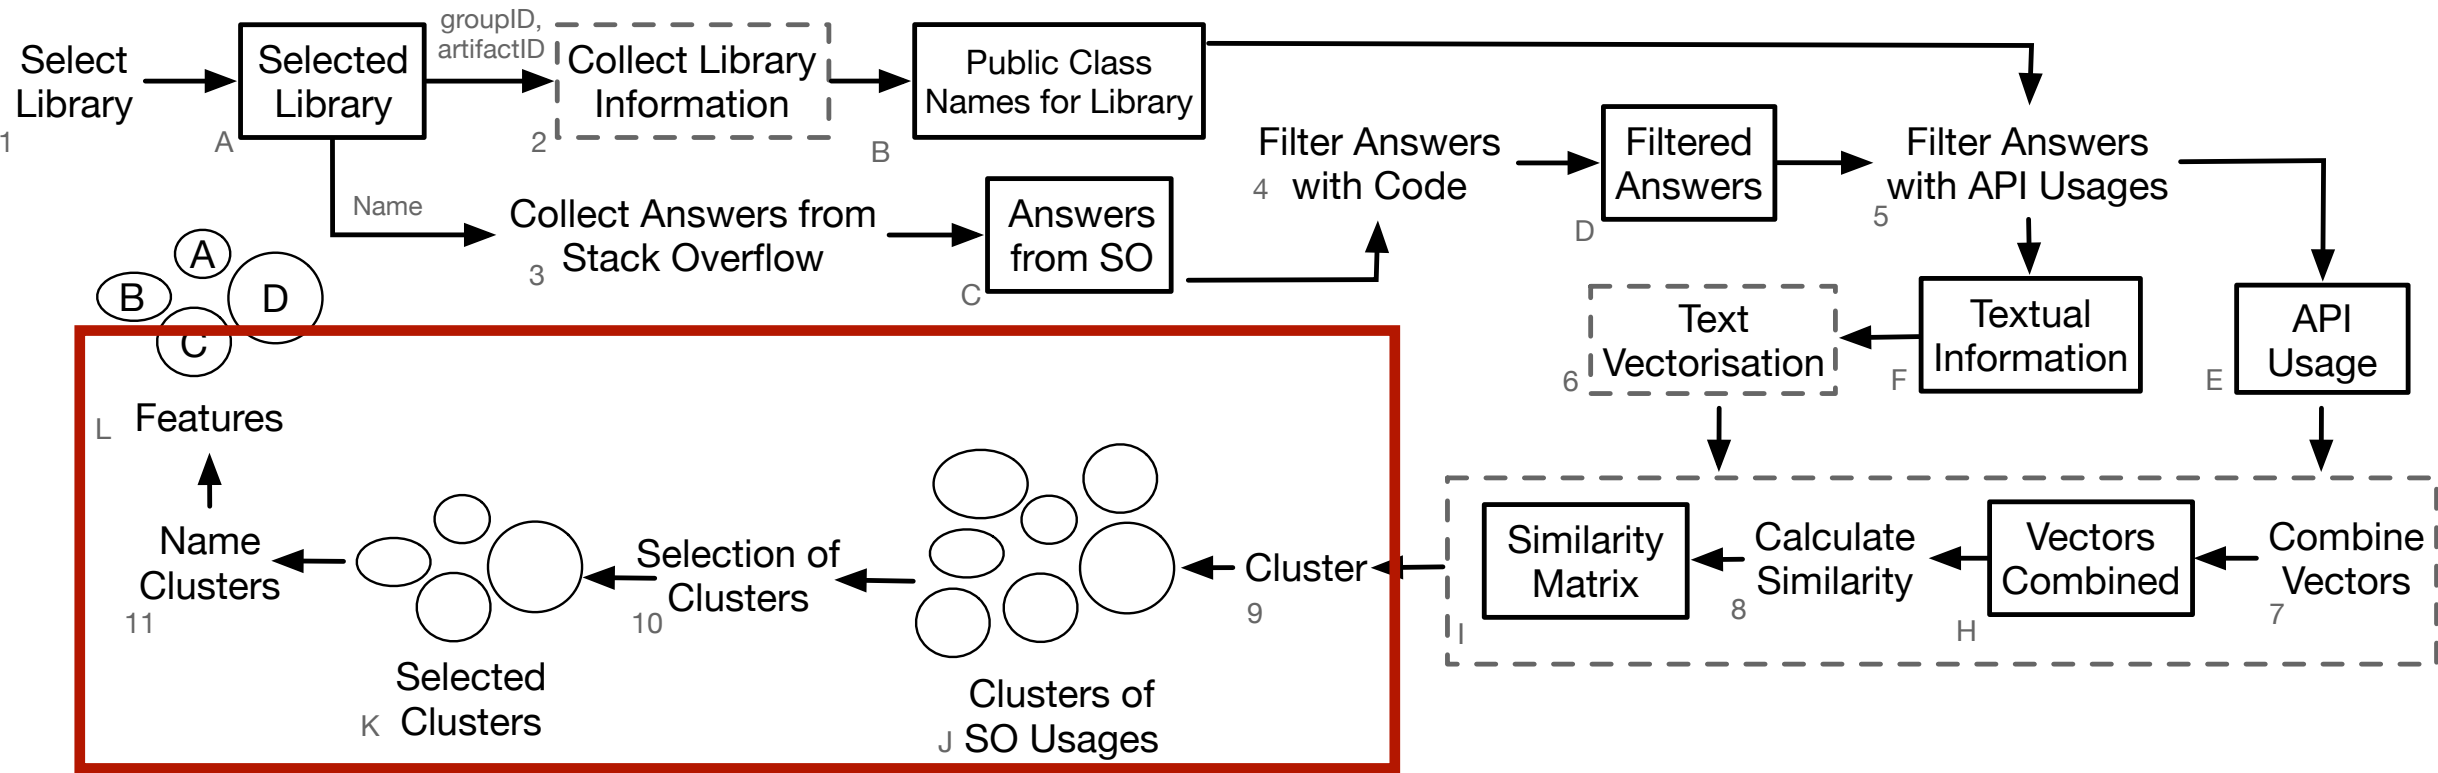


unique.methods	freq.methods	percent.methods
filter	45	100.00
and	2	5.00
cells	2	5.00
clear	2	5.00
ofNullable	2	5.00

[4] - Breunig, Markus M., Hans-Peter Kriegel, Raymond T. Ng, and Jörg Sander. "LOF: identifying density-based local outliers." In *Proceedings of the 2000 ACM SIGMOD international conference on Management of data*, pp. 93-104. 2000.

[5] - C. D. Manning, M. Surdeanu, J. Bauer, J. R. Finkel, S. Bethard, and D. McClosky, "The Stanford CoreNLP natural language processing toolkit," in *The 52nd Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pp. 55-60. 2014.

Approach



Formed clusters would ideally have frequent API calls, characterising the code-part of the cluster.

LOF: Identifying Density-Based Local Outliers

Markus M. Breunig[†], Hans-Peter Kriegel[†], Raymond T. Ng[‡], Jörg Sander[†]

[†] Institute for Computer Science
University of Munich
Oettingenstr. 67, D-80538 Munich, Germany

{ breunig | kriegel | sander }
@dbs.informatik.uni-muenchen.de

Department of Computer Science
University of British Columbia
Vancouver, BC V6T 1Z4 Canada

rng@cs.ubc.ca

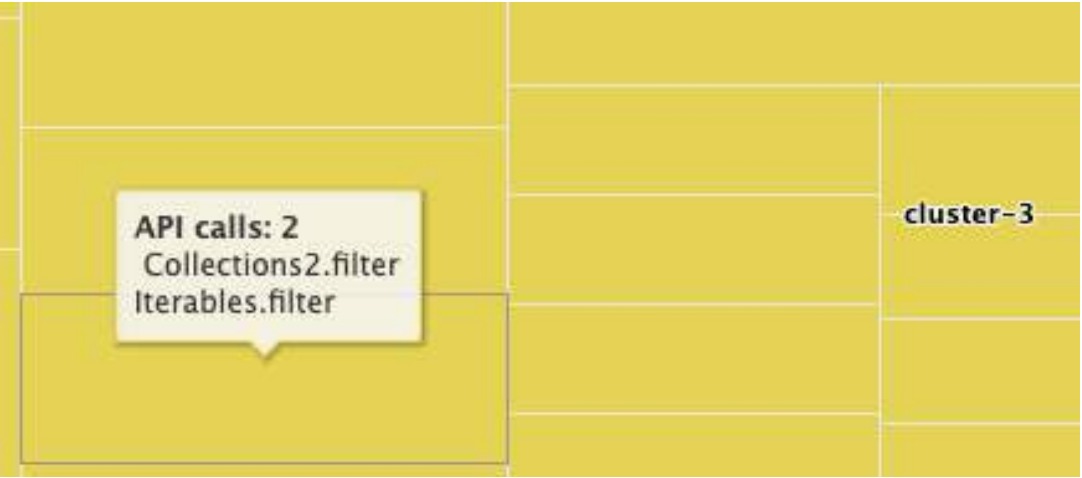
ABSTRACT

For many KDD applications, such as detecting criminal activities in E-commerce, finding the rare instances or the outliers, can be more interesting than finding the common patterns. Existing work in outlier detection regards being an outlier as a binary property. In this paper, we contend that for many scenarios, it is more meaningful to assign to each object a *degree* of being an outlier. This degree is called the *local outlier factor* (LOF) of an object. It is *local* in that the degree depends on how isolated the object is with respect to the surrounding neighborhood. We give a detailed formal analysis showing that LOF enjoys many desirable properties. Using real-world datasets, we demonstrate that LOF can be used to find outliers which appear to be meaningful, but can otherwise not be identified with existing approaches. Finally, a careful performance evaluation of our algorithm confirms we show that our approach of finding local outliers can be practical.

mon cases. Finding such exceptions and outliers, however, has not yet received as much attention in the KDD community as some other topics have, e.g. association rules.

Recently, a few studies have been conducted on outlier detection for large datasets (e.g. [18], [1], [13], [14]). While a more detailed discussion on these studies will be given in section 2, it suffices to point out here that most of these studies consider being an outlier as a binary property. That is, either an object in the dataset is an outlier or not. For many applications, the situation is more complex. And it becomes more meaningful to assign to each object a *degree* of being an outlier.

Also related to outlier detection is an extensive body of work on clustering algorithms. From the viewpoint of a clustering algorithm, outliers are objects not located in clusters of a dataset, usually called noise. The set of noise produced by a clustering algorithm, however, is highly dependent on the particular algorithm and on its



- filter list
- use collection
- use predicate

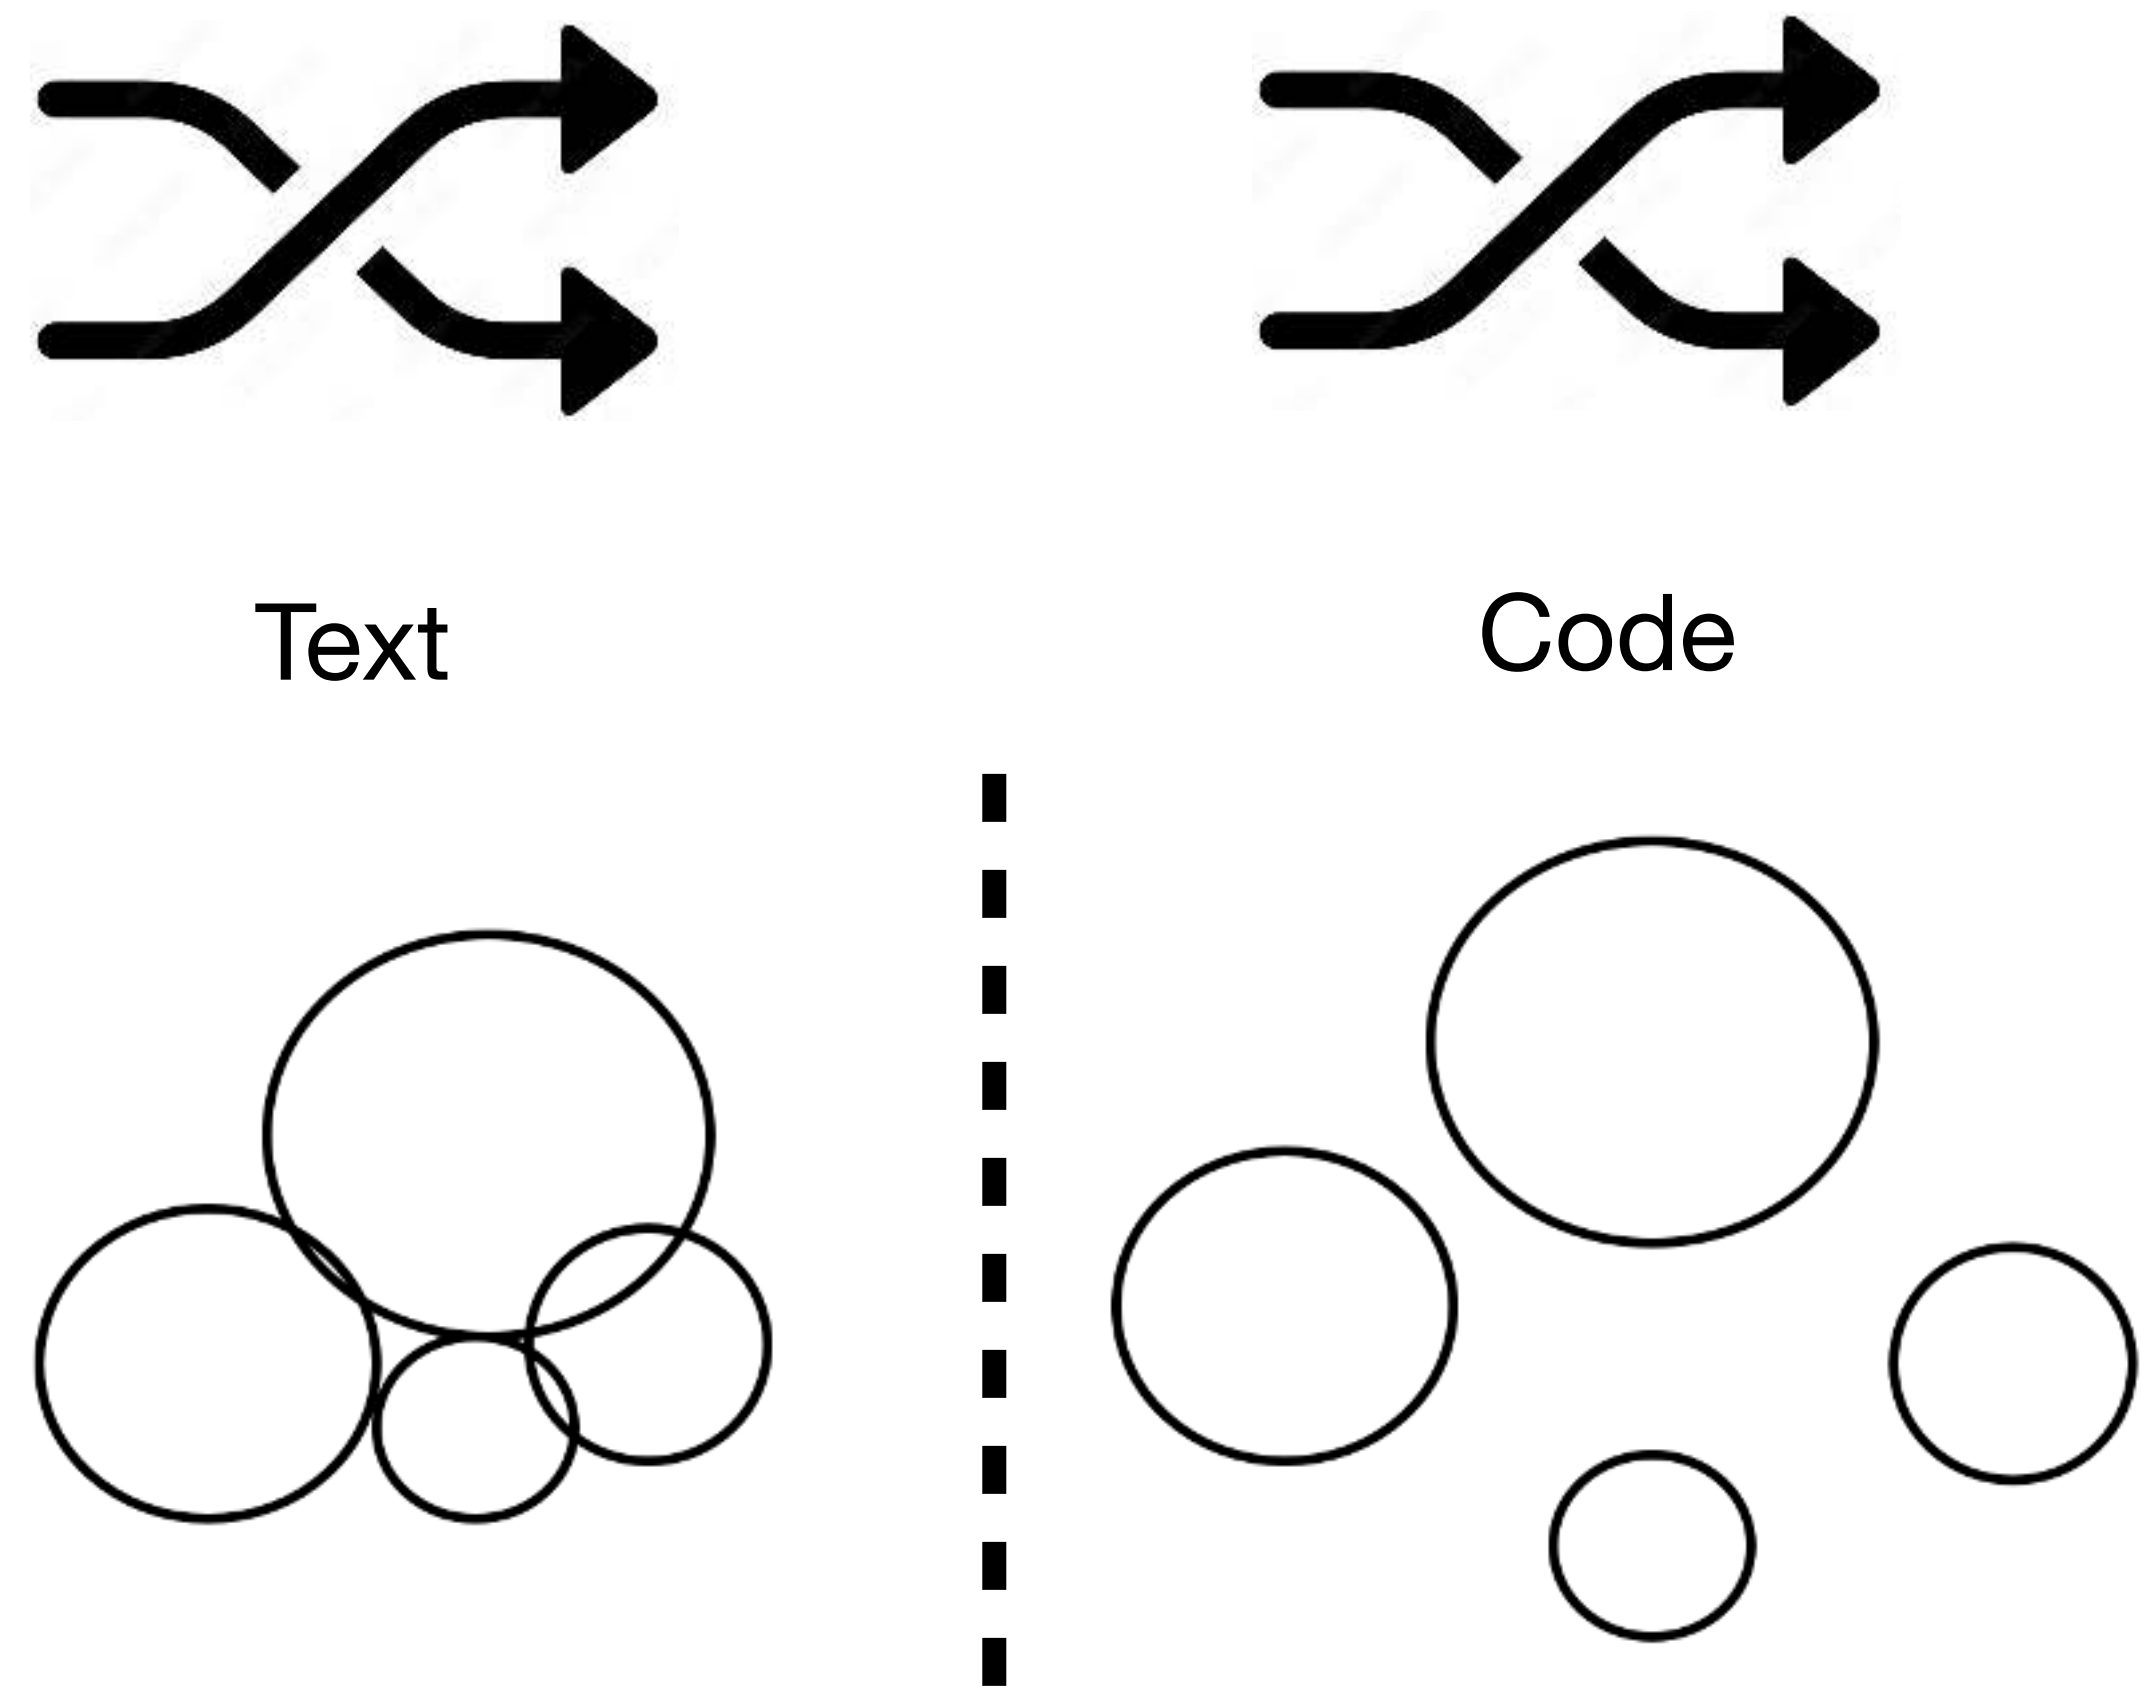
unique.methods	freq.methods	percent.methods
filter	45	100.00
and	2	5.00
cells	2	5.00
clear	2	5.00
ofNullable	2	5.00

[4] - Breunig, Markus M., Hans-Peter Kriegel, Raymond T. Ng, and Jörg Sander. "LOF: identifying density-based local outliers." In *Proceedings of the 2000 ACM SIGMOD international conference on Management of data*, pp. 93-104. 2000.

[5] - C. D. Manning, M. Surdeanu, J. Bauer, J. R. Finkel, S. Bethard, and D. McClosky, "The Stanford CoreNLP natural language processing toolkit," in *The 52nd Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pp. 55-60. 2014.

Evaluation

RQ1. Which combination of SO answer attributes produces the most cohesive clusters?



SC	Proposed Interpretation
0.71–1.00	A strong structure has been found
0.51–0.70	A reasonable structure has been found
0.26–0.50	The structure is weak and could be artificial; please try additional methods on this data set
≤ 0.25	No substantial structure has been found

Evaluation

RQ1. Which combination of SO answer attributes produces the most cohesive clusters?

TABLE II

TOP- K SCORES FOR ALL COMBINATIONS OF LIBRARY ATTRIBUTES.
BETWEEN PARENTHESES IS THE OBTAINED SCORE.

Library	K=1	K=2	K=3
Guava	Methods (0.6)	Methods CC (0.6)	Methods + Methods CC (0.59)
HttpClient	API Calls (0.46)	Methods (0.45)	Methods CC (0.45)
JFreeChart	Methods (0.38)	Methods CC (0.38)	Methods + Methods CC (0.38)
JSoup	API Calls (0.63)	Methods CC (0.51)	Methods (0.51)
PDFBox	Methods + Methods CC (0.37)	Methods CC (0.36)	Methods (0.36)
Apache-POI	Methods (0.52)	Methods CC (0.52)	Methods + Methods CC (0.51)
Quartz	Methods + Methods CC (0.45)	Methods CC (0.42)	API Calls (0.42)

The **Methods** combination from the API usage information attributes produced the most cohesive clusters.

Evaluation

RQ2. How similar are the automatically uncovered features to documented tutorial features?

Guava Collections Cookbook

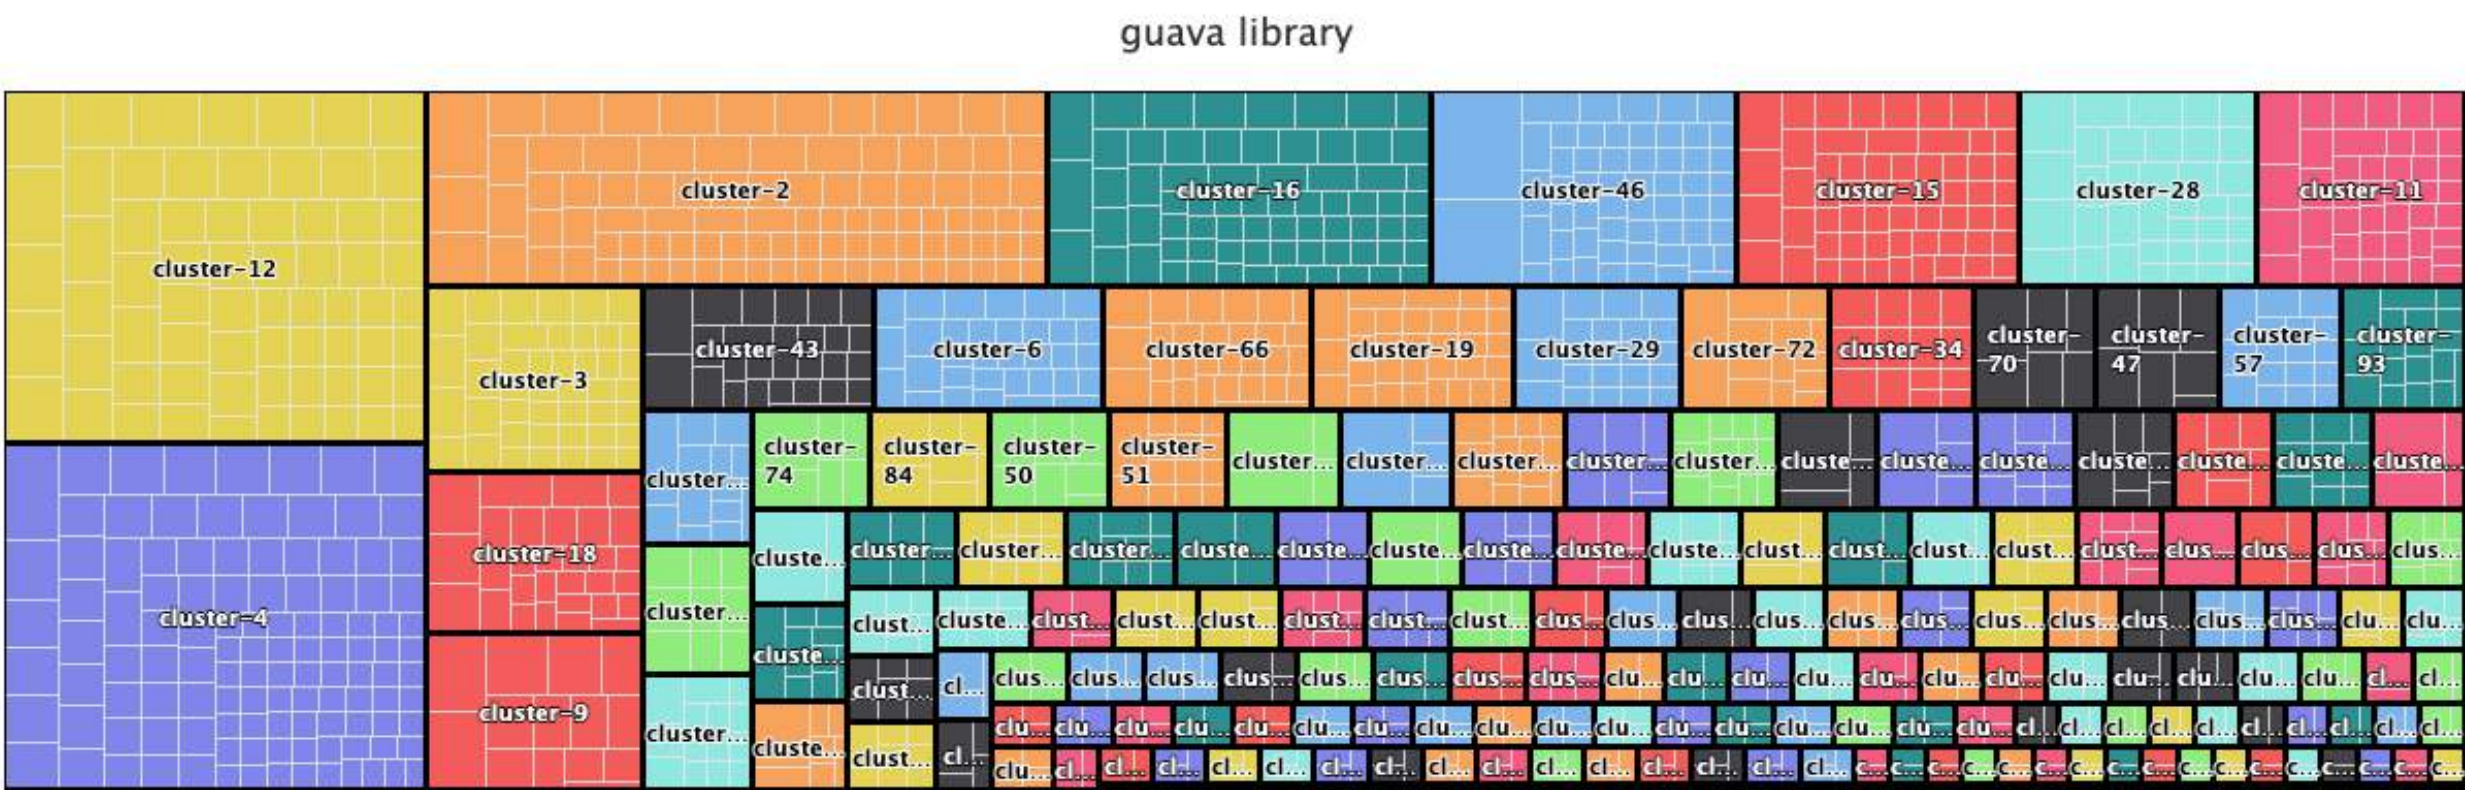
adding an iterable to a collection

```
Iterable<String> iter = Lists.newArrayList();
Collection<String> collector = Lists.newArrayList();
Iterables.addAll(collector, iter);
```

check if collection contains element(s) according to a custom matching rule

```
Iterable<String> theCollection = Lists.newArrayList("a", "bc", "def");
boolean contains = Iterables.any(theCollection, new Predicate<String>() {
    @Override
    public boolean apply(final String input) {
        return input.length() == 1;
    }
});
assertTrue(contains);
```

Generated features



Evaluation

RQ2. How similar are the automatically uncovered features to documented tutorial features?

TABLE III
ANALYSIS OF THE MATCHED FEATURES PER LIBRARY.

Library	No. Feat.	Match.	Not Found	Acc.	R-Acc.	Relv.	Over.
Guava	30	22	7	0.73	0.97	0.71	-0.28
HttpClient	12	9	3	0.75	1.0	0.71	-0.23
JFreeChart	9	8	0	0.89	0.89	0.73	-0.26
JSoup	15	11	0	0.73	0.73	0.95	-0.51
PDFBox	15	8	4	0.53	0.73	0.82	-0.30
Apache-POI	20	10	3	0.50	0.59	0.75	-0.39
Quartz	22	12	5	0.55	0.71	0.65	-0.58

The majority of the tutorial features is part of our uncovered features.

Evaluation

RQ2. How similar are the automatically uncovered features to documented tutorial features?

TABLE III
ANALYSIS OF THE MATCHED FEATURES PER LIBRARY.

Library	No. Feat.	Match.	Not Found	Acc.	R-Acc.	Relv.	Over.
Guava	30	22	7	0.73	0.97	0.71	-0.28
HttpClient	12	9	3	0.75	1.0	0.71	-0.23
JFreeChart	9	8	0	0.89	0.89	0.73	-0.26
JSoup	15	11	0	0.73	0.73	0.95	-0.51
PDFBox	15	8	4	0.53	0.73	0.82	-0.30
Apache-POI	20	10	3	0.50	0.59	0.75	-0.39
Quartz	22	12	5	0.55	0.71	0.65	-0.58

When considered against the usages in the SO data, the matched accuracy increases.

Evaluation

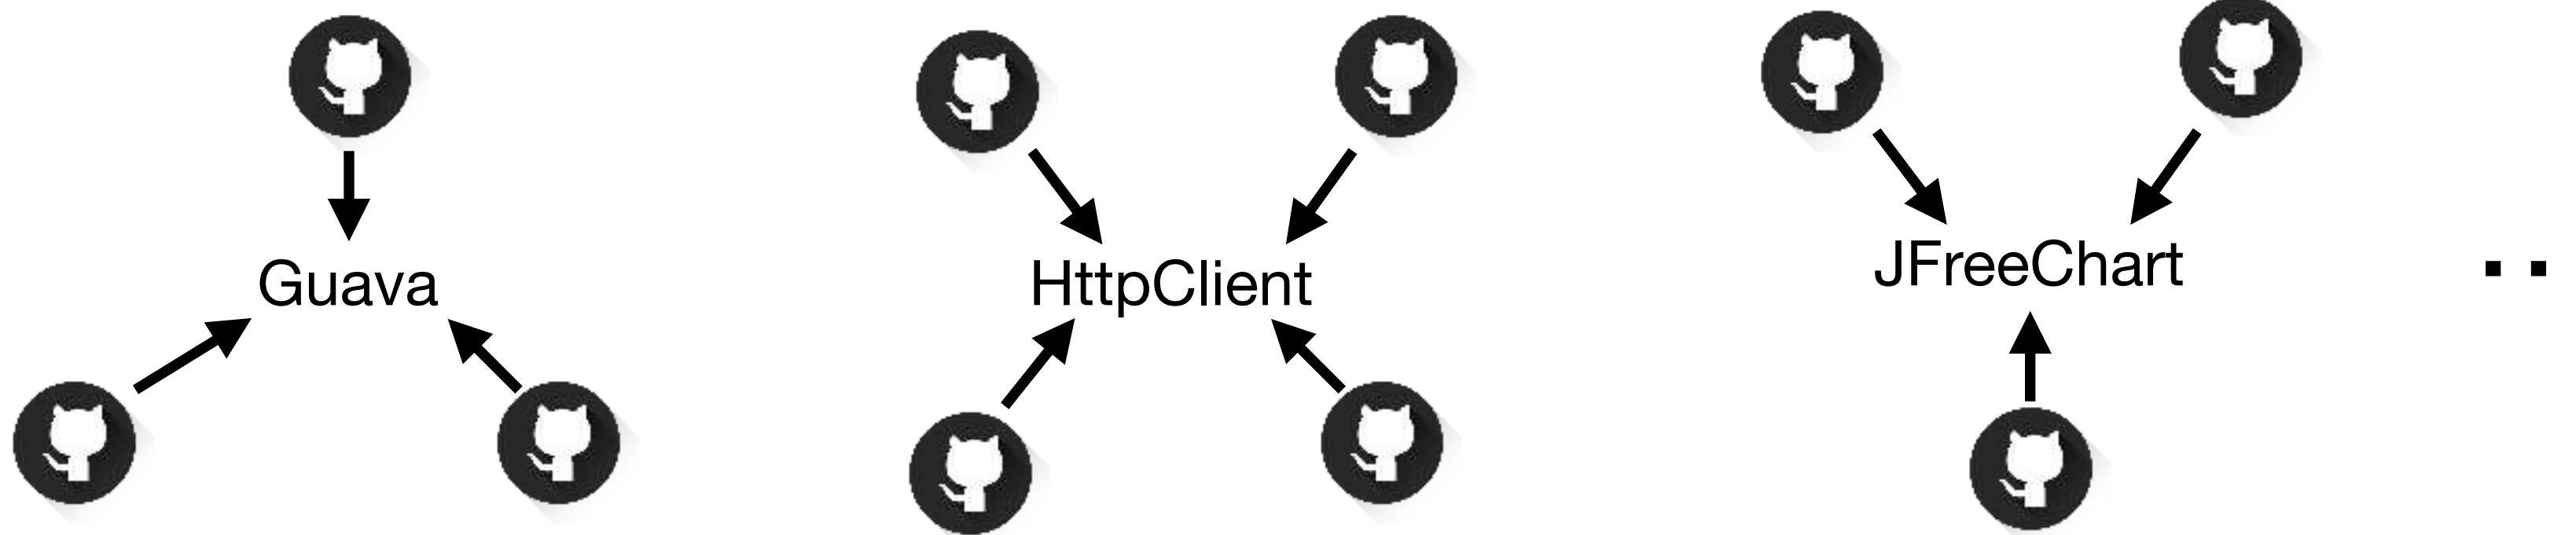
RQ2. How similar are the automatically uncovered features to documented tutorial features?

TABLE III ANALYSIS OF THE MATCHED FEATURES PER LIBRARY.							
Library	No. Feat.	Match.	Not Found	Acc.	R-Acc.	Relv.	Over.
Guava	30	22	7	0.73	0.97	0.71	-0.28
HttpClient	12	9	3	0.75	1.0	0.71	-0.23
JFreeChart	9	8	0	0.89	0.89	0.73	-0.26
JSoup	15	11	0	0.73	0.73	0.95	-0.51
PDFBox	15	8	4	0.53	0.73	0.82	-0.30
Apache-POI	20	10	3	0.50	0.59	0.75	-0.39
Quartz	22	12	5	0.55	0.71	0.65	-0.58

On average, our uncovered features are very similar to those in tutorials denoted by the relevance score.

Evaluation

RQ3. To which extent do the uncovered features that do not match documented tutorial features correspond to actual API usage in client projects?



Evaluation

RQ3. To which extent do the uncovered features that do not match documented tutorial features correspond to actual API usage in client projects?

TABLE V
NEWLY MATCHED FEATURES FROM GITHUB CLIENT PROJECTS.

Library	No. Feat.	M-RQ2	U-RQ2	M-RQ3	%	Relv.	Over.
Guava	110	22	14	91	95	0.45	0.00
HttpClient	38	9	7	24	77	0.50	0.05
JFreeChart	70	8	5	55	85	0.35	0.02
JSoup	81	11	11	52	74	0.60	-0.21
PDFBox	44	8	7	32	86	0.42	0.01
Apache-POI	81	10	9	53	74	0.50	-0.10
Quartz	31	12	5	19	73	0.48	0.00

The majority of our uncovered features are also used in GitHub client projects.

Evaluation

RQ3. To which extent do the uncovered features that do not match documented tutorial features correspond to actual API usage in client projects?

TABLE V
NEWLY MATCHED FEATURES FROM GITHUB CLIENT PROJECTS.

Library	No. Feat.	M-RQ2	U-RQ2	M-RQ3	%	Relv.	Over.
Guava	110	22	14	91	95	0.45	0.00
HttpClient	38	9	7	24	77	0.50	0.05
JFreeChart	70	8	5	55	85	0.35	0.02
JSoup	81	11	11	52	74	0.60	-0.21
PDFBox	44	8	7	32	86	0.42	0.01
Apache-POI	81	10	9	53	74	0.50	-0.10
Quartz	31	12	5	19	73	0.48	0.00

The decrease in the relevance might be caused by a less-focused usage of the features in GitHub.

Examples of uncovered features

Good features

```
// convert image, convert pdf  
PDDocument.load(...);  
PDDocument.getPage(...);  
PDPageContentStream.drawImage(...);  
PDDocument.save(...);  
// merge file, reuse PDFMergerUtility, ...  
PDFMergerUtility.mergeDocuments(...);
```

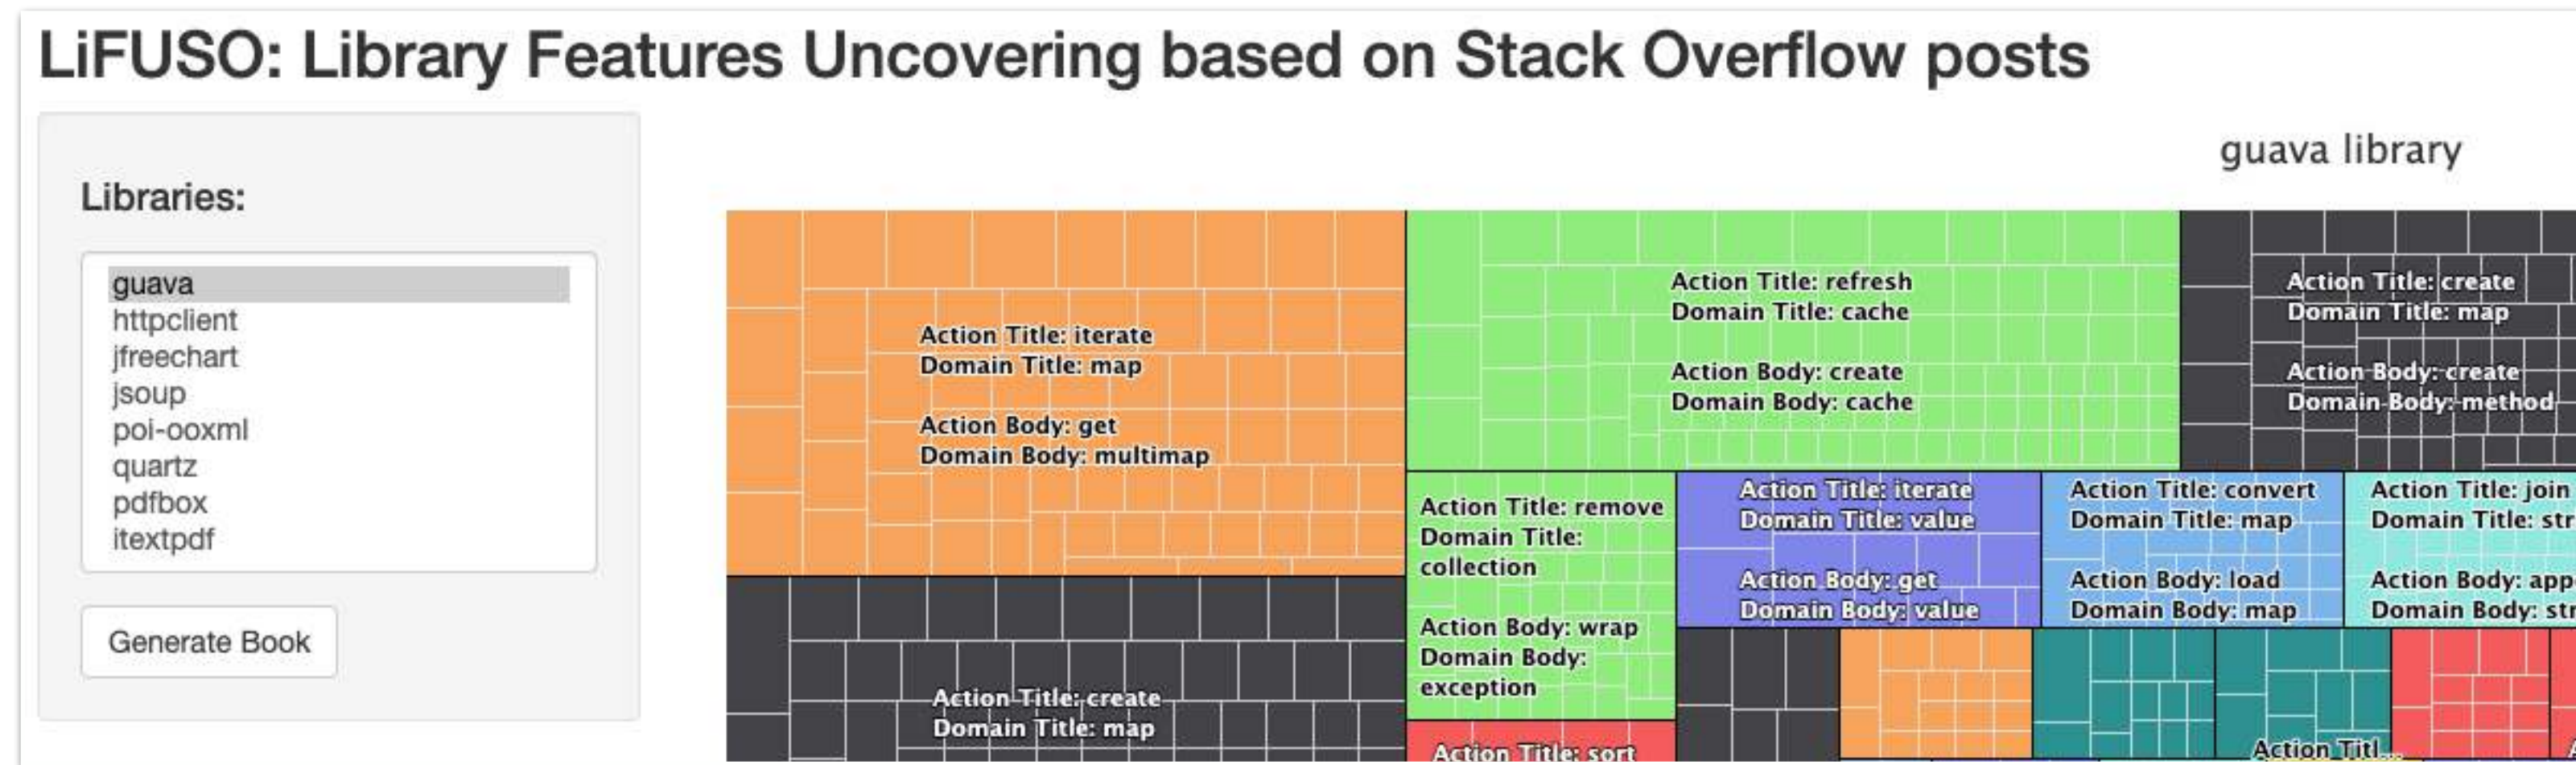
PDFBox

Bad features

```
// Feature from HttpClient  
CloseableHttpResponse.close();  
// Feature from JFreeChart  
CategoryPlot.setAxisOffset(...);  
// Feature from Apache-POI  
POIXMLDocument.hasOOXMLHeader(...);  
// Feature from Quartz  
JobDetail.equals(...);
```

Several libraries

Tool LiFUSO

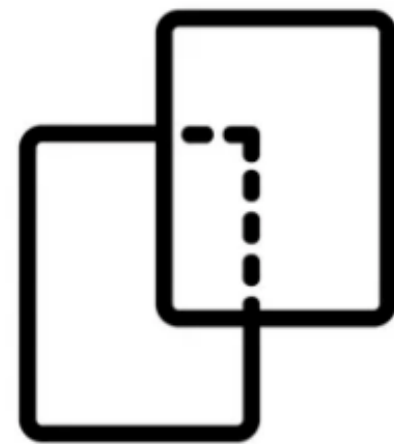
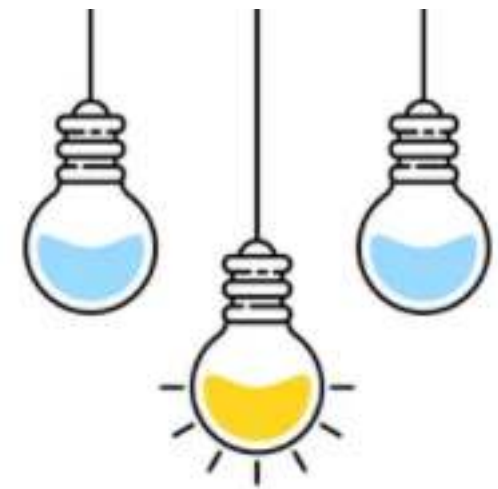


<https://github.com/softwarelanguageslab/lifuso>



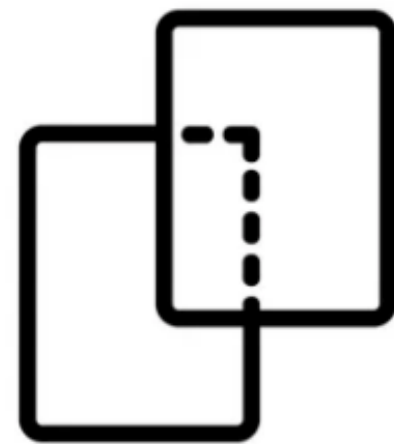
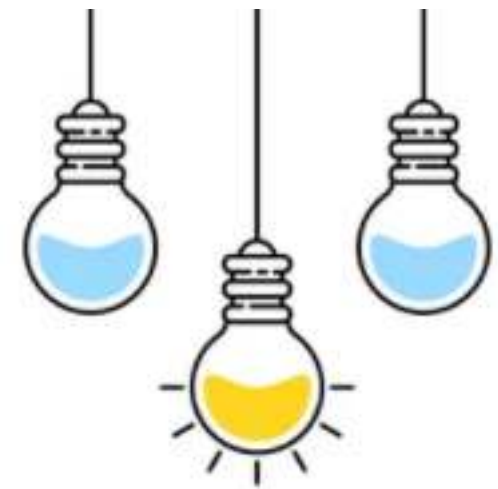
Potential application

Maintainers



Potential application

Maintainers



Users



Conclusion

Introduction

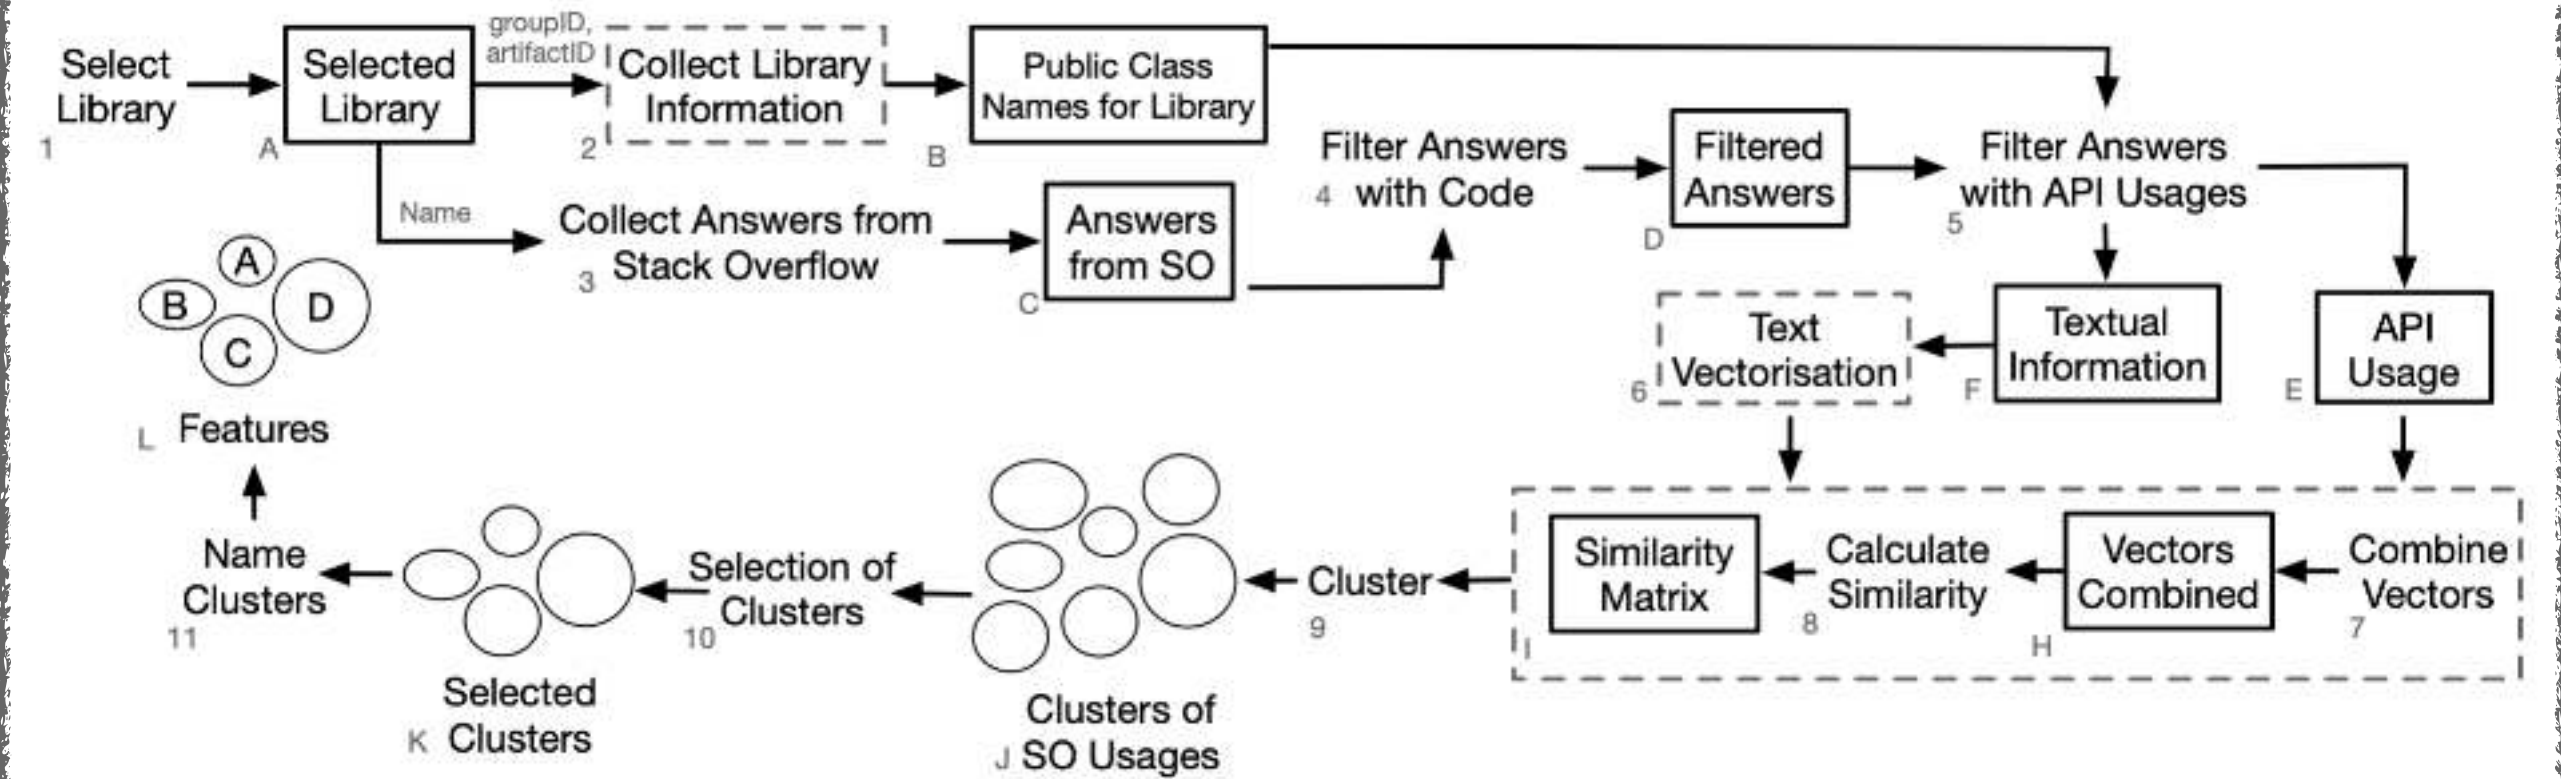
JSoup library

```
String text = EntityUtils.toString(getResponse.getEntity());
Document doc = Jsoup.parse(text);
Element form = doc.select("form").first();

for (Element input : form.select("input")) {
    String value = input.attr("value");
}
```

Aspect Oriented	Collections
Actor Frameworks	Configuration Libraries
Application Metrics	Core Utilities
Build Tools	Date and Time Utilities
Bytecode Libraries	Dependency Injection
Command Line Parsers	Embedded SQL Databases
Cache Implementations	HTML Parsers
Cloud Computing	HTTP Clients
Code Analyzers	I/O Utilities
	...

Approach



Evaluation

RQ1. Which combination of SO answer attributes produces the most cohesive clusters?

TABLE II TOP-K SCORES FOR ALL COMBINATIONS OF LIBRARY ATTRIBUTES BETWEEN PARENTHESES IS THE OBTAINED SCORE.			
Library	K=1	K=2	K=3
Guava	Methods (0.6)	Methods CC (0.6)	Methods + Methods CC (0.59)
HttpClient	API Calls (0.46)	Methods (0.45)	Methods CC (0.45)
JPicClient	Methods (0.38)	Methods CC (0.38)	Methods + Methods CC (0.38)
JSoup	API Calls (0.61)	Methods CC (0.51)	Methods (0.51)
PDFBox	Methods + Methods CC (0.37)	Methods CC (0.36)	Methods (0.36)
Apache-POI	Methods (0.52)	Methods CC (0.52)	Methods + Methods CC (0.51)
Quartz	Methods + Methods CC (0.45)	Methods CC (0.42)	API Calls (0.42)

The **Methods** combination from the API usage information attributes produced the most cohesive clusters.

RQ2. How similar are the automatically uncovered features to documented tutorial features?

TABLE III ANALYSIS OF THE MATCHED FEATURES PER LIBRARY.							
Library	No. Feat.	Match.	Not Found	Acc.	R-Acc.	Rel.	Over.
Guava	30	22	7	0.73	0.97	0.71	-0.28
HttpClient	12	9	3	0.75	1.0	0.71	-0.25
JPicClient	9	8	0	0.89	0.89	0.73	-0.26
JSoup	15	11	0	0.73	0.73	0.95	-0.51
PDFBox	15	8	4	0.53	0.73	0.82	-0.50
Apache-POI	20	10	3	0.50	0.59	0.75	-0.59
Quartz	22	15	5	0.55	0.71	0.65	-0.58

When considered against the usages in the SO data, the matched accuracy increases.

RQ3. To which extent do the uncovered features that do not match documented tutorial features correspond to actual API usage in client projects?

TABLE V NEWLY MATCHED FEATURES FROM LIBRARY CLIENT PROJECTS							
Library	No. Feat.	M-Q1	M-Q2	M-Q3	%	Rel.	Over.
Guava	110	22	14	91	58	0.55	0.00
HttpClient	58	9	7	54	55	0.50	0.05
JPicClient	20	8	5	25	85	0.55	0.00
JSoup	81	11	1	52	74	0.60	-0.13
PDFBox	22	8	7	23	66	0.52	0.00
Apache-POI	81	10	9	51	54	0.50	-0.10
Quartz	54	12	5	19	35	0.48	0.00

TABLE VI OVERALL PERFORMANCE OF LIFUSO							
Guava HttpClient JPicClient JSoup PDFBox POI Quartz							
Guava	0.58	0.50	0.51	0.52	0.57	0.53	0.54
Methods	0.58	0.58	0.55	0.57	0.57	0.54	0.54

A closer look to the overflow metric indicates that GitHub usages might encompass more than one of our uncovered features.

Examples of uncovered features

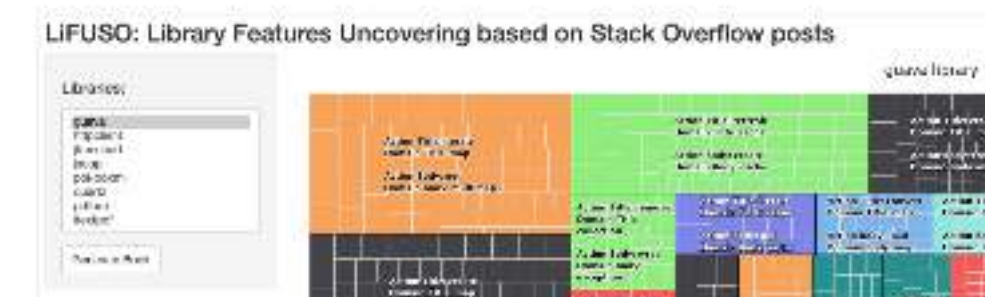
Good features

PDFBox

Bad features

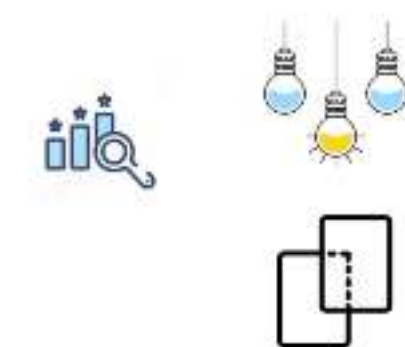
Several libraries

Tool LIFUSO



Potential application

Maintainers



Users

